

Name: Fan Feng

Assignment: 3

Course: CPSC 433/533

- [25 pt in total] Protocol and functions
 - a. [5 pt] Method: HTTP 1.0 GET method (URL mapping)
 - b. [5 pt] Header: Last-Modified, If-Modified and User-Agent
 - c. [10 pt in total] Support CGI program
 - c.1. [4 pt] Detect executable
 - c.2. [6 pt] Setup Environment variables and direct I/O
 - d. [5 pt in total] Heartbeat monitoring
- [40 pt in total] Implementation structure
 - a. [3 pt] Sequential
 - b. [3 pt] Per request thread
 - c. [3 pt] Thread pool with shared welcome thread
 - d. [3 pt] Thread pool with a shared queue and busy wait
 - e. [3 pt] Thread pool with a shared queue and suspension
 - f. Asynchronous I/O using select [12 pt]
 - g. Asynchronous I/O using Java 7 Async channel future/listener [13 pt]
- [15 pt in total] Performance benchmarking
 - a. [6 pt] Clear description of benchmarking methodology
 - b. [9 pt] Best throughput exceeds 10 Mbps
- [20 pt in total] Comparison and exploration of other designs
 - a. 2 points for each question
- [35 pt in total potential] Bonus points
 - a. [10 pt] Beat or close to Apache performance
 - b. [25 pt] The fastest server

Part 1a:

SHTTPTestClient.java

RequestSender.java

Statistics.java

Part 1b:

Configuration.java

FileCache.java

HTTPServerSequential.java

HTTPServerPerThread.java

RequestHandler.java

HeartbeatMonitor.java

Design of heartbeat monitoring:

Monitor <MonitorClassName> is an optional parameter in config.txt.

In RequestHandler.java, when mapping URL to file, if a request queries a virtual URL (/load), the server randomly accepts/rejects the request with a pre-defined acceptance rate.

If a request is accepted, the server responds with status code 200 and “Heartbeat Monitor” as response body. Otherwise, status code 503 is returned.

See mapURL2File() in RequestHandler.java and HeartbeatMonitor.java for details.

Part 2a:

Thread pool with service threads competing on welcome socket:

HTTPServerSharedWelcomeSocket.java

ServiceThreadSharedWelcomeSocket.java

Thread pool with shared queue busy wait:

HTTPServerSharedQueueBusyWait.java

ServiceThreadSharedQueueBusyWait.java

Thread pool with shared queue suspension:

HTTPServerSharedQueueSuspension.java

ServiceThreadSharedQueueSuspension.java

Part 2b:

Part 2c:

Part 2d: Comparison of Designs

(1): Comparison with xsocket

(i).

<http://xsocket.sourceforge.net/core/tutorial/V2/TutorialCore.htm>:

xsocket allows (number of CPUs + 1) dispatchers.

Line 105 in IoSocketDispatcherPool.java:

The dispatchers share workload in a round-robin fashion.

(ii).

Line 208 – 258 in IoSocketDispatcher.java:

While the dispatcher is open, the selector times out for 5 seconds, performs the register handler tasks, and calls `handleReadWriteKey()` on new events. After that all socket handlers are deregistered, and the selector is closed.

(iii).

In `EchoServerTest.java`: `EchoServer server = new EchoServer(0);`

In `EchoServer.java`: `private IServer server = null;`

`IHandler hdl = new EchoHandler();`

`server = new Server(listenPort, hdl);`

Line 210 in `HandlerAdapter.java`:

`performOnData(connection, taskQueue, ignoreException, handler);`

Line 242 in `HandlerAdapter.java`: `handler.onData(connection);`

(iv).

xsocket uses `IdleTimeoutHandler` to implement idle timeout of a connection.

(v).

The library tests one class at a time. Specifically, `EchoServerTest.java` tests server side and both sides respectively.

We can test the server with idle timeout by establishing a connection to it without sending anything. The server should be able to kill the idle connection automatically.

(2): Comparison with Netty

(i). <https://netty.io/wiki/user-guide-for-4.x.html>

There are two event loop groups. The first one, often called 'boss', accepts an incoming connection. The second one, often called 'worker', handles the traffic of the accepted connection once the boss accepts the connection and registers the accepted connection to the worker.

Synchronization among them is achieved by multithreading.

(ii). <https://netty.io/4.0/api/io/netty/channel/ChannelFuture.html>

`sync()` waits for this future until it is done. One can implement it using the `isDone()` method from the superinterface `java.util.concurrent.Future`.

(iii). <https://github.com/netty/netty/wiki/Using-as-a-generic-library>

One key difference between `ByteBuf` and `ByteBuffer` is that `ByteBuf`'s life cycle is bound to its reference count, whereas `ByteBuffer`'s life cycle is not, and it relies on Java's garbage collection.

(iv).

<https://netty.io/4.0/xref/io/netty/example/http/helloworld/HttpHelloWorldServerInitializer.html>

<https://netty.io/4.0/xref/io/netty/example/http/snoop/HttpSnoopServerInitializer.html>

HTTP Hello World Server: (SslHandler), HttpServerCodec, HttpHelloWorldServerHandler

HTTP Snoop Server: (SslHandler), HttpRequestDecoder, HttpResponseEncoder, HttpSnoopServerHandler

(v).

<https://netty.io/5.0/api/io/netty/channel/ChannelPipeline.html>

A ChannelPipeline contains a list of Channel Handlers which handles or intercepts inbound events and outbound operations of a Channel.

An I/O event is handled by a ChannelHandler and is forwarded by the ChannelHandler which handled the event to the ChannelHandler which is placed right next to it.

A ChannelHandler can also trigger an arbitrary I/O event if necessary. To forward or trigger an event, a ChannelHandler calls the event propagation methods defined in ChannelHandlerContext.


Part 2e: Performance Benchmarking

Benchmarking methodology:

For each server, run different number of threads (1, 2, 3, 4, 5, 10, 15, 20, 30, 40, 50, 60) for 60 seconds. Record the data throughput, and plot data throughput against number of threads.

Run Apache server on zoo by making -server zoo.cs.yale.edu, and changing URLs to follow /classes/cs433/web/www-root/html-small/doc1.html. Add data throughput against number of threads to the above plot.

Conclusion:



Sequential server and per-thread server perform the worst. The three thread pool servers have better performance, and perform roughly the same.

Data throughput in the plot is in bytes per second. The best data throughput equals about 750 KB/s = 6Mbps.