Network Applications: Email Security, DNS

Y. Richard Yang

http://zoo.cs.yale.edu/classes/cs433/

9/18/2018

Outline

- Admin and recap
- Email
  - Basic email systems design
  - Email security
- DNS

Admin

- 72 discretionary late hours for assignments across the semester

Recap: The Big Picture of the Internet

- Hosts and routers:
  - ~ 1 bill. hosts
  - organized into ~50K networks
  - backbone links 100 Gbps
- Software:
  - datagram switching with virtual circuit support
  - layered network architecture
    - use end-to-end arguments to determine the services provided by each layer
    - the 5-layer hourglass architecture of the Internet

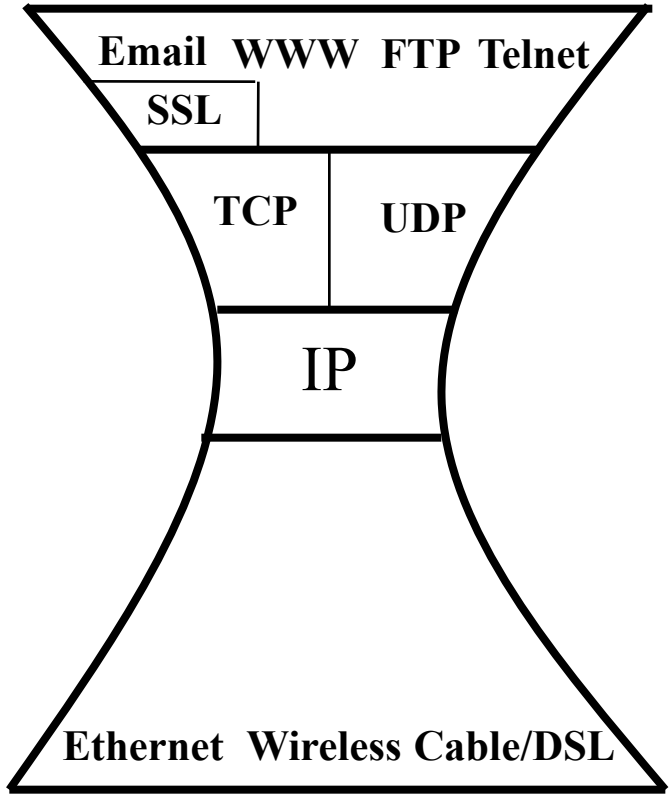

Formats of main protocols

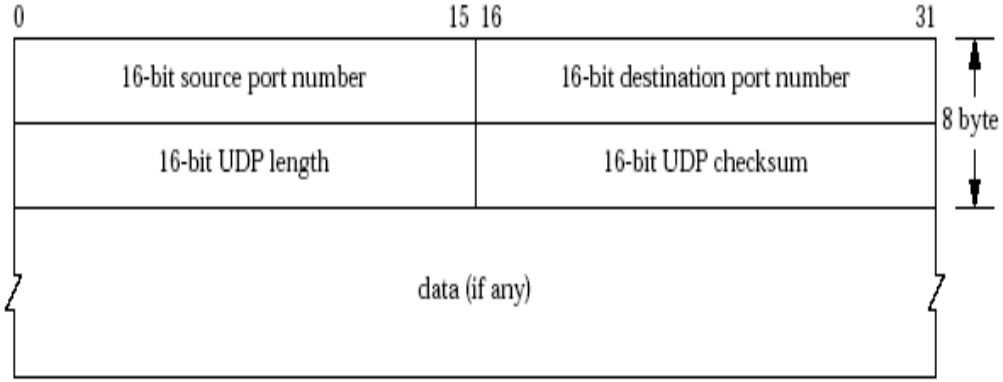

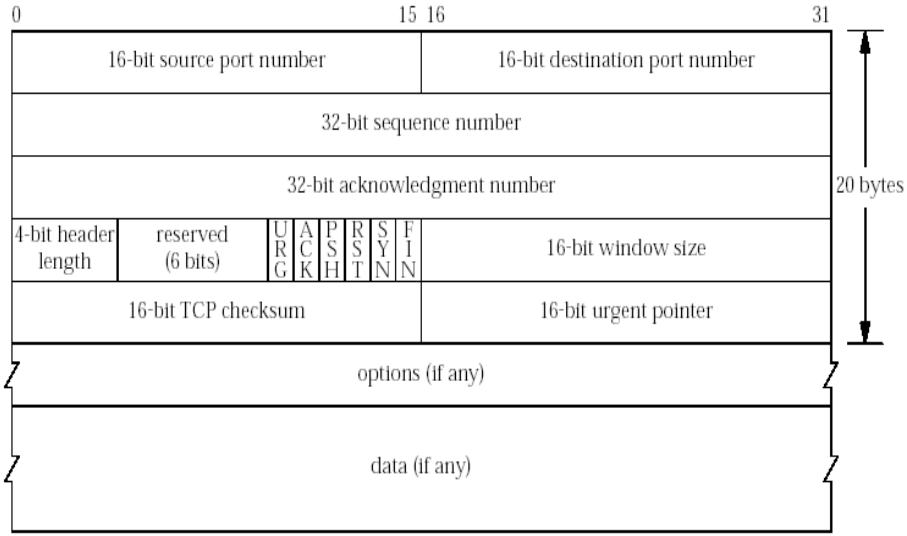

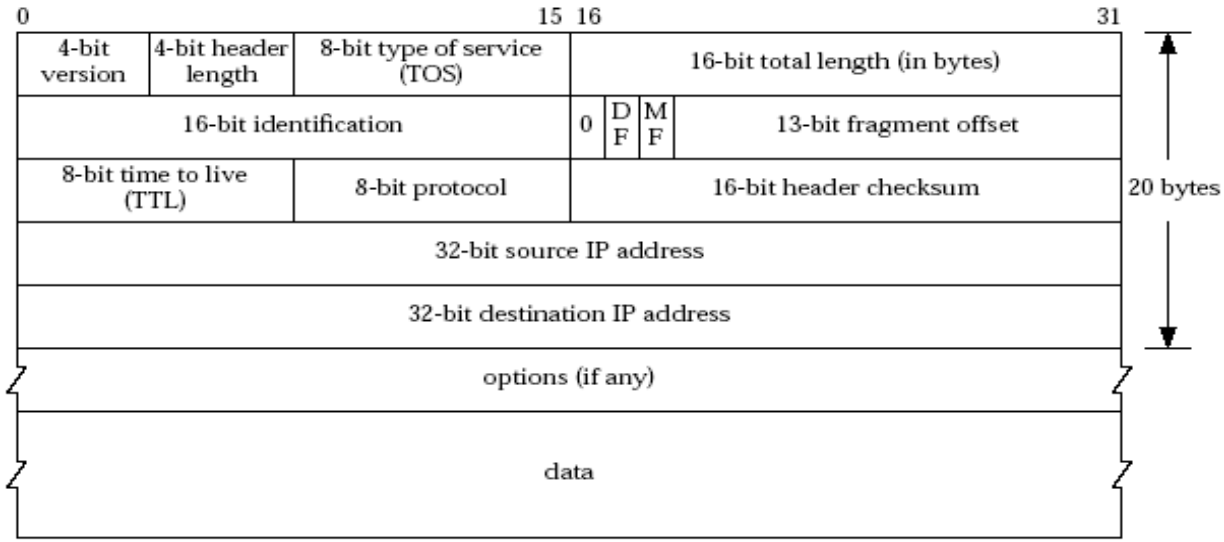

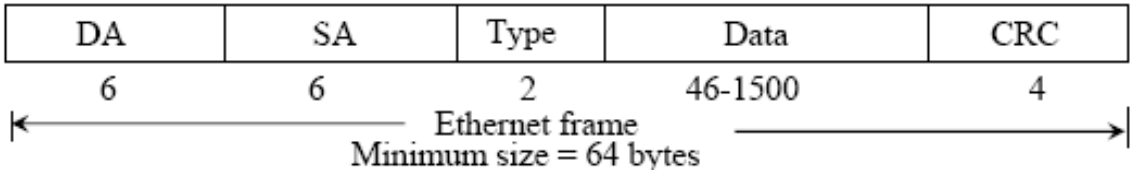

Recap: Client-Server Paradigm

- The basic paradigm of network applications is the client-server (C-S) paradigm

- Some key design questions to ask about a C-S application:
  - extensibility
  - scalability
  - robustness
  - security

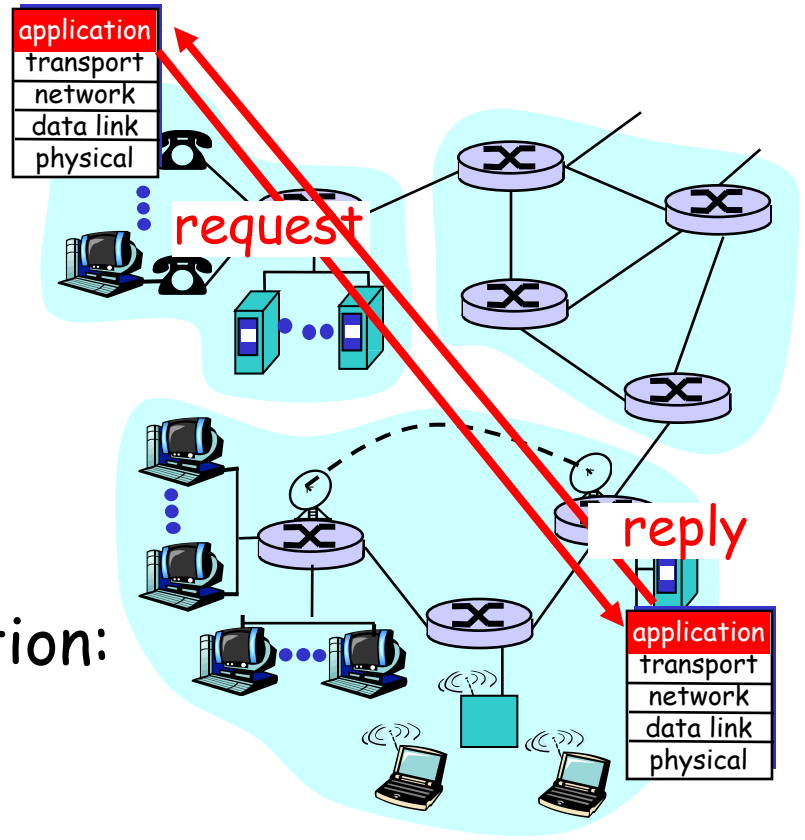

Recap: Email Design Features

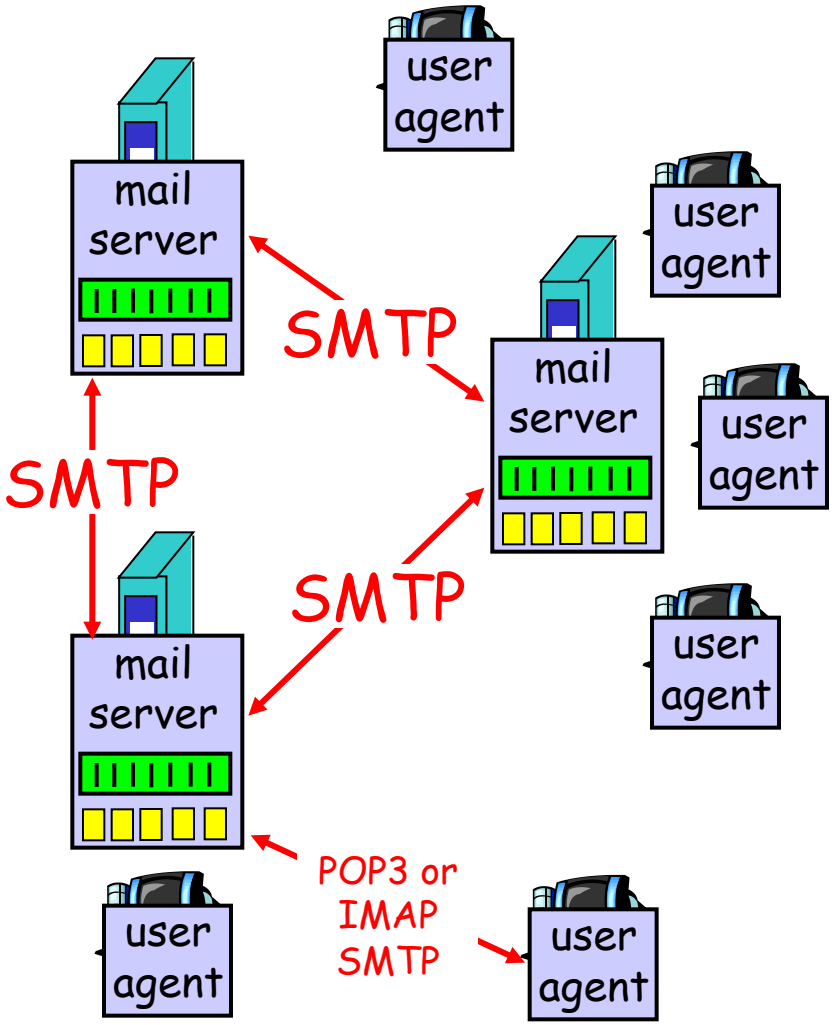

Some key design features of Email

- Separate protocols for different functions
  - email access (e.g., POP3, IMAP)
  - email transport (SMTP)
- A SMTP transaction consists of an envelope and a message body
  - separation of envelope and message body (end-to-end arguments)
    - envelope: simple/basic requests to implement transport control;
    - message body: fine-grain control through ASCII header and message body
      - MIME type as self-describing data type
- Status code in response makes message easy to parse

Recap: Evaluation of SMTP

Key questions to ask about a C-S application

- extensible?
  separate envelope and msg;
  self-describing message;
  ehlo negotiation

- scalable?
  have not seen mechanism yet

- robust?
  have not seen mechanism yet

- security?
  authentication/authorization
  (spoof, spam) are major issues
  of mail transport

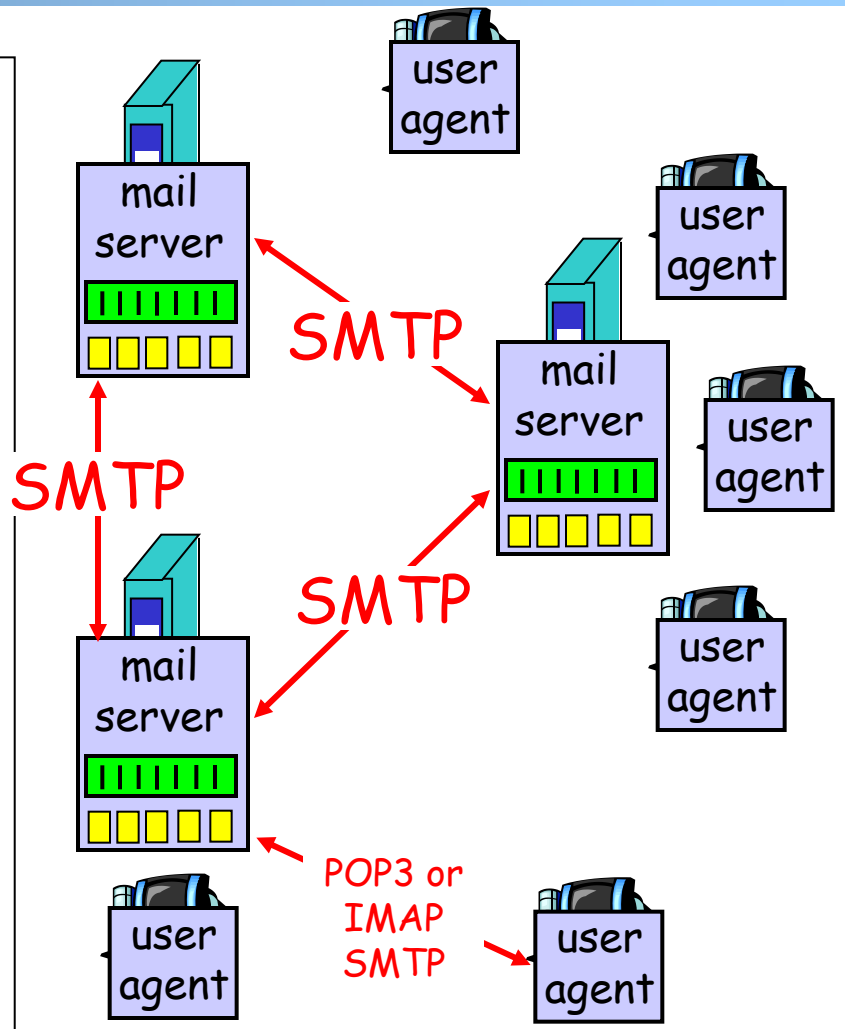

Spam Trend

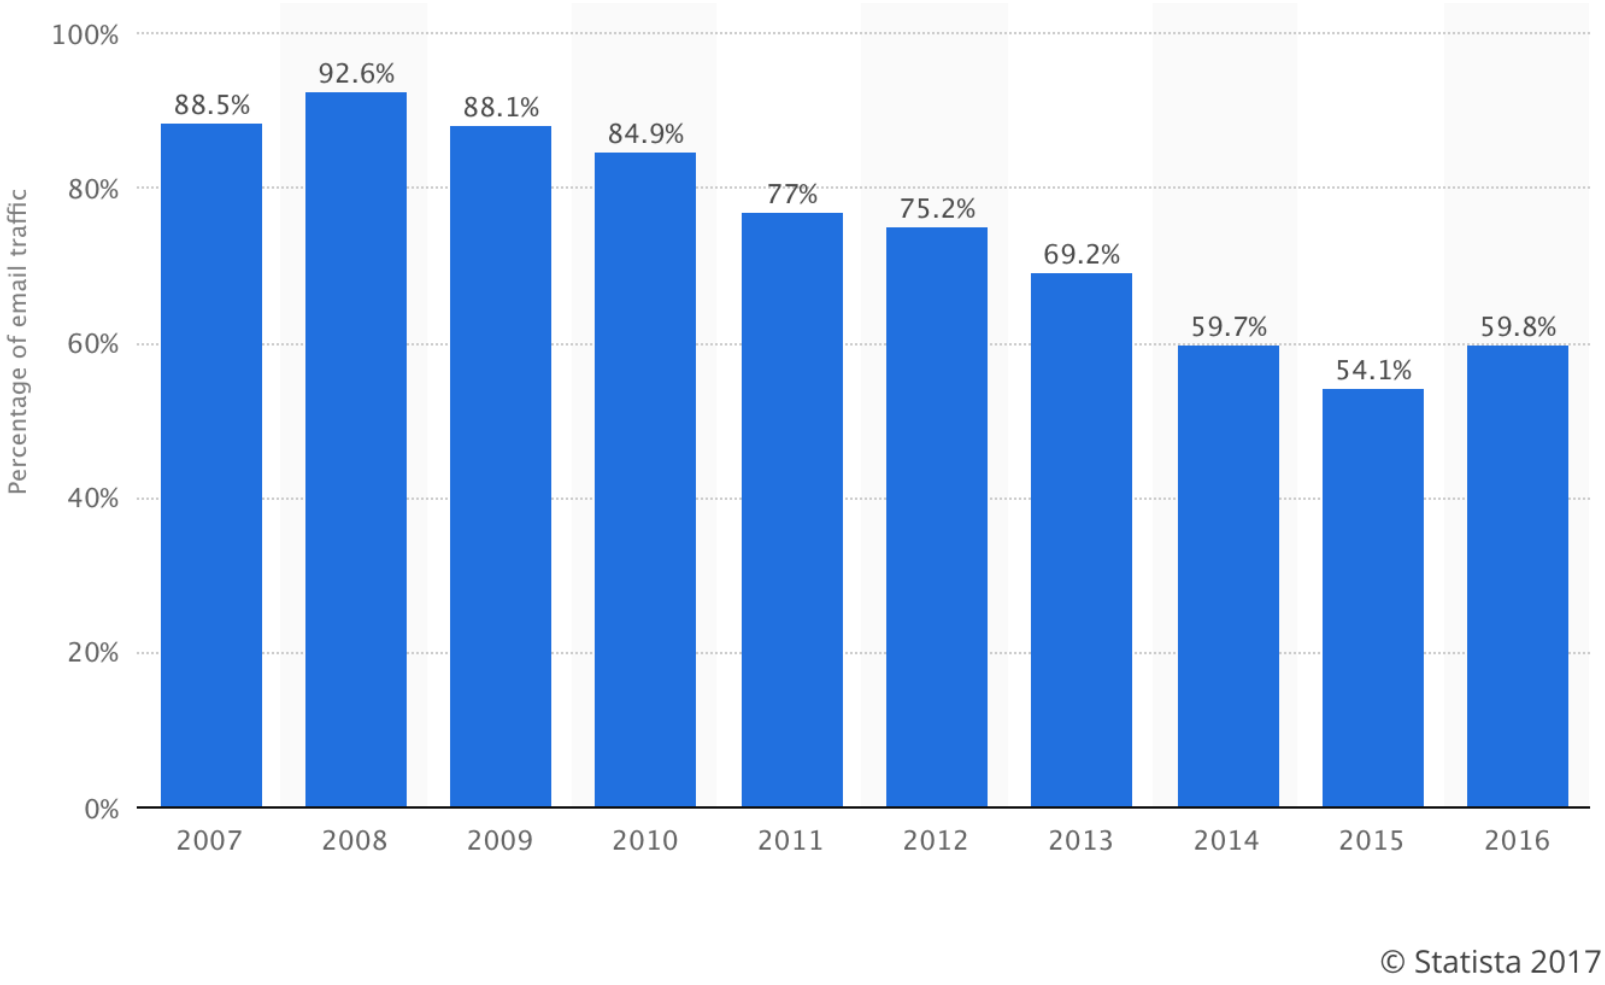

Source: https://www.statista.com/statistics/420400/spam-email-traffic-share-annual/

Current SMTP Authentication Approaches

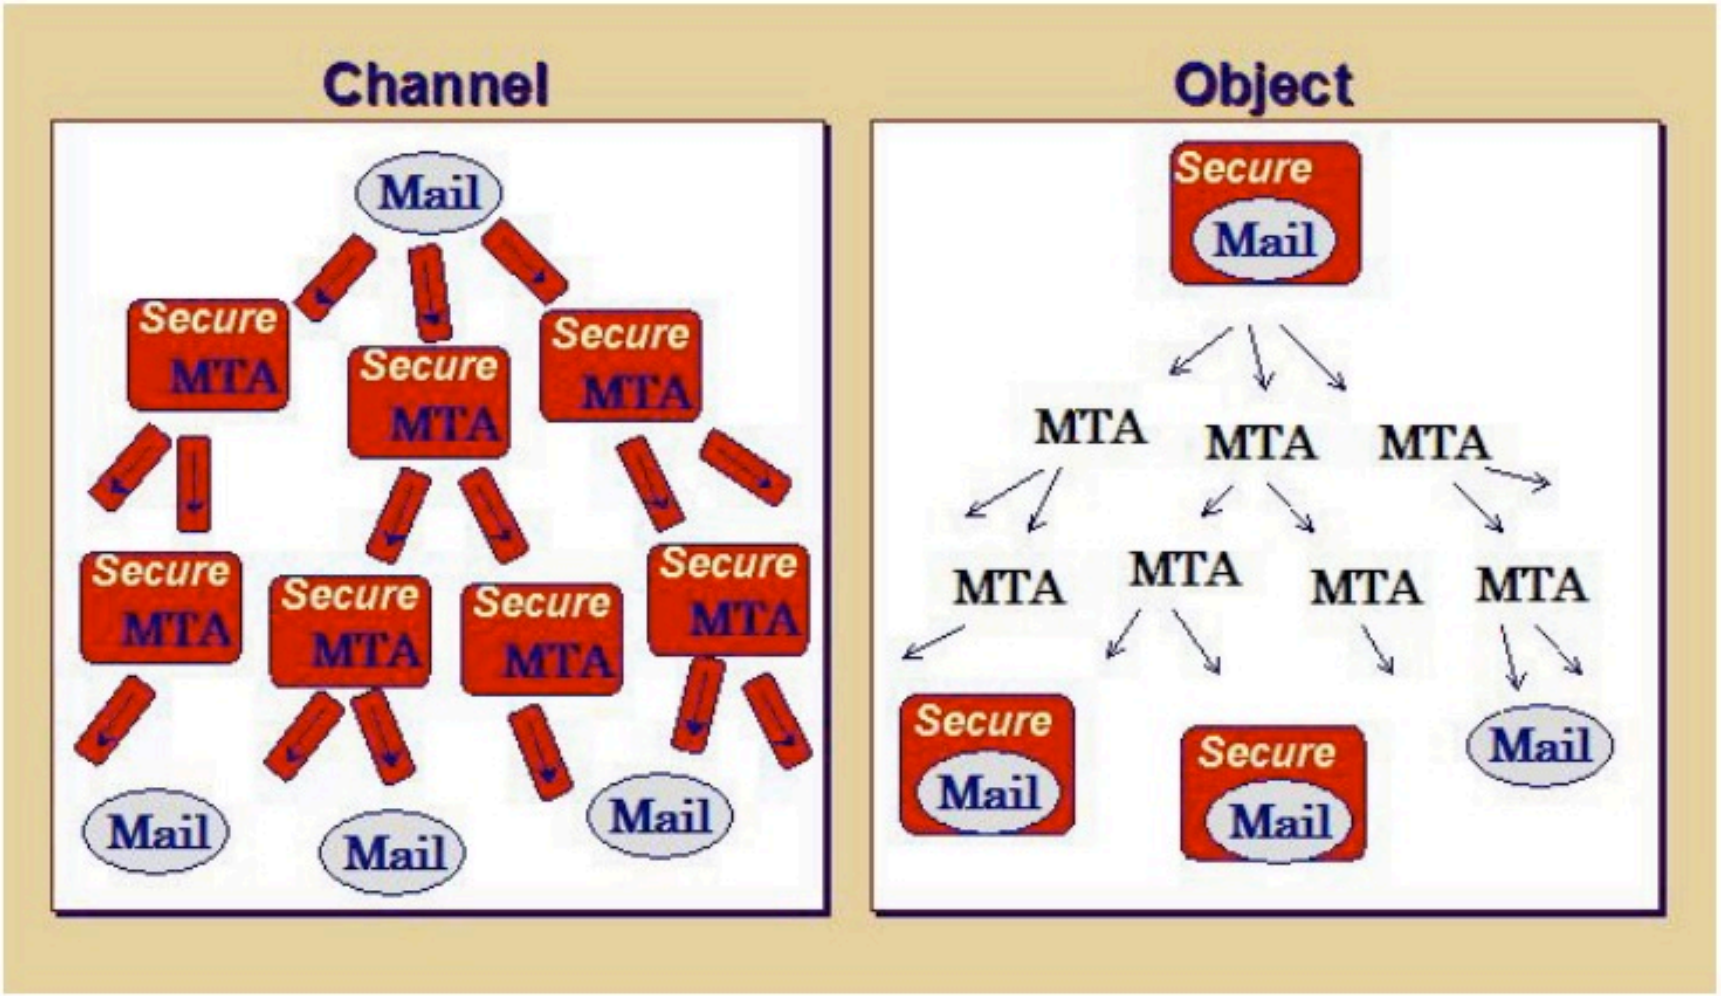

Sender Policy Frame (SPF)

DomainKeys Identified Mail (DKIM)

https://tools.ietf.org/html/rfc7208

Sender Policy Framework (SPF RFC7208)

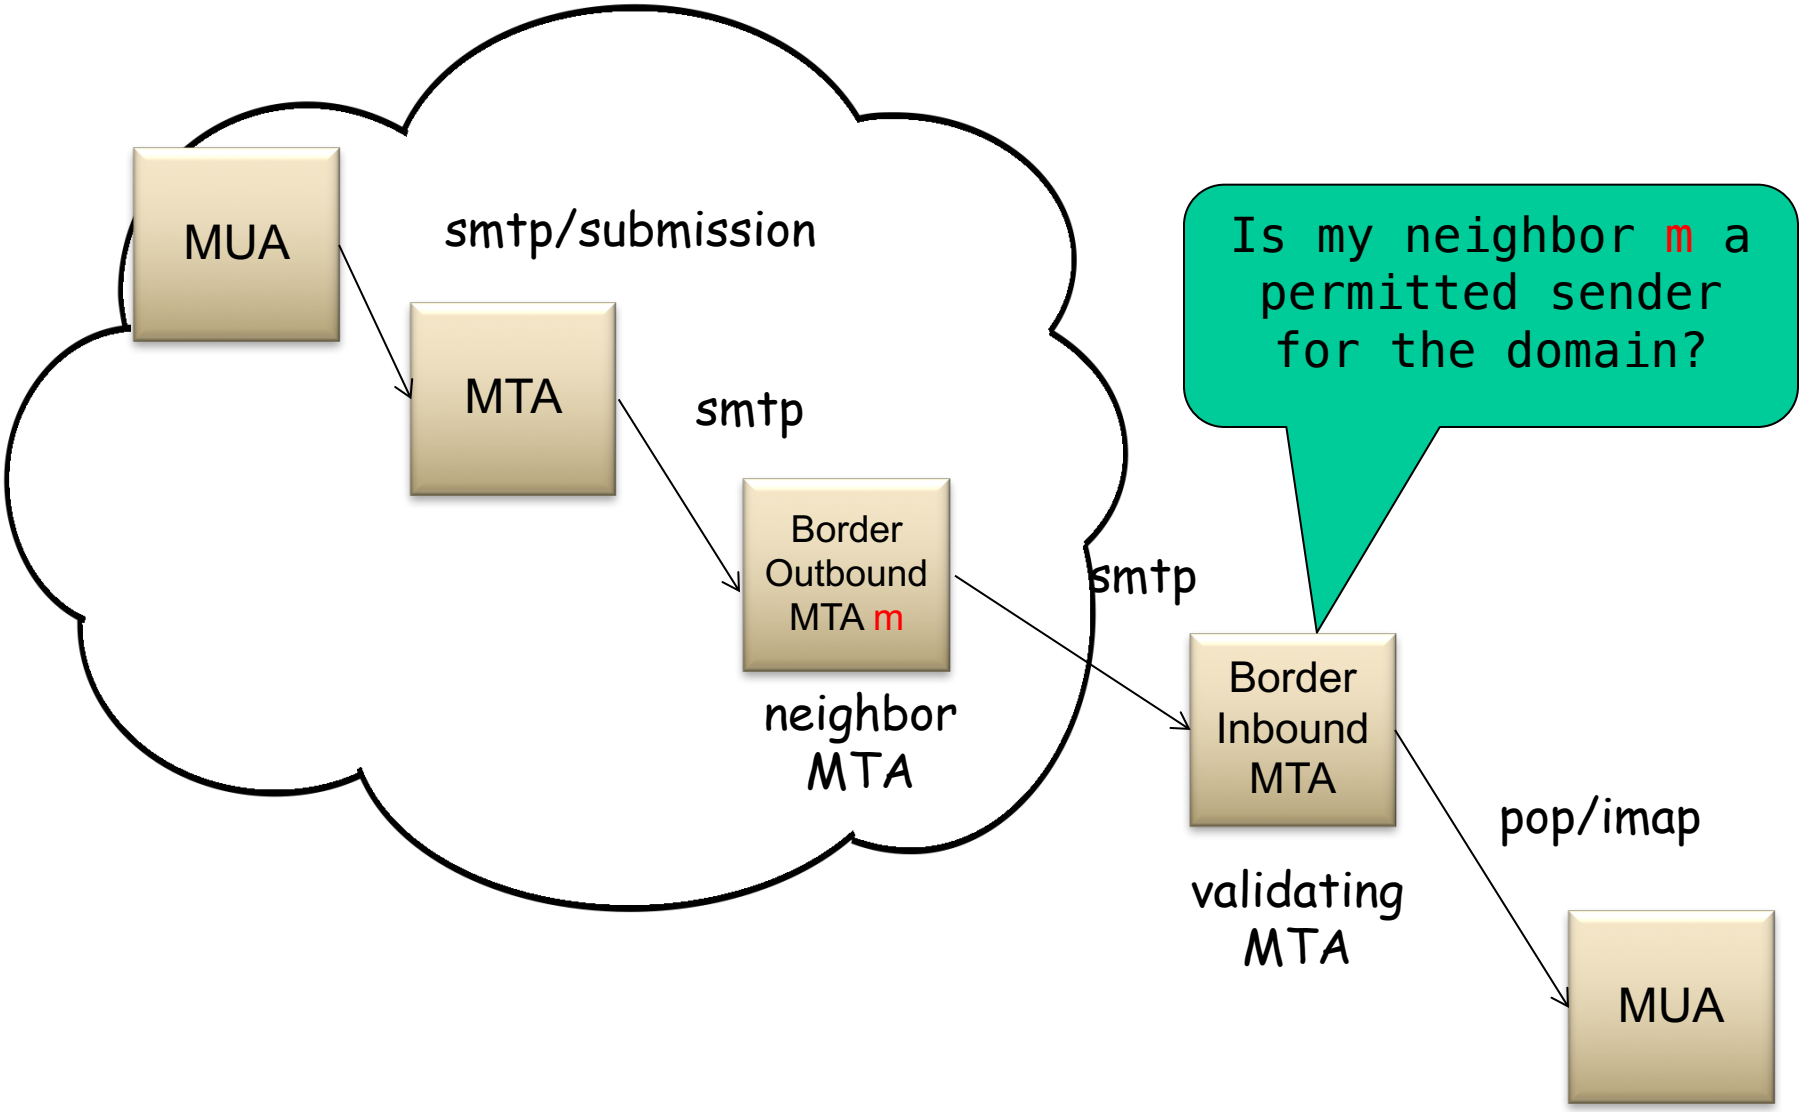

DomainKeys Identified Mail (DKIM)

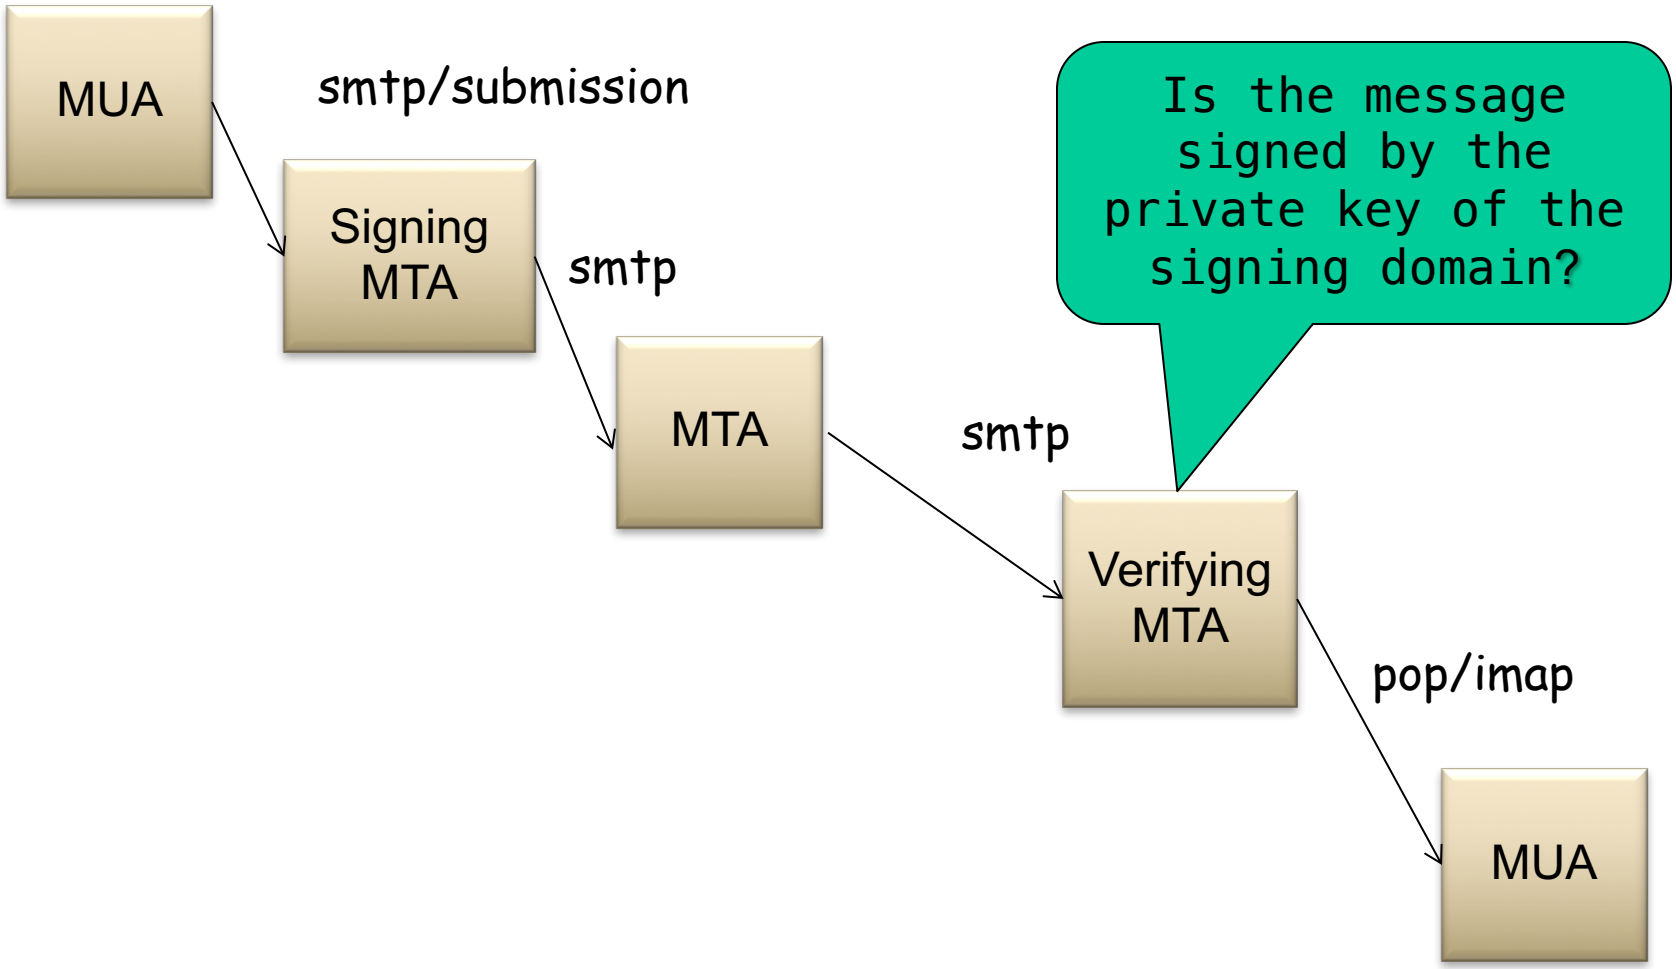

Exercise

- Capture and look at SFP and DKIM in email messages

See pop3-trace.txt

DomainKeys Identified Mail (DKIM; RFC 5585; RFC6376)

- A domain-level digital signature authentication framework for email, using public key crypto, typically RSA
- Basic idea of RSA type public key signature
  - Owner has both public and private keys
  - Owner uses private key to sign a message to generate a signature
  - Others with public key can verify signature
  - Assumption: difficult to get private key even w/ public key distributed

Example: RSA

1. Choose two large prime numbers *p*, *q*. (e.g., 1024 bits each)
2. Compute *n* = *pq*, *z* = *(p-1)(q-1)*
3. Choose *e* (with *e* < *n*) that has no common factors with z. (*e*, *z* are "relatively prime").
4. Choose *d* such that *ed-1* is exactly divisible by *z*. (in other words: *ed* mod *z* = *1*).
5. *Public* key is *(n,e)*. *Private* key is *(n,d)*.

RSA: Signing/Verification

0. Given (n,e) and (n,d) as computed above

1. To sign message, m, compute h = hash(m), then sign with private key

   $s = h^d \bmod n$ (i.e., remainder when h^d is divided by n)

2. To verify signature s, compute

   $h' = s^e \bmod n$ (i.e., remainder when s^e is divided by n)

Magic happens!

$$h = (h^d \bmod n)^e \bmod n$$

The magic is a simple application of Euler's generalization of Fermat's little theorem

Domain-based Message Authentication, Reporting, and Conformance (DMARC) [RFC7489]

- Remaining issue: How to handle unauthenticated messages?

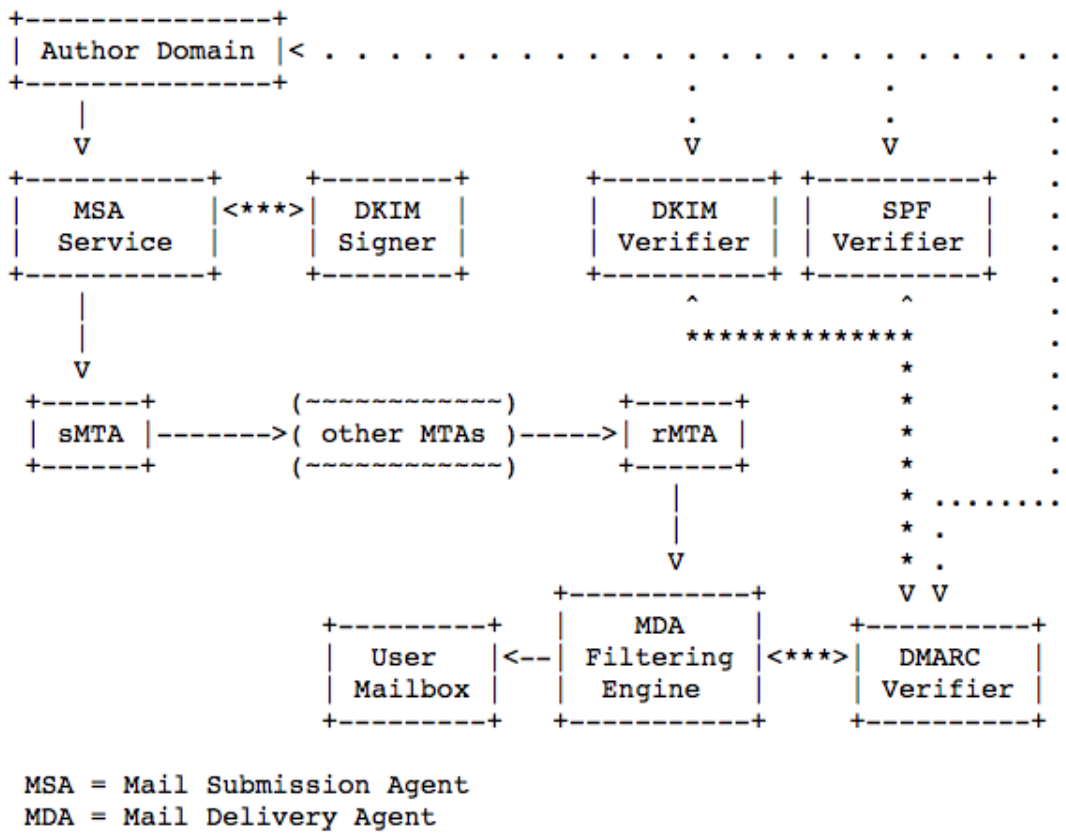

See pop3-trace.txt

Summary: Some Key Remaining Issues about Email

- Basic: How to find the email server of a domain?
- Scalability/robustness: how to find multiple servers for the email domain?
- Security
  - SPF: How does SPF know if its neighbor MTA is a permitted sender of the domain?
  - DKIM: How does DKIM retrieve the public key of the author domain?
  - DMARC: How to find the security policy?

Scalability/Robustness

- Both scalability and robustness require that multiple email servers serve the same email address

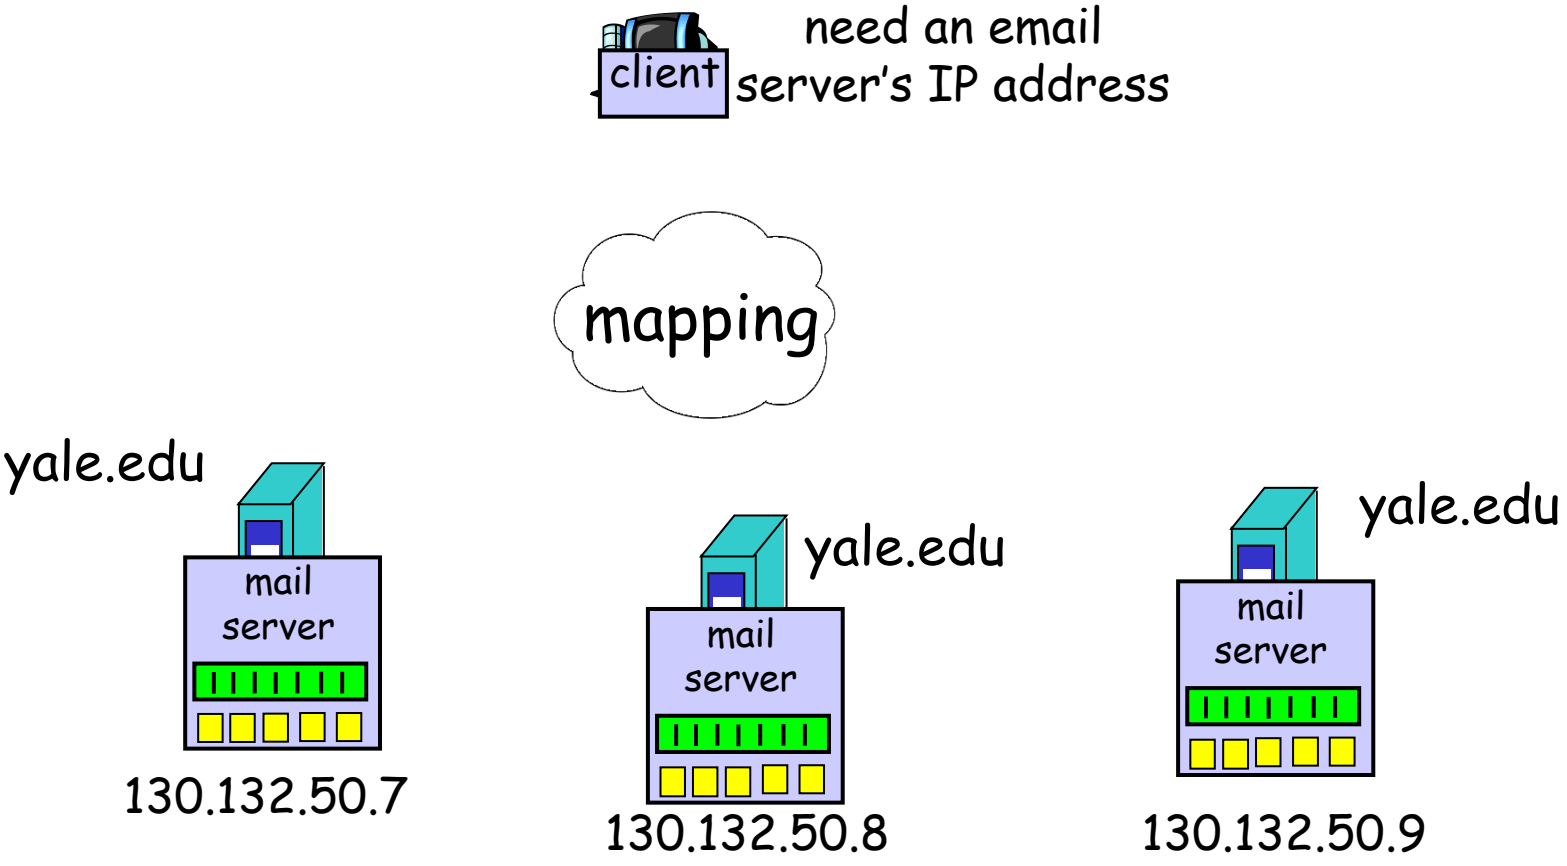

Mapping Functions Design Alternatives

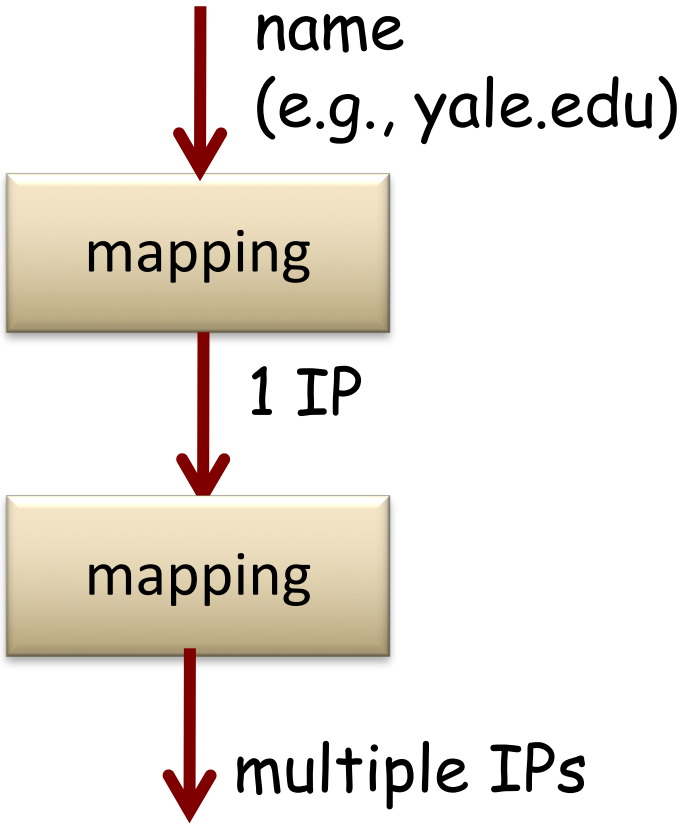

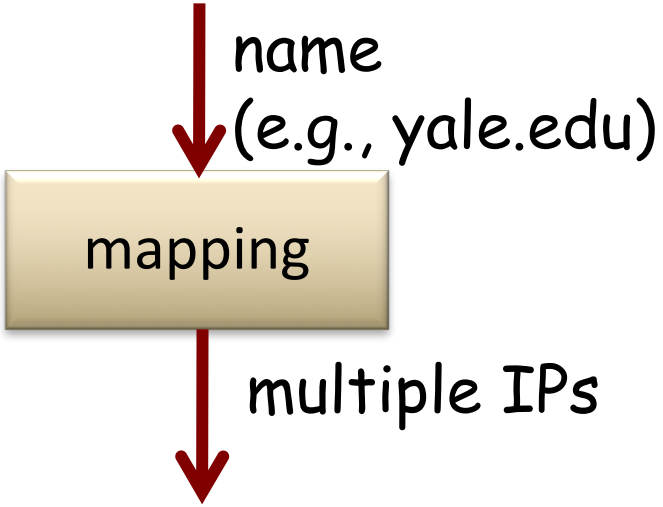

Mapping Functions Design Alternatives

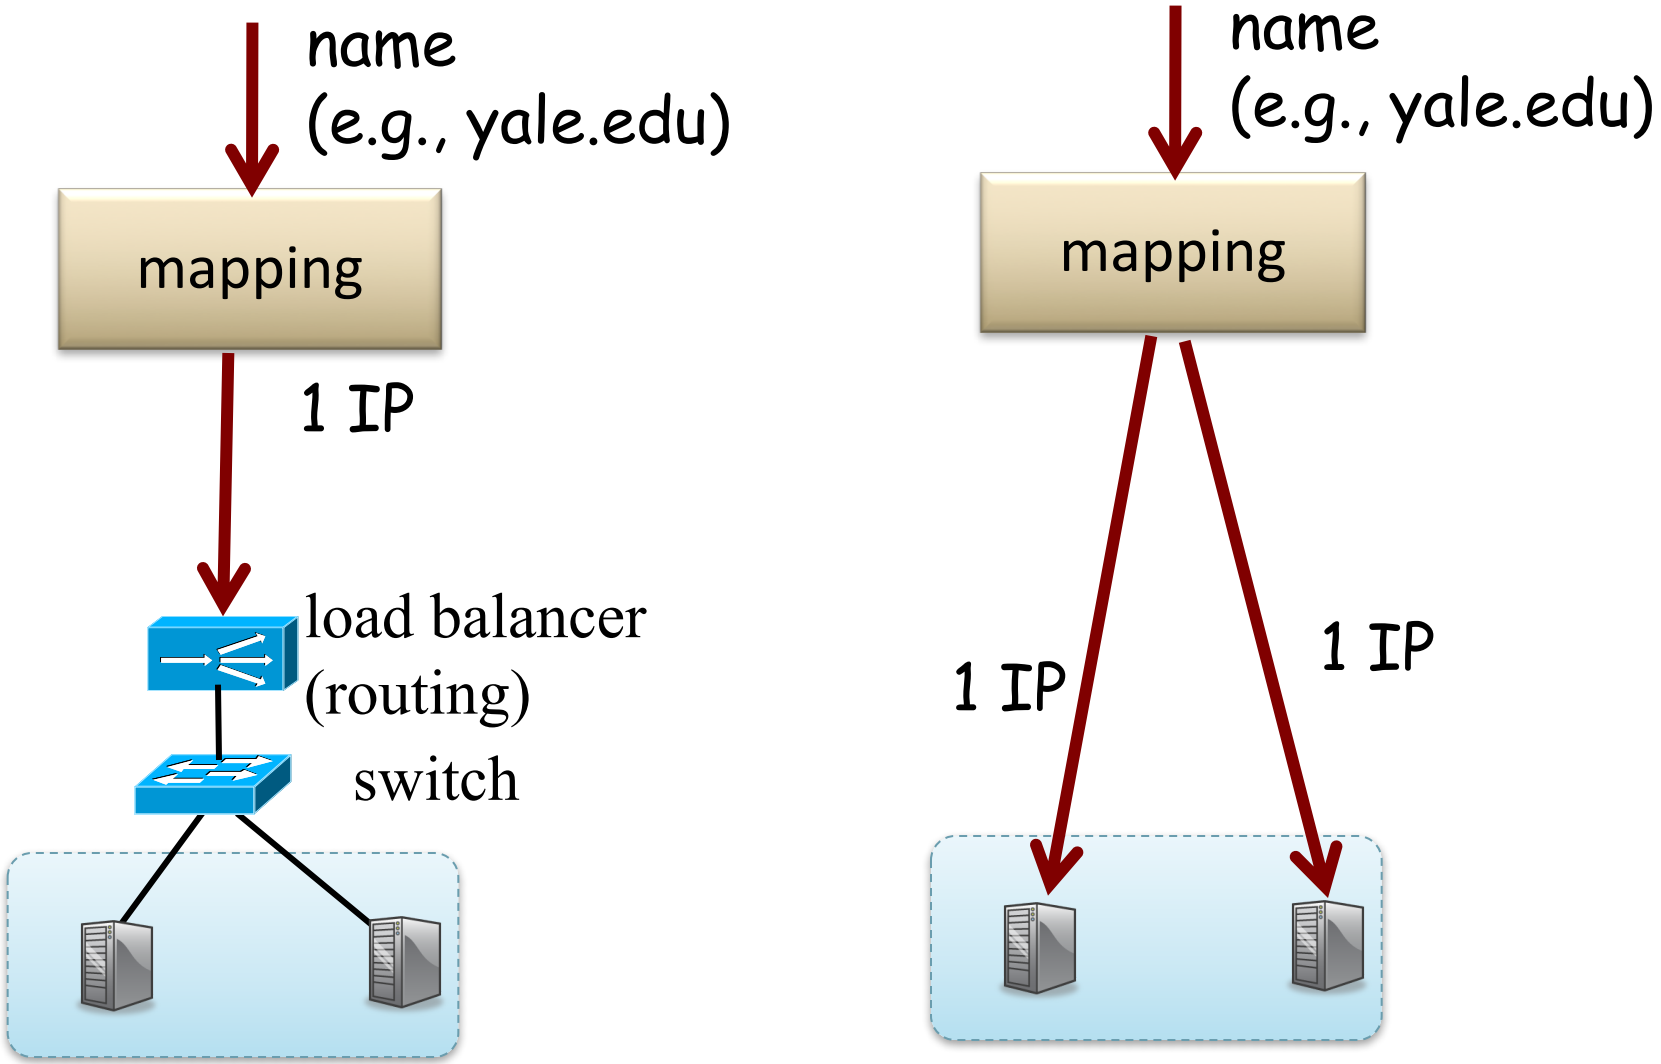

Outline

- Recap
- Email security (authentication)
- DNS

DNS: Domain Name System

- Function
  - map between (domain name, service) to value, e.g.,
    - (www.cs.yale.edu, addr) -> 128.36.229.30
    - (yale.edu, email) -> chai.mail.yale.edu rosehip.mail.yale.edu

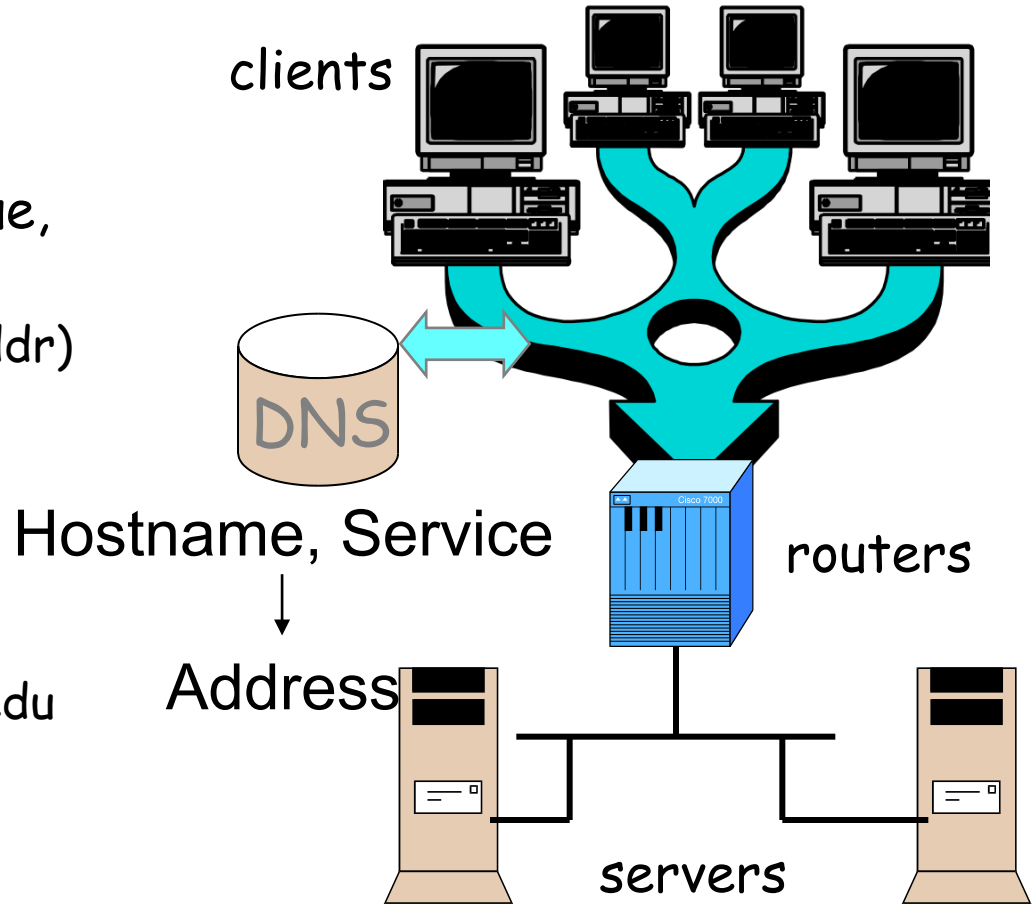

DNS Records

http://www.iana.org/assignments/dns-parameters/dns-parameters.xhtml

DNS: stores resource records (RR)

RR format: **(name, type, value, ttl)**

- Type=A
  - **name** is hostname
  - **value** is IP address
- Type=NS
  - **name** is domain (e.g. yale.edu)
  - **value** is the name of the authoritative name server for this domain
- Type=TXT
  - general txt
- Type=CNAME
  - **name** is an alias of a "canonical" (real) name
  - **value** is canonical name
- Type=MX
  - **value** is hostname of mail server associated with **name**
- Type=SRV
  - general extension for services
- Type=PTR
  - a pointer to another name

Try DNS: Examples

- dig [@dnsserver] <name> <type>
  - try yale.edu and various types
  - dig www.yale.edu ANY
  - dig -x IP

  - try www.yale.edu

Observations

- A name/type can return multiple answers
- DNS may rotate the answered servers
- ...

SPF Exercise

- telnet to netra.cs.yale.edu smtp
- Some test cases
  - From: yry@yale.edu
  - From: yry@harvard.edu
- dig <domain> txt to retrieve spf

DKIM Exercise

- Send email from gmail and check message

DKIM Example

- DKIM:
  Msg: DKIM-Signature: v=1; a=rsa-sha256;
  c=relaxed/relaxed;
  d=accounts.google.com; s=20161025;
  h=mime-version:date:feedback-id:message-id:subject:from:to; …
  Query:
  20161025._domainkey.accounts.google.com
  txt
- DKIM introduces a session key to allow multiple public keys
  - <session>._domainkey.<domain>

Outline

- Recap
- Email security (authentication)
- DNS
  - Interface
  - Architecture design

DNS Design: Dummy Design

- DNS itself can be considered as a client-server system as well
- How about a dummy design: introducing one super Internet DNS server?

THE DNS server of the Internet

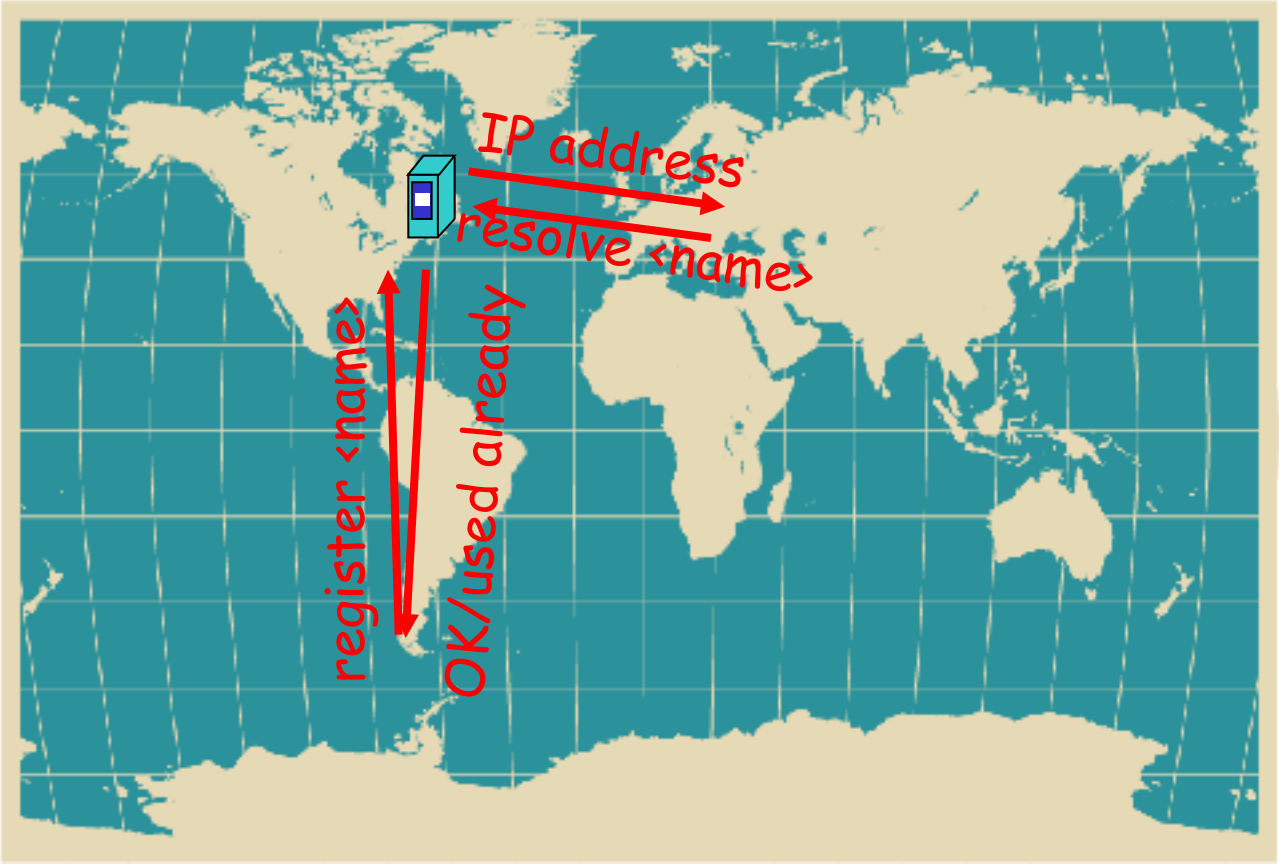

Problems of a Single DNS Server

- Scalability and robustness bottleneck
- Administrative bottleneck

DNS: Distributed Management of the Domain Name Space

- A distributed database managed by authoritative name servers
  - divided into zones, where each zone is a sub-tree of the global tree
  - each zone has its own **authoritative name servers**
  - an authoritative name server of a zone may delegate a subset (i.e. a sub-tree) of its zone to another name server

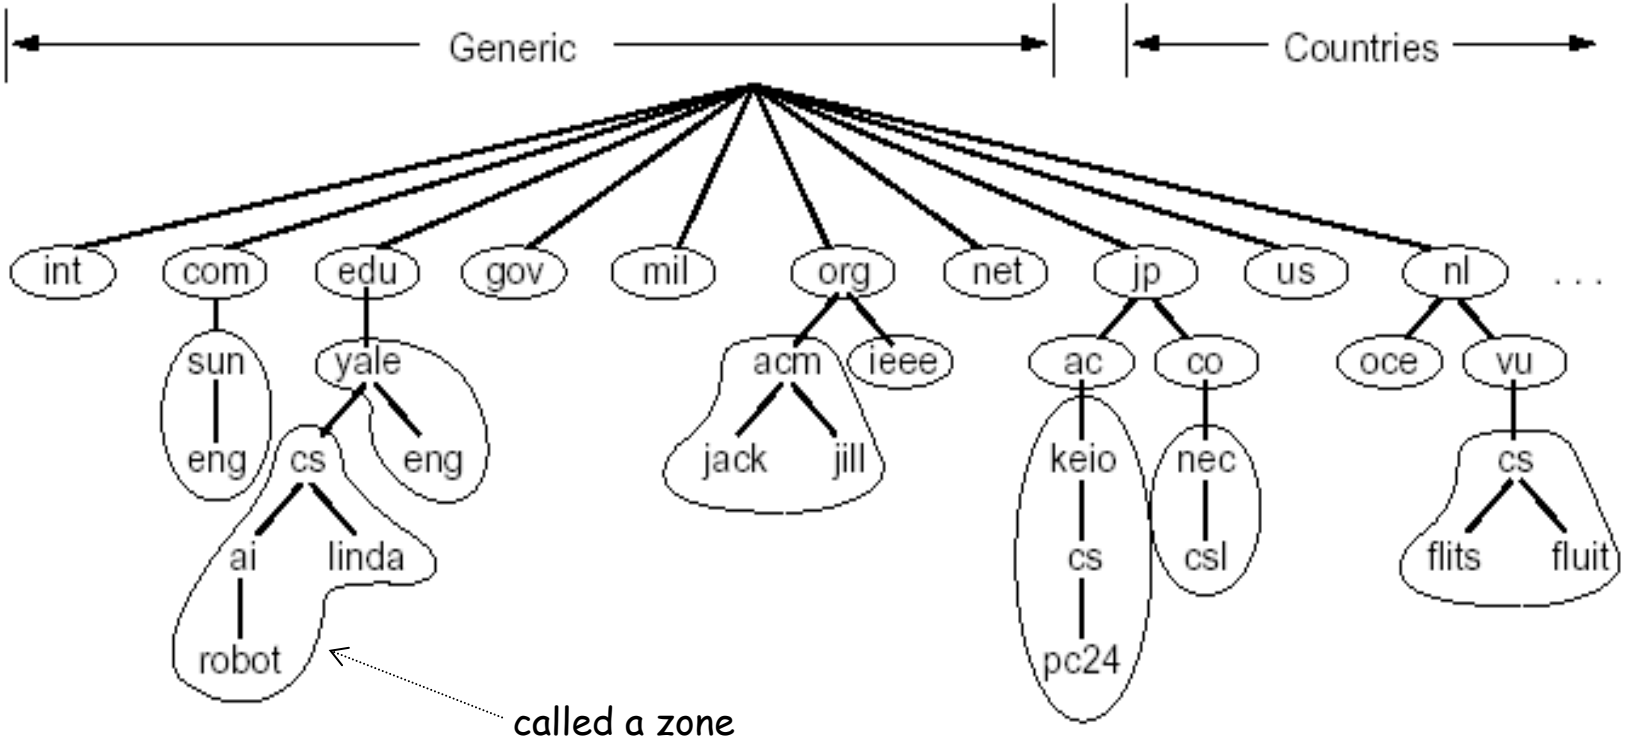

Email Architecture + DNS

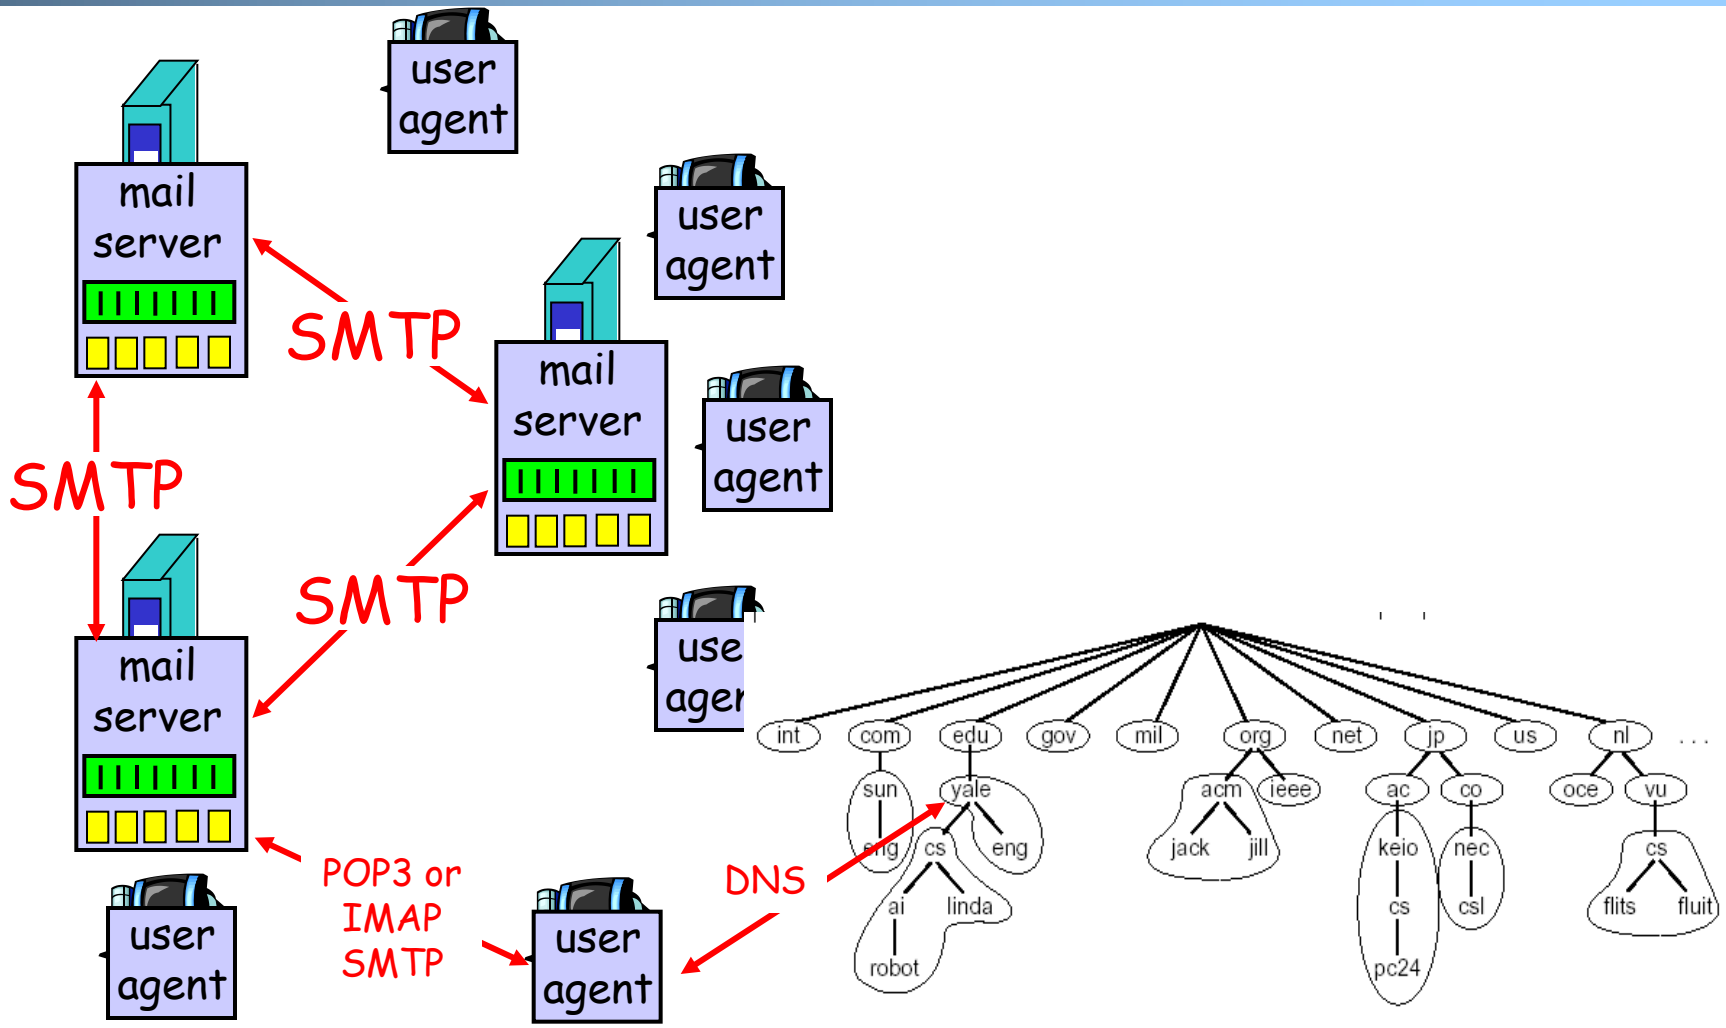

Root Zone and Root Servers

- The root zone is managed by the root name servers
  - 13 root name servers worldwide

See http://root-servers.org/ for more details

Linking the Name Servers

- Each name server knows the addresses of the root servers
- Each name server knows the addresses of its immediate children (i.e., those it delegates)

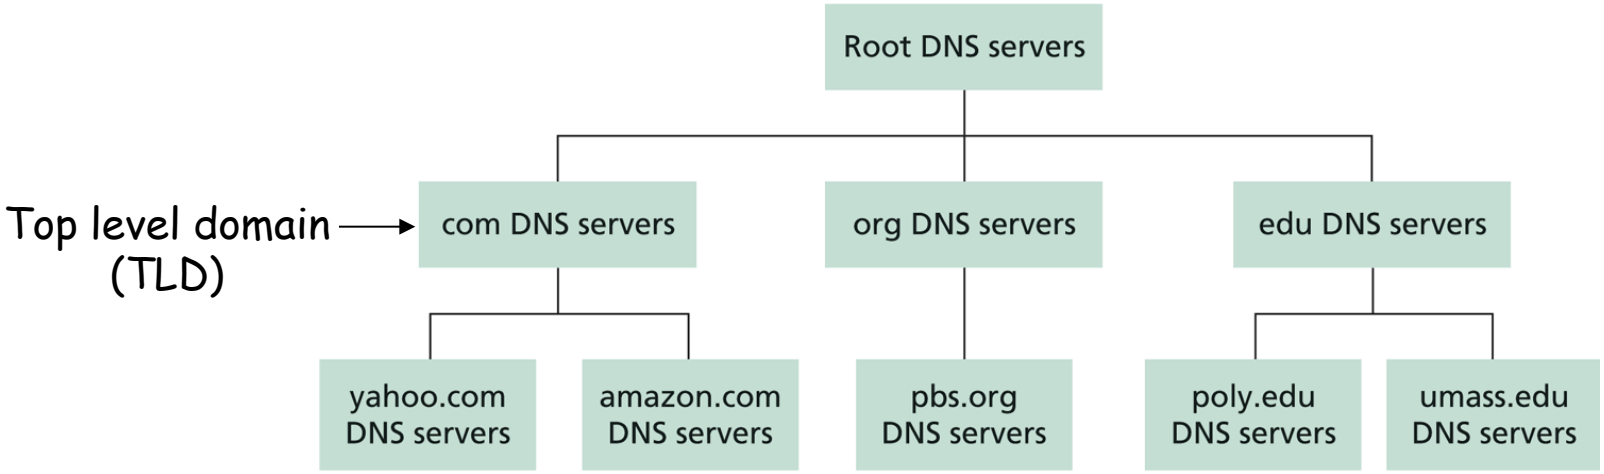

Q: how to query a hierarchy?

DNS Message Flow: Two Types of Queries

Recursive query:

- The contacted name server resolves the name completely

Iterated query:

- Contacted server replies with name of server to contact
  - "I don't know this name, but ask this server"

Two Extreme DNS Message Flows

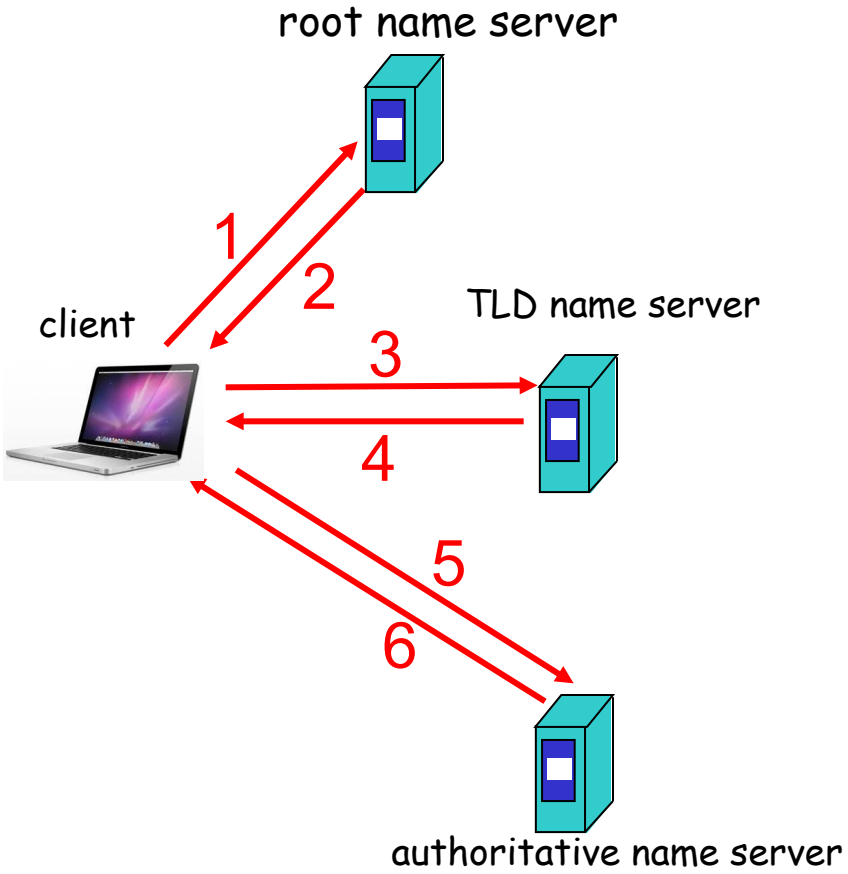

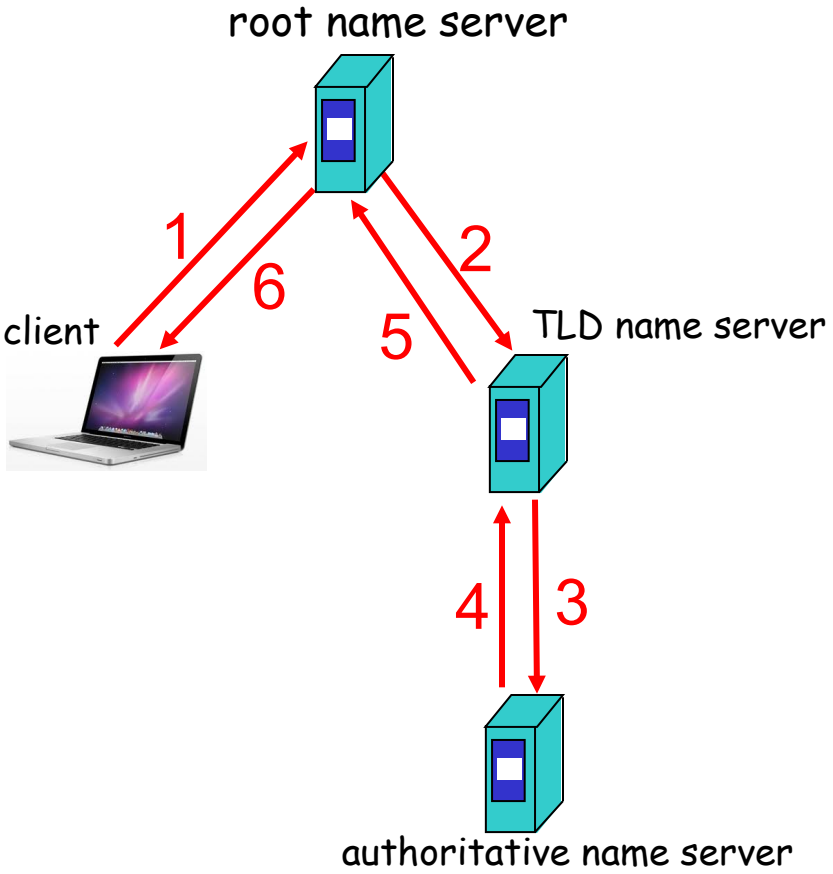

Issues of the two approaches?

cicada.cs.yale.edu

cicada.cs.yale.edu

Typical DNS Message Flow: The Hybrid Case

- Host knows only local name server
- Local name server is learned from DHCP, or configured, e.g. /etc/resolv.conf
- Local DNS server helps clients resolve DNS names

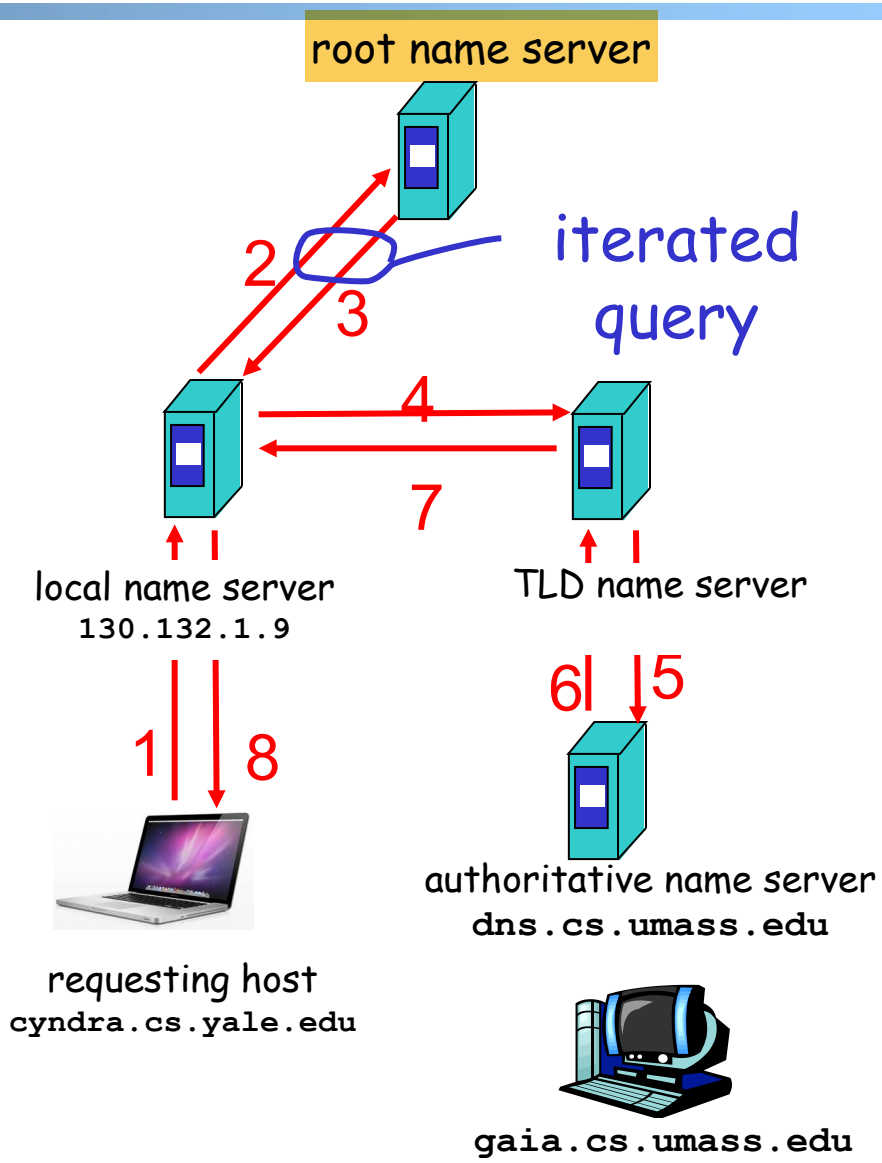

Typical DNS Message Flow: The Hybrid Case

- Host knows only local name server
- Local name server is learned from DHCP, or configured, e.g. /etc/resolv.conf
- Local DNS server helps clients resolve DNS names
- Benefits of local name servers (often called **resolvers**)
  - simplifies client
  - caches/reuses results

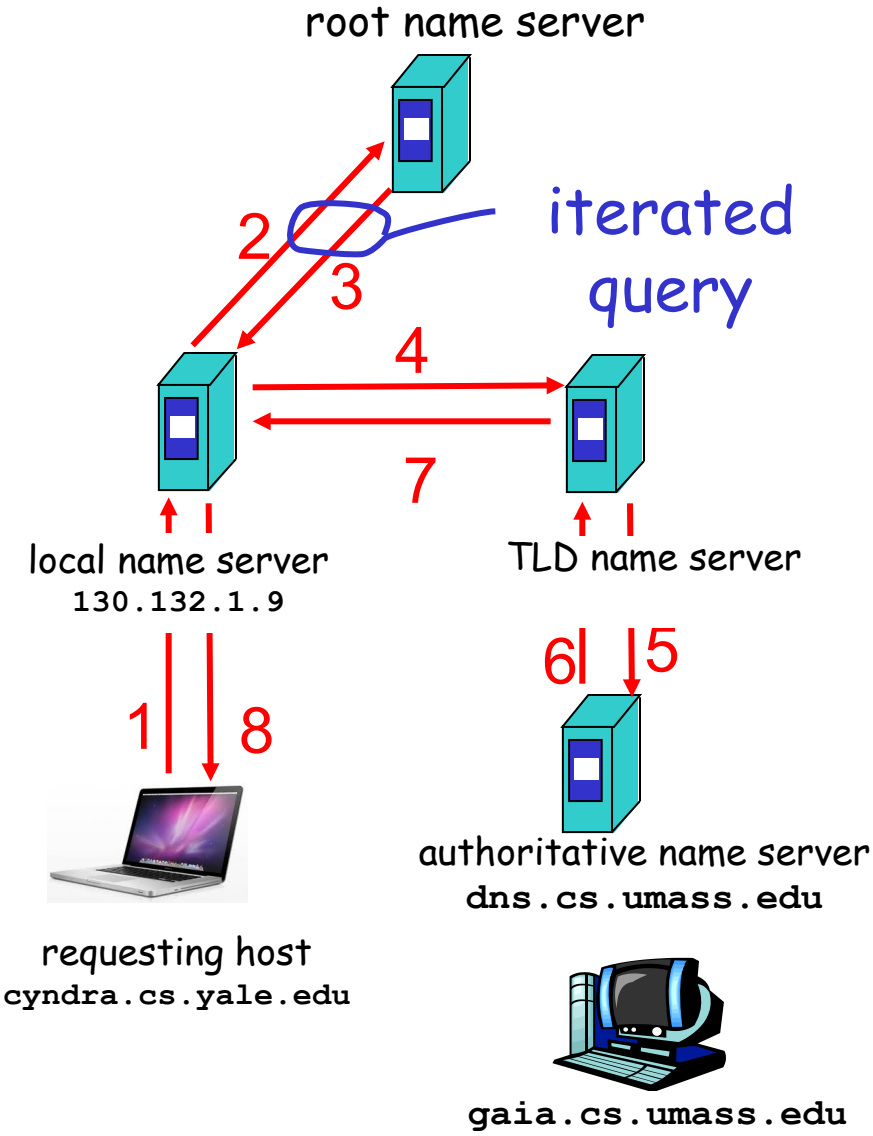

Outline

- Recap
- Email security (authentication)
- DNS
  - Interface
  - Architecture design
  - Message design

DNS Message Format?

Basic encoding decisions: UDP/TCP, how to encode domain name, how to encode answers…

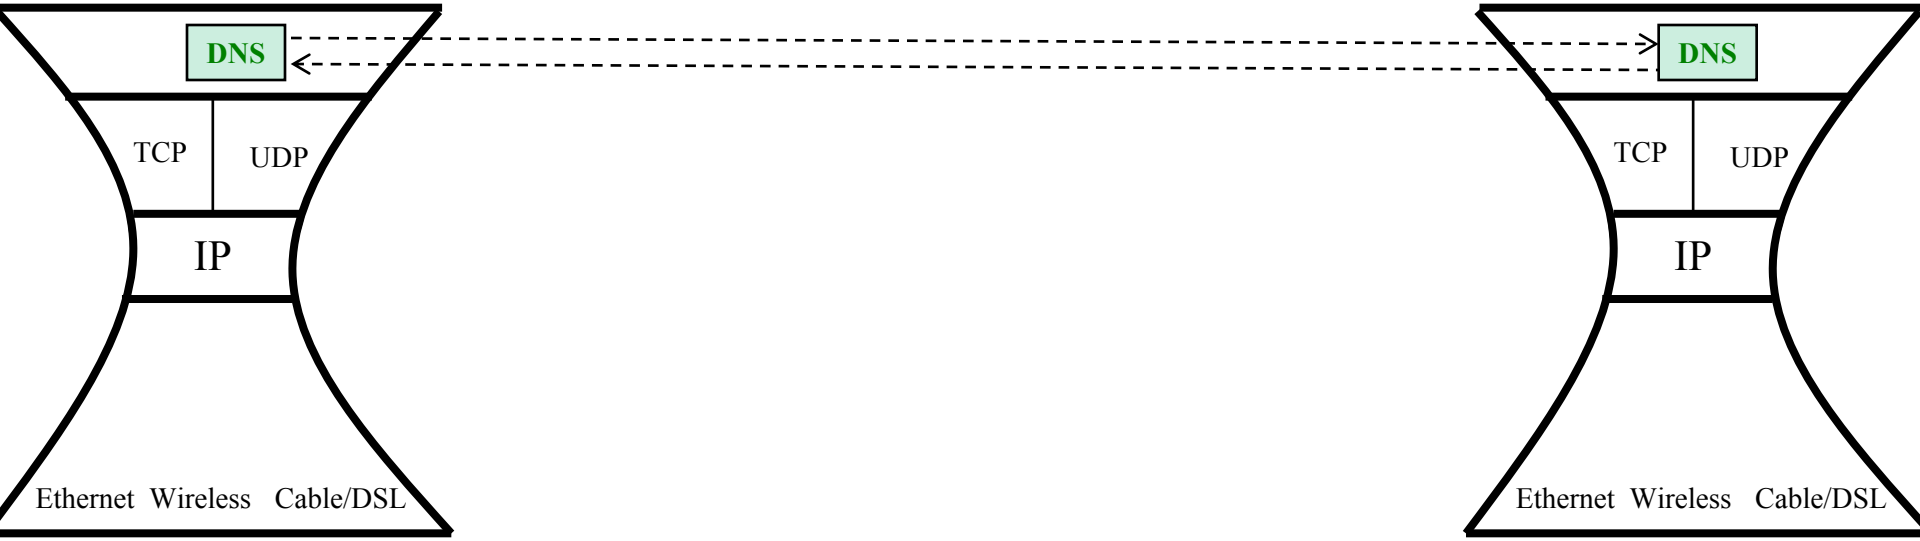

Observing DNS Messages

- Capture the messages
  - DNS server is at port 53
    - Display and clear DNS cache
      - https://support.apple.com/en-us/HT202516 (e.g., MAC sudo killall -HUP mDNSResponder)
  - visit gmail.com
  - dig +tcp to see TCP mode
  - Try to load the dns-capture file from class Schedule page, if you do not want live capture

https://www.ietf.org/rfc/rfc1035.txt

DNS Protocol, Messages

DNS protocol : typically over UDP (can use TCP):

query and *reply* messages, both with the same *message format*

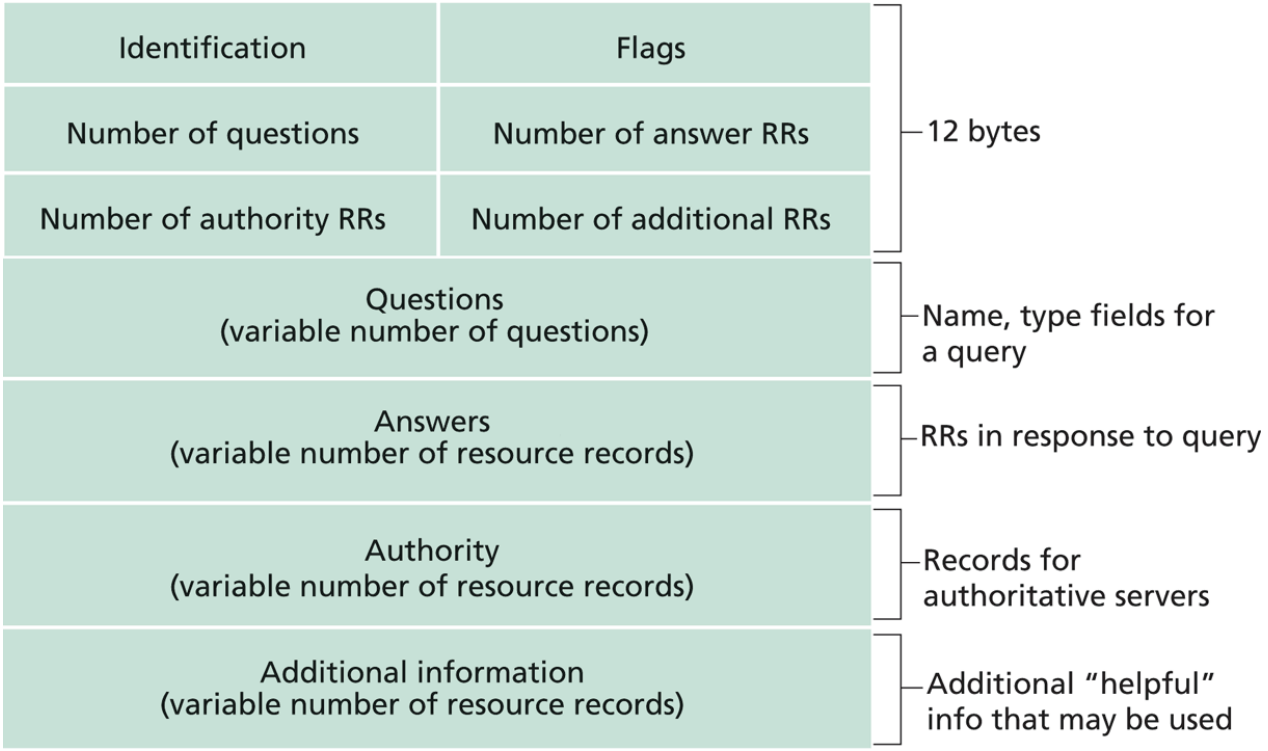

DNS Details

- Header (Sec. 4.1.1 of https://www.ietf.org/rfc/rfc1035.txt)
- Encoding of questions (Sec. 4.1.2):
  - [Label-length label-chars]
- Encoding of answers (Sec. 4.1.3)
  - Pointer format (http://www.iana.org/assignments/dns-parameters/dns-parameters.xhtml)

- See example DNS packets

Name Encoding

```
▼ Queries
  ▼ gmail.com: type A, class IN
      Name: gmail.com
      [Name Length: 9]
      [Label Count: 2]
      Type: A (Host Address) (1)
      Class: IN (0x0001)
```

```
0000  00 21 d7 75 74 00 6c 40  08 98 57 82 08 00 45 00   .!.ut.l@ ..W...E.
0010  00 37 16 b7 00 00 40 11  2e c6 ac 1b 05 91 82 84   .7....@. ........
0020  01 09 81 9b 00 35 00 23  93 65 63 32 01 00 00 01   .....5.# .ec2....
0030  00 00 00 00 00 00 05 67  6d 61 69 6c 03 63 6f 6d   .......g mail.com
0040  00 00 01 00 01                                     .....
```

Message Compression (Label Pointer)

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  1   1|                 OFFSET                |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

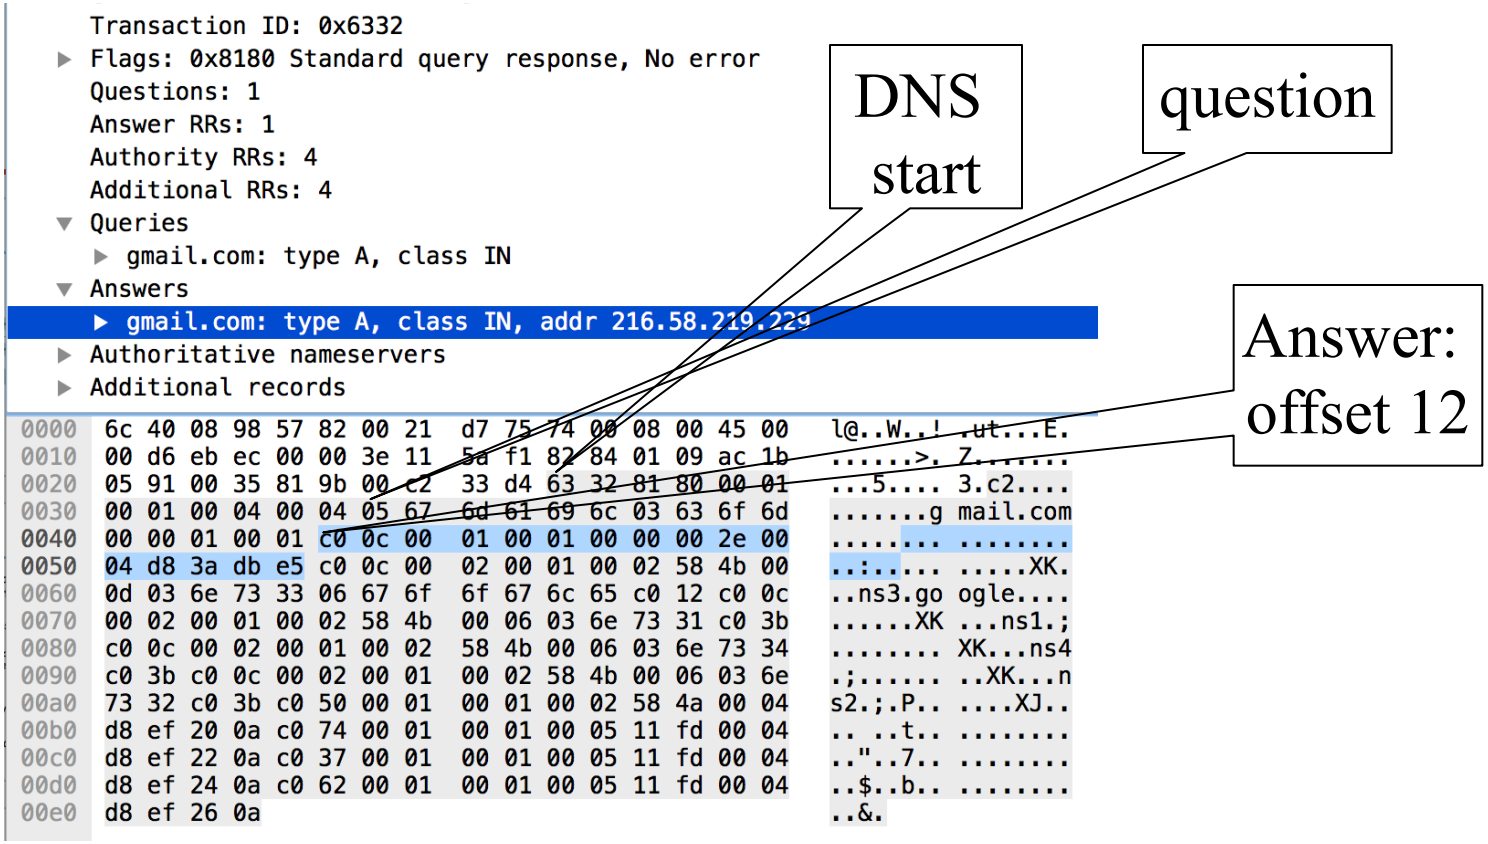

https://www.ietf.org/rfc/rfc1035.txt

Summary: DNS Protocol, Messages

Many features: typically over UDP (can use TCP); *query* and *reply* messages with the same *message format*; *length/content encoding* of *names*; simple *compression*; *additional info* as *server push*

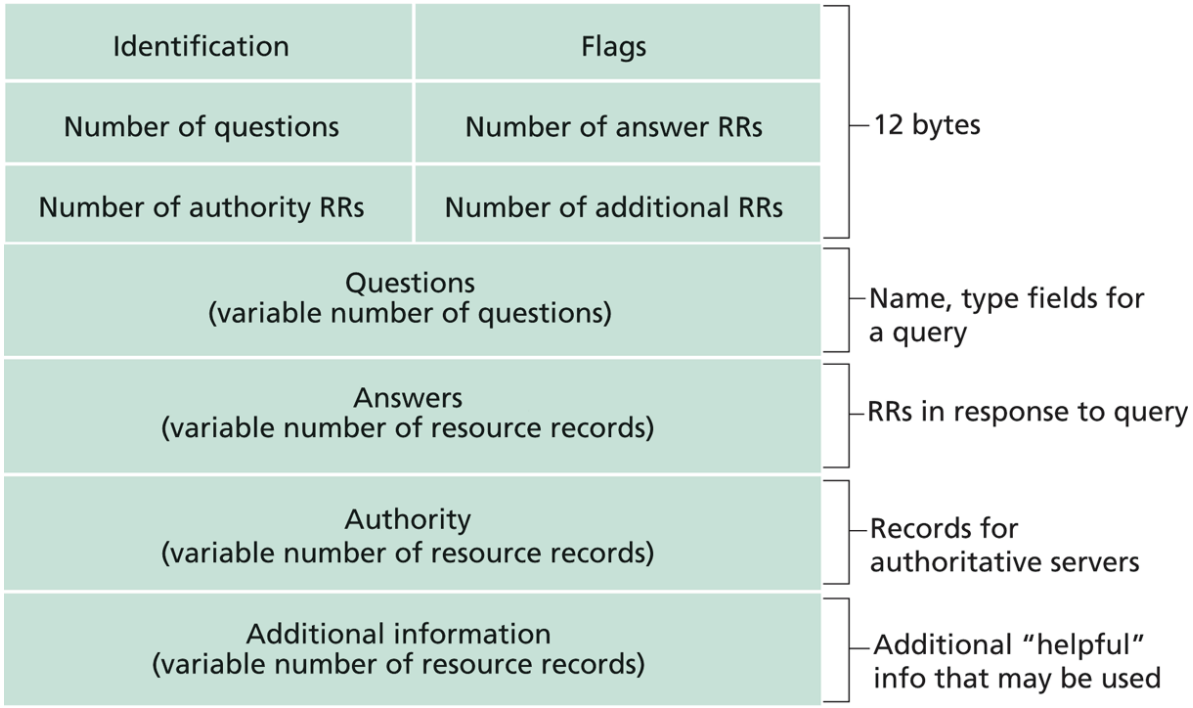

Discussion: What DNS did Right