
Network Applications: DNS; Socket Programming

Y. Richard Yang

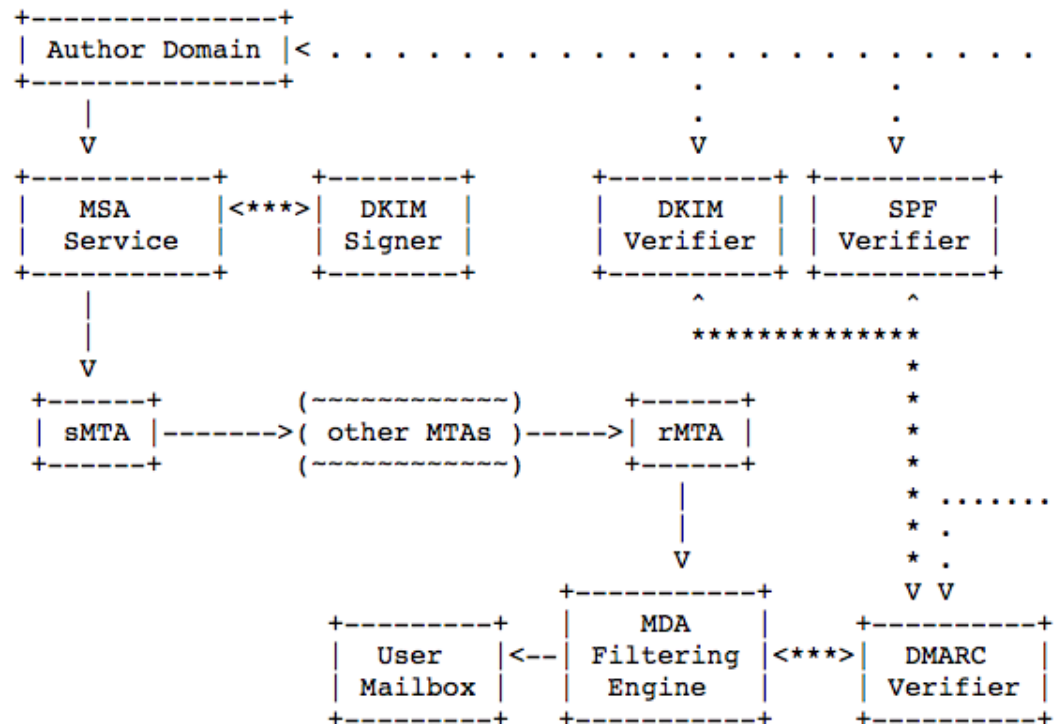
<http://zoo.cs.yale.edu/classes/cs433/>

9/20/2018

Admin

- ❑ Assignment Two to be linked on the Schedule page
- ❑ Pace slow down?

Recap: Domain-based Message Authentication, Reporting, and Conformance (DMARC) [RFC7489]




MSA = Mail Submission Agent
MDA = Mail Delivery Agent

Recap: Domain Name System (DNS)

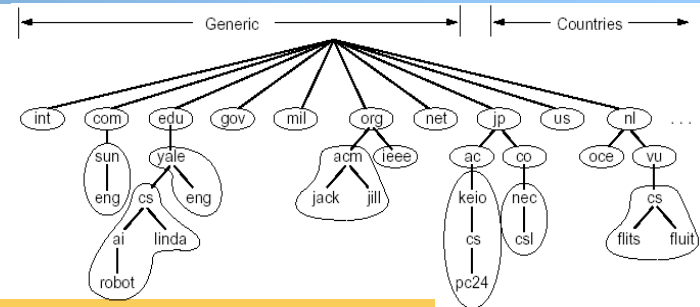
□ Function

- map between (domain name, service) to value, e.g.,
 - (www.cs.yale.edu, Addr)
→ 128.36.229.30
 - (cs.yale.edu, Email)
→ netra.cs.yale.edu



Summary: DNS Design Features

- ❑ Hierarchical name space and hierarchical delegation avoids administrative bottleneck/central control, improving manageability and scalability
- ❑ Multiple domain servers improve scalability/robustness
- ❑ Native caching (control) reduces workload and improves robustness
- ❑ Flexible recursive and iterative query allows structure such as local resolver to simplify client and enable caching
- ❑ Using UDP to reduce overhead but also support TCP using the same format
- ❑ Same query and response format can make simplify basically servers
- ❑ Domain name encoding compression reduces query/response overhead
- ❑ Proactive answers of anticipated queries (server push) reduce # queries on server and latency on client



Today: approximately 1.3 million authoritative name servers listed in the .COM, .NET and .ORG zone files.

Grown from a few thousand entries to over 100 million entries. – That's scaling!

Many Other Uses of DNS

- ❑ DNSBL (black list) or RBL (realtime)
 - Spec: <https://tools.ietf.org/html/rfc5782>
 - See changes:
<https://www.spamhaus.org/sbl/latest/>
 - Query dig <reverse>.zen.spamhaus.org
 - <https://www.spamhaus.org/zen/>

Problems/Remaining Issues of DNS

- ❑ Security of DNS itself
- ❑ Limited extensibility
 - limited query model
 - Mixed, limited query cmd and query type
 - See <https://www.iana.org/assignments/dns-parameters/dns-parameters.xhtml#dns-parameters-4>
- ❑ Largely a read data store, although theoretically you can update the values of the records, it is rarely enabled
- ❑ Each local domain needs servers, but an ad hoc domain may not have a DNS server

Outline

- ❑ Admin and recap
- DNS
 - Interface
 - Architecture design
 - Message design
 - Extensions/alternatives
 - service discovery

Context

- ❑ What do we need to extend standard DNS to support service discovery, say to implement Bonjour-type service discovery (discover local printers, local appletv, file-share...)?

DNS-Service Discovery Component: Multicast DNS [RFC6762]

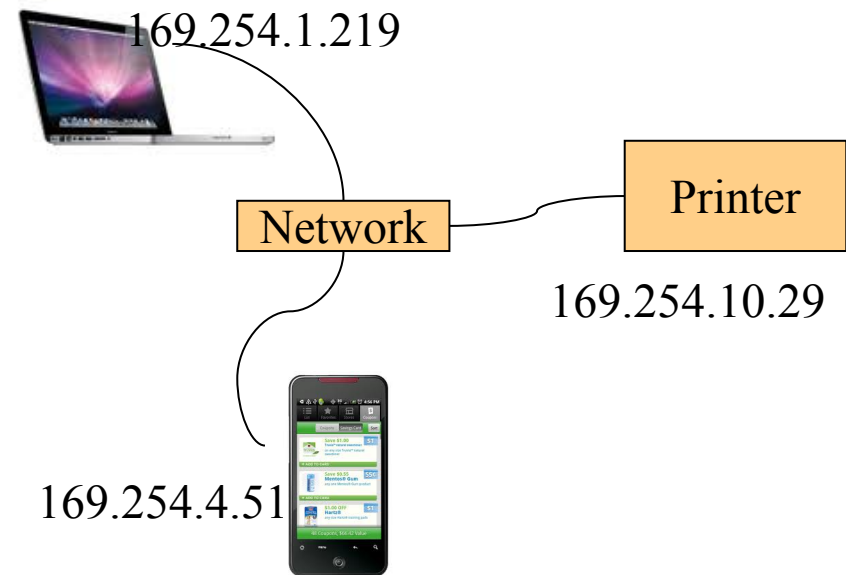
❑ Utilize IP multicast (broadcast medium)

○ link-local addressing

- send to multicast address: 224.0.0.251 (address as a group name, and any host can specify that it joins the group)

❑ Implication:

- each node (host) can become a responder
- each node (host) can use multicast to announce (write) its values



DNS-Service Discovery Component: DNS-based Service Discovery [RFC6763]

- ❑ Avoid continuous adding to DNS Resource Record Type, use
 - ptr as the only type
 - introduce an extensible service naming convention (service in name)

"My Test" _printer._tcp dns-sd.org.

Name of instance
providing service

<type_service>.
<transport>

domain (. means
default, which is local

- ❑ Example: `dig _http._tcp.dns-sd.org. ptr`

DNS-SD Local Service Discovery

- ❑ Use the dns-sd command on Mac as example
 - Advertise (register) an LPR printer on port 515

```
dns-sd -R "My Test" _printer._tcp . 515  
pdl=application/postscript
```

Name of
instance
providing
the service

<type_service>.
<transport>

domain (. means
default, which is
local

port

Txt for
additional
data

Capture packets using wireshark

24	15.749230	172.27.21.251	224.0.0.251	MDNS	232	Standard query response 0x0000 PTR, cache flush Ys-MacBook-Pro.local PTR, cache flush Ys-MacBook-Pro.local NSEC, ca...
25	15.750136	fe80::1c99:22de:9a...	ff02::fb	MDNS	252	Standard query response 0x0000 PTR, cache flush Ys-MacBook-Pro.local PTR, cache flush Ys-MacBook-Pro.local NSEC, ca...
26	15.946407	172.27.21.251	224.0.0.251	MDNS	172	Standard query 0x0000 ANY My Test._printer._tcp.local, "QU" question ANY Ys-MacBook-Pro.local, "QU" question SRV 0 ...
27	15.946465	fe80::1c99:22de:9a...	ff02::fb	MDNS	192	Standard query 0x0000 ANY My Test._printer._tcp.local, "QU" question ANY Ys-MacBook-Pro.local, "QU" question SRV 0 ...
28	16.197838	172.27.21.251	224.0.0.251	MDNS	172	Standard query 0x0000 ANY My Test._printer._tcp.local, "QM" question ANY Ys-MacBook-Pro.local, "QM" question SRV 0 ...
29	16.197896	fe80::1c99:22de:9a...	ff02::fb	MDNS	192	Standard query 0x0000 ANY My Test._printer._tcp.local, "QM" question ANY Ys-MacBook-Pro.local, "QM" question SRV 0 ...
30	16.450462	172.27.21.251	224.0.0.251	MDNS	172	Standard query 0x0000 ANY My Test._printer._tcp.local, "QM" question ANY Ys-MacBook-Pro.local, "QM" question SRV 0 ...
31	16.450508	fe80::1c99:22de:9a...	ff02::fb	MDNS	192	Standard query 0x0000 ANY My Test._printer._tcp.local, "QM" question ANY Ys-MacBook-Pro.local, "QM" question SRV 0 ...
32	16.700950	172.27.21.251	224.0.0.251	MDNS	291	Standard query response 0x0000 TXT, cache flush PTR _printer._tcp.local PTR My Test._printer._tcp.local SRV, cache ...
33	16.700998	fe80::1c99:22de:9a...	ff02::fb	MDNS	311	Standard query response 0x0000 TXT, cache flush PTR _printer._tcp.local PTR My Test._printer._tcp.local SRV, cache ...
34	16.805250	172.27.21.251	224.0.0.251	MDNS	232	Standard query response 0x0000 PTR, cache flush Ys-MacBook-Pro.local PTR, cache flush Ys-MacBook-Pro.local NSEC, ca...
35	16.805318	fe80::1c99:22de:9a...	ff02::fb	MDNS	252	Standard query response 0x0000 PTR, cache flush Ys-MacBook-Pro.local PTR, cache flush Ys-MacBook-Pro.local NSEC, ca...
36	17.703216	172.27.21.251	224.0.0.251	MDNS	291	Standard query response 0x0000 TXT, cache flush PTR _printer._tcp.local PTR My Test._printer._tcp.local SRV, cache ...
37	17.704185	fe80::1c99:22de:9a...	ff02::fb	MDNS	311	Standard query response 0x0000 TXT, cache flush PTR _printer._tcp.local PTR My Test._printer._tcp.local SRV, cache ...
38	18.808877	fe80::1c99:22de:9a...	ff02::fb	MDNS	469	Standard query response 0x0000 PTR, cache flush Ys-MacBook-Pro.local PTR, cache flush Ys-MacBook-Pro.local TXT, cac...
39	18.809057	172.27.21.251	224.0.0.251	MDNS	449	Standard query response 0x0000 PTR, cache flush Ys-MacBook-Pro.local PTR, cache flush Ys-MacBook-Pro.local TXT, cac...

Exercise

- ❑ Use the `dns-sd` command on Mac as example
 - ❑ Browse web pages on local machines

```
dns-sd -B _http._tcp
```

- Advertise (register) a web page on local machine

```
dns-sd -R "My Test" _http._tcp . 80  
path=/path-to-page.html
```

- Kill the command

Network Service Discovery in Android

- ❑ Based on DNS-SD/mDNS
- ❑ Foundation for peer-to-peer/Wi-Fi Direct in Android
- ❑ See <https://developer.android.com/training/connect-devices-wirelessly/nsd.html> for programming using nsd

General Service/Naming

Discovery Paradigm: Linda

- ❑ “Distributed workspace” by David Gelernter in the 80's at Yale
- ❑ Very influential in naming and resource discovery
- ❑ Key issues
 - How to name services/resources
 - How to write/update into name space
 - How to resolve names

The Linda Paradigm

- ❑ Naming scheme:
 - arbitrary tuples (heterogeneous-type vectors)
- ❑ Name resolution:
 - Nodes write into shared memory
 - Nodes read matching tuples from shared memory
 - exact matching is required for extraction

Linda: Core API

- ❑ `out()`: writes tuples to shared space
 - example: `out("abc", 1.5, 12)`.
 - result: insert ("abc", 1.5, 12) into space
- ❑ `read()`: retrieves tuple copy matching arg list (blocking)
 - example: `read("abc", ? A, ? B)`
 - result: finds ("abc", 1.5, 12) and sets local variables `A = 1.5, B = 12`. Tuple ("abc", 1.5, 12) is still resident in space.
- ❑ `in()`: retrieves and deletes matching tuple from space (blocking)
 - example: same as above except ("abc", 1.5, 12) is deleted
- ❑ `eval(expression)`: similar to `out` except that the tuple argument to `eval` is evaluated
 - example: `eval("ab",-6,abs(-6))` creates tuple ("ab", -6, 6)

Linda Extension: JavaSpaces

- ❑ Industry took Linda principles and made modifications
 - add transactions, leases, events
 - store Java objects instead of tuples
 - a very comprehensive service discovery system

- ❑ Definitive book, “JavaSpaces Principles, Patterns, and Practice”
 - 2 of 3 authors got Ph.D.’s from Yale

Additional Pointers

- ❑ Grapevine: Xerox PARC early 1980's Birrell, Levin, Needham, Schroeder CACM 25(1)
- ❑ The MAIN name system, an exercise in centralized computing, Deegan, Crowcroft and Warfield, ACM SIGCOMM 35(5), Oct 2005

Outline

- ❑ Admin and recap
- ❑ DNS
- Implementation/programming: UDP programming

Socket Programming

Socket API

- ❑ introduced in BSD4.1 UNIX, 1981
- ❑ Two types of sockets
 - connectionless (UDP)
 - connection-oriented (TCP)


socket

an interface (a “door”) into which one application process can **both send and receive** messages to/from another (remote or local) application process

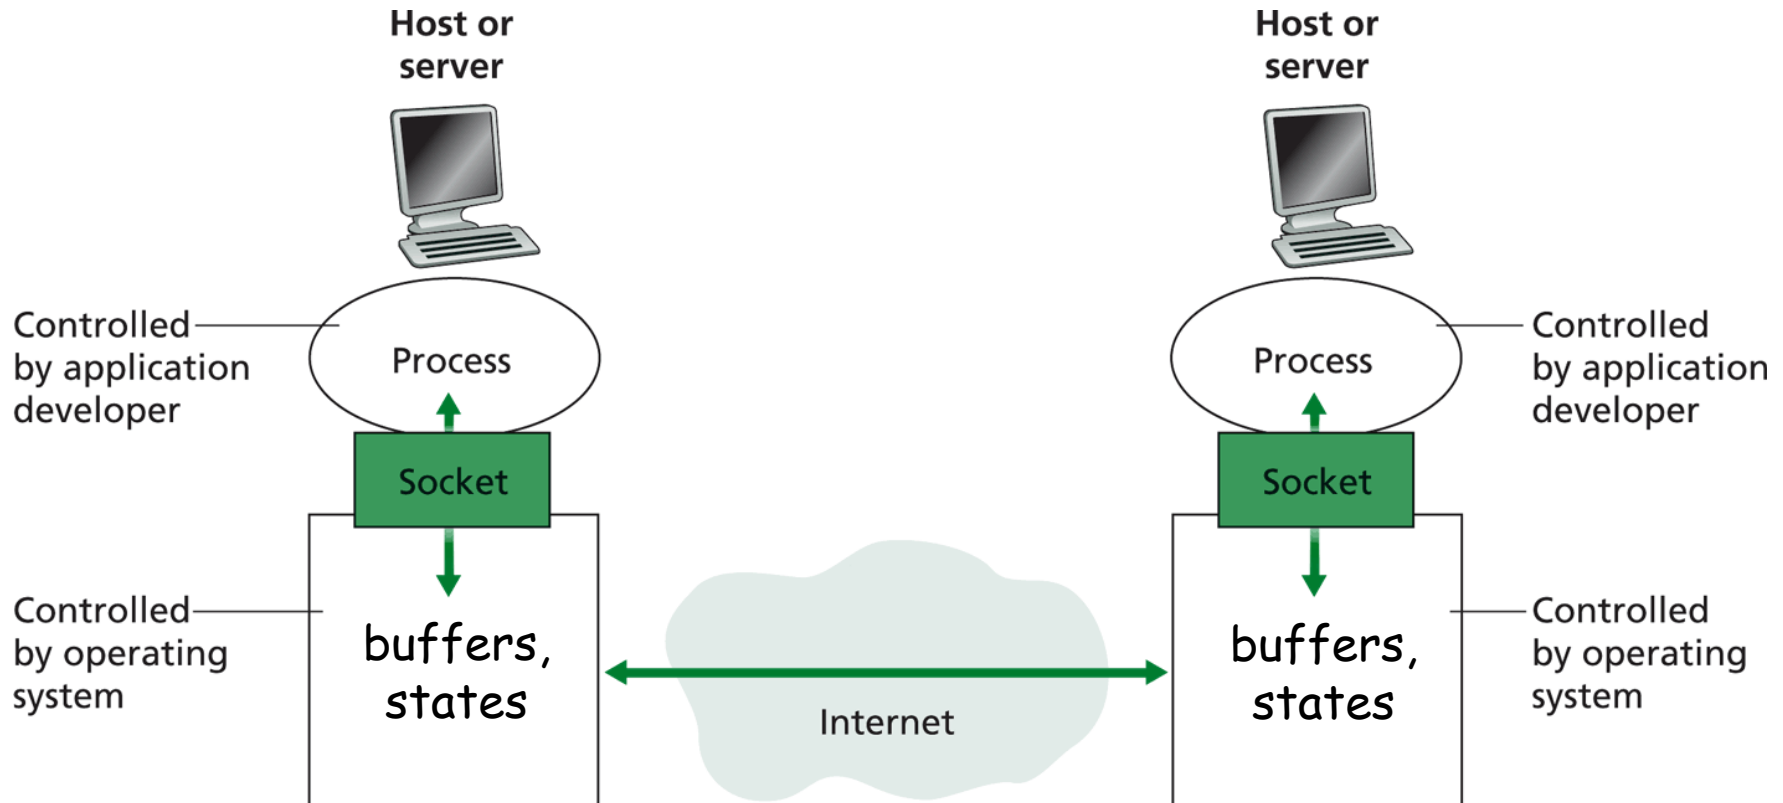
Services Provided by Transport

- User data protocol (UDP)
 - multiplexing/demultiplexing

- Transmission control protocol (TCP)
 - multiplexing/demultiplexing
 - reliable data transfer
 - rate control: flow control and congestion control



Big Picture: Socket



Outline

- ❑ Admin and recap
- ❑ DNS
- ❑ Basic network application programming
 - Overview
 - UDP (Datagram Socket)

Discussion

- What might the UDP API look like if you were to design it?

DatagramSocket (Java) (Basic)

- ❑ **DatagramSocket()**
constructs a datagram socket and binds it to any available port on the local host
 - ❑ **DatagramSocket(int lport)**
constructs a datagram socket and binds it to the specified port on the local host machine.
// more methods on multiplexing control: bind, connect; see demos
 - ❑ **DatagramPacket(byte[] buf, int length)**
constructs a DatagramPacket for receiving packets of length length.
 - ❑ **DatagramPacket(byte[] buf, int length, InetAddress address, int port)**
constructs a datagram packet for sending packets of length length to the specified port number on the specified host.
 - ❑ **receive(DatagramPacket p)**
receives a datagram packet from this socket.
 - ❑ **send(DatagramPacket p)**
sends a datagram packet from this socket.
- // socket state control
- ❑ **close()**
closes this datagram socket.

Connectionless UDP: Big Picture (Java version)

Server (running on `serv`)

create socket,
port=`x`, for
incoming request:
`serverSocket =`
`DatagramSocket(x)`

read request from
`serverSocket`

generate reply, create
datagram using client
host address, port number

write reply to
`serverSocket`

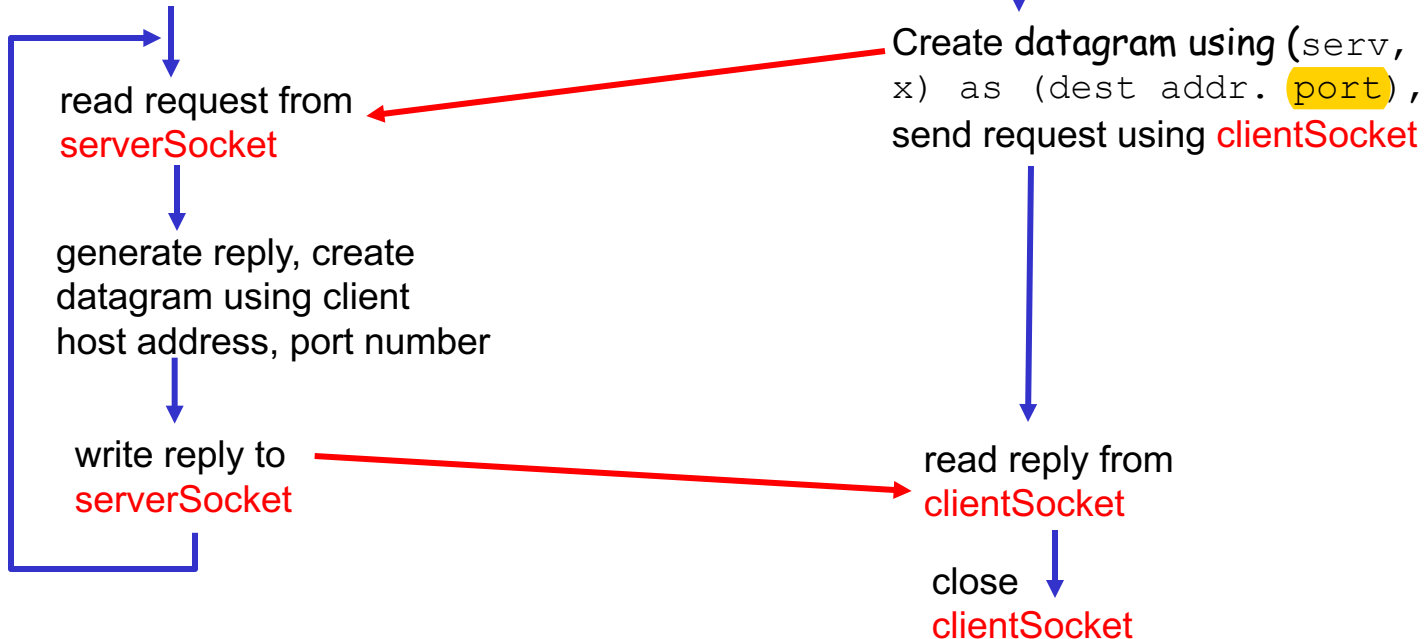
Client

create socket,
`clientSocket =`
`DatagramSocket()`

Create datagram using (`serv`,
`x`) as (dest addr. `port`),
send request using `clientSocket`

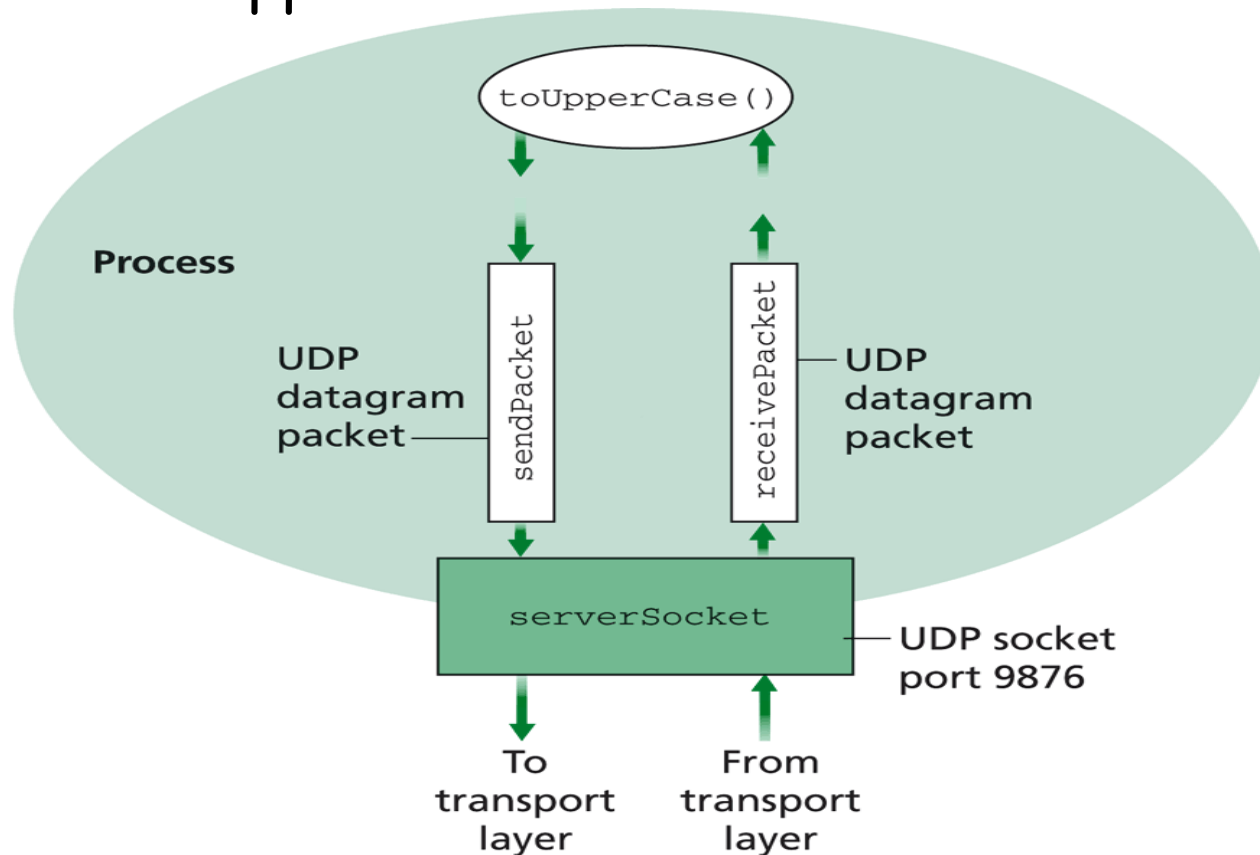
read reply from
`clientSocket`

close
`clientSocket`



Example: UDPServer.java

- A simple UDP server which changes any received sentence to upper case.




Java Server (UDP): Create Socket

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception  
    {
```

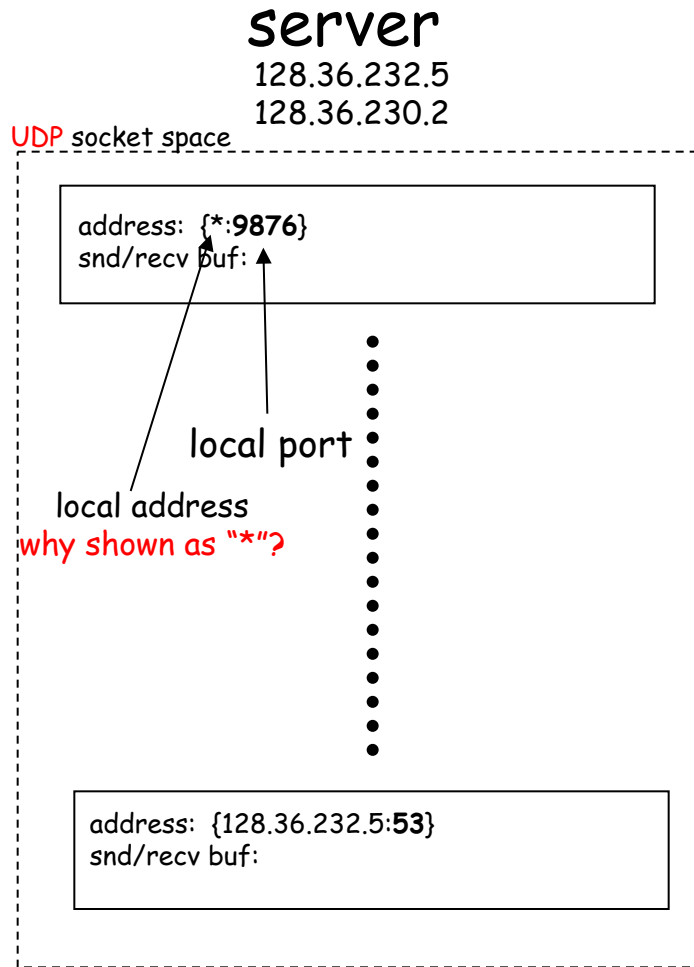
Create
datagram socket
bind at port 9876



```
        DatagramSocket serverSocket = new DatagramSocket(9876);
```

Check socket state:
%netstat -a -p udp -n

System State after the Call



"*" indicates that the socket binds to **all** IP addresses of the machine:

% `ifconfig -a`

Binding to Specific IP Addresses

server

Public address: 128.36.59.2

Local address: 127.0.0.1

UDP socket space

address: {127.0.0.1:9876}
snd/rcv buf:

address: {128.36.59.2:9876}
snd/rcv buf:

address: {*:6789}
snd/rcv buf:

•
•
•
•
•

address: {128.36.232.5:53}
snd/rcv buf:

```
InetAddress sIP1 =  
    InetAddress.getByName("localhost");  
DatagramSocket ssock1 = new  
    DatagramSocket(9876, sIP1);
```

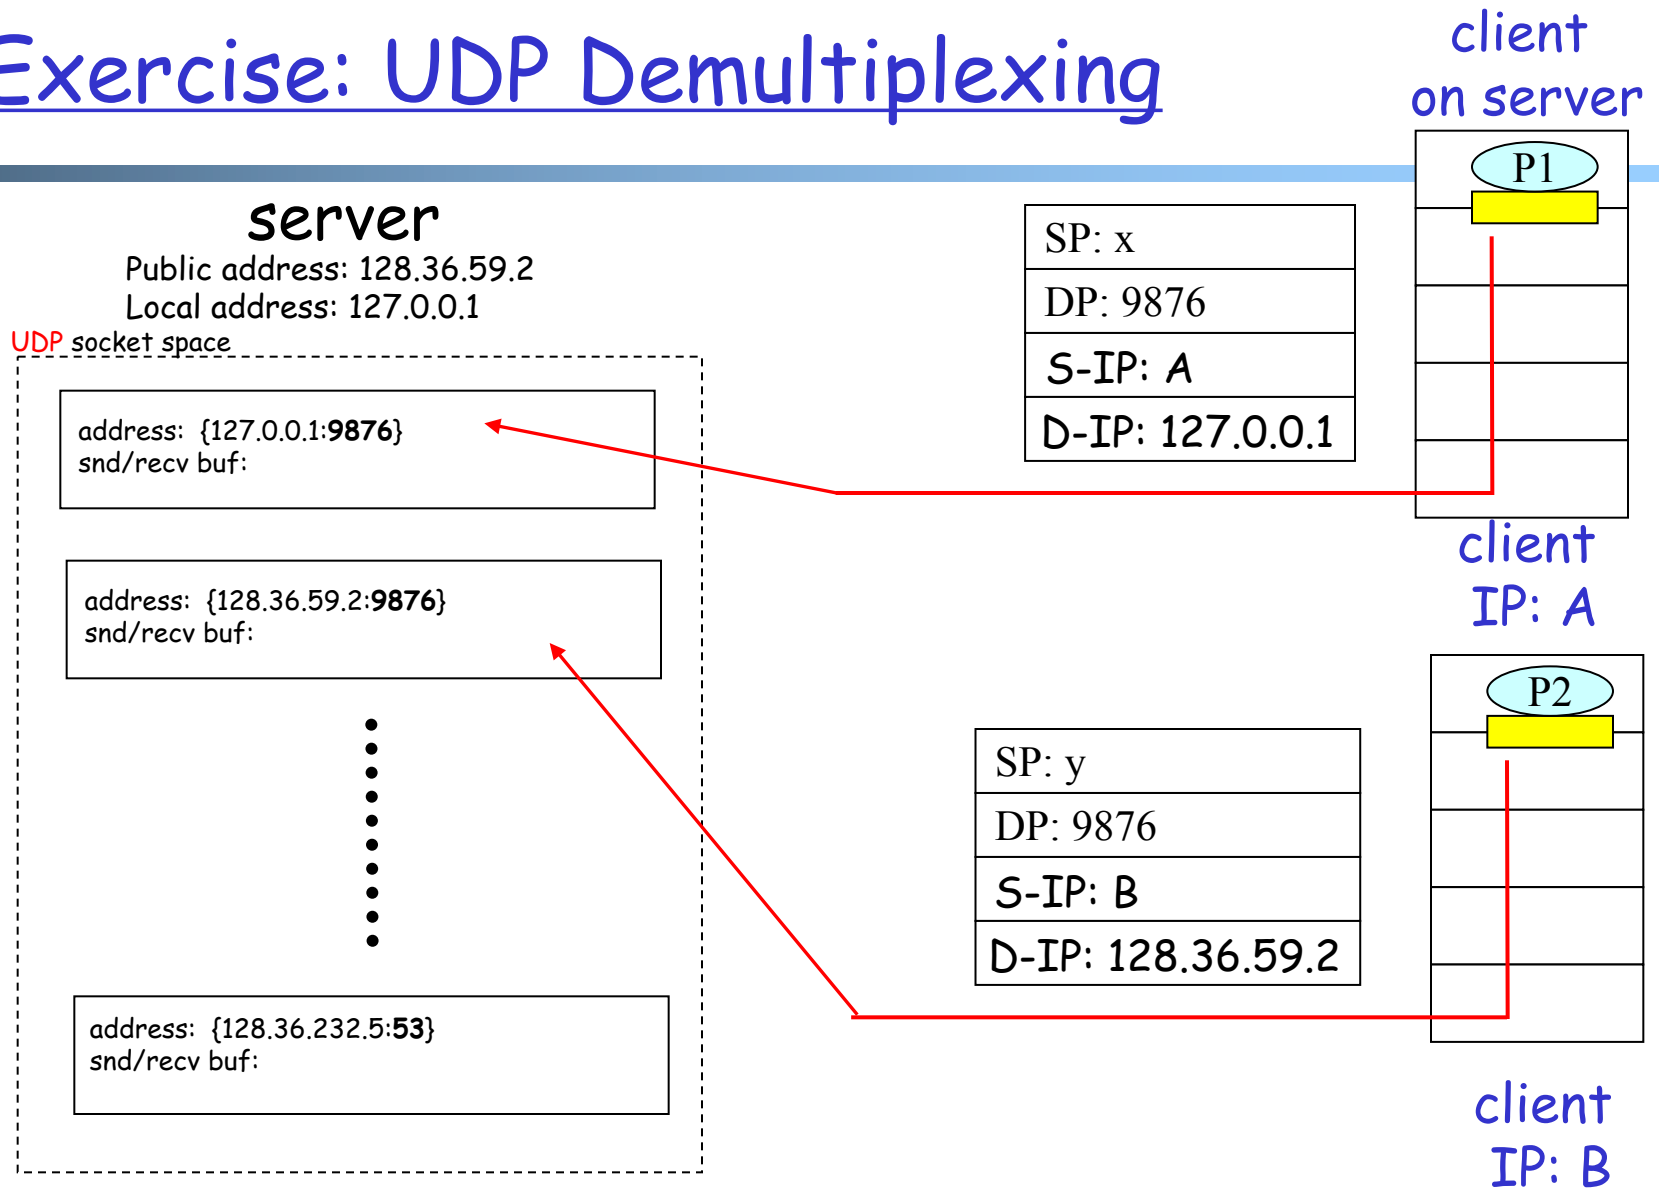
```
InetAddress sIP2 =  
    InetAddress.getByName("128.36.59.2");  
DatagramSocket ssock2 = new  
    DatagramSocket(9876, sIP2);
```

```
DatagramSocket serverSocket = new  
    DatagramSocket(6789);
```


Exercise: UDPPortScanner

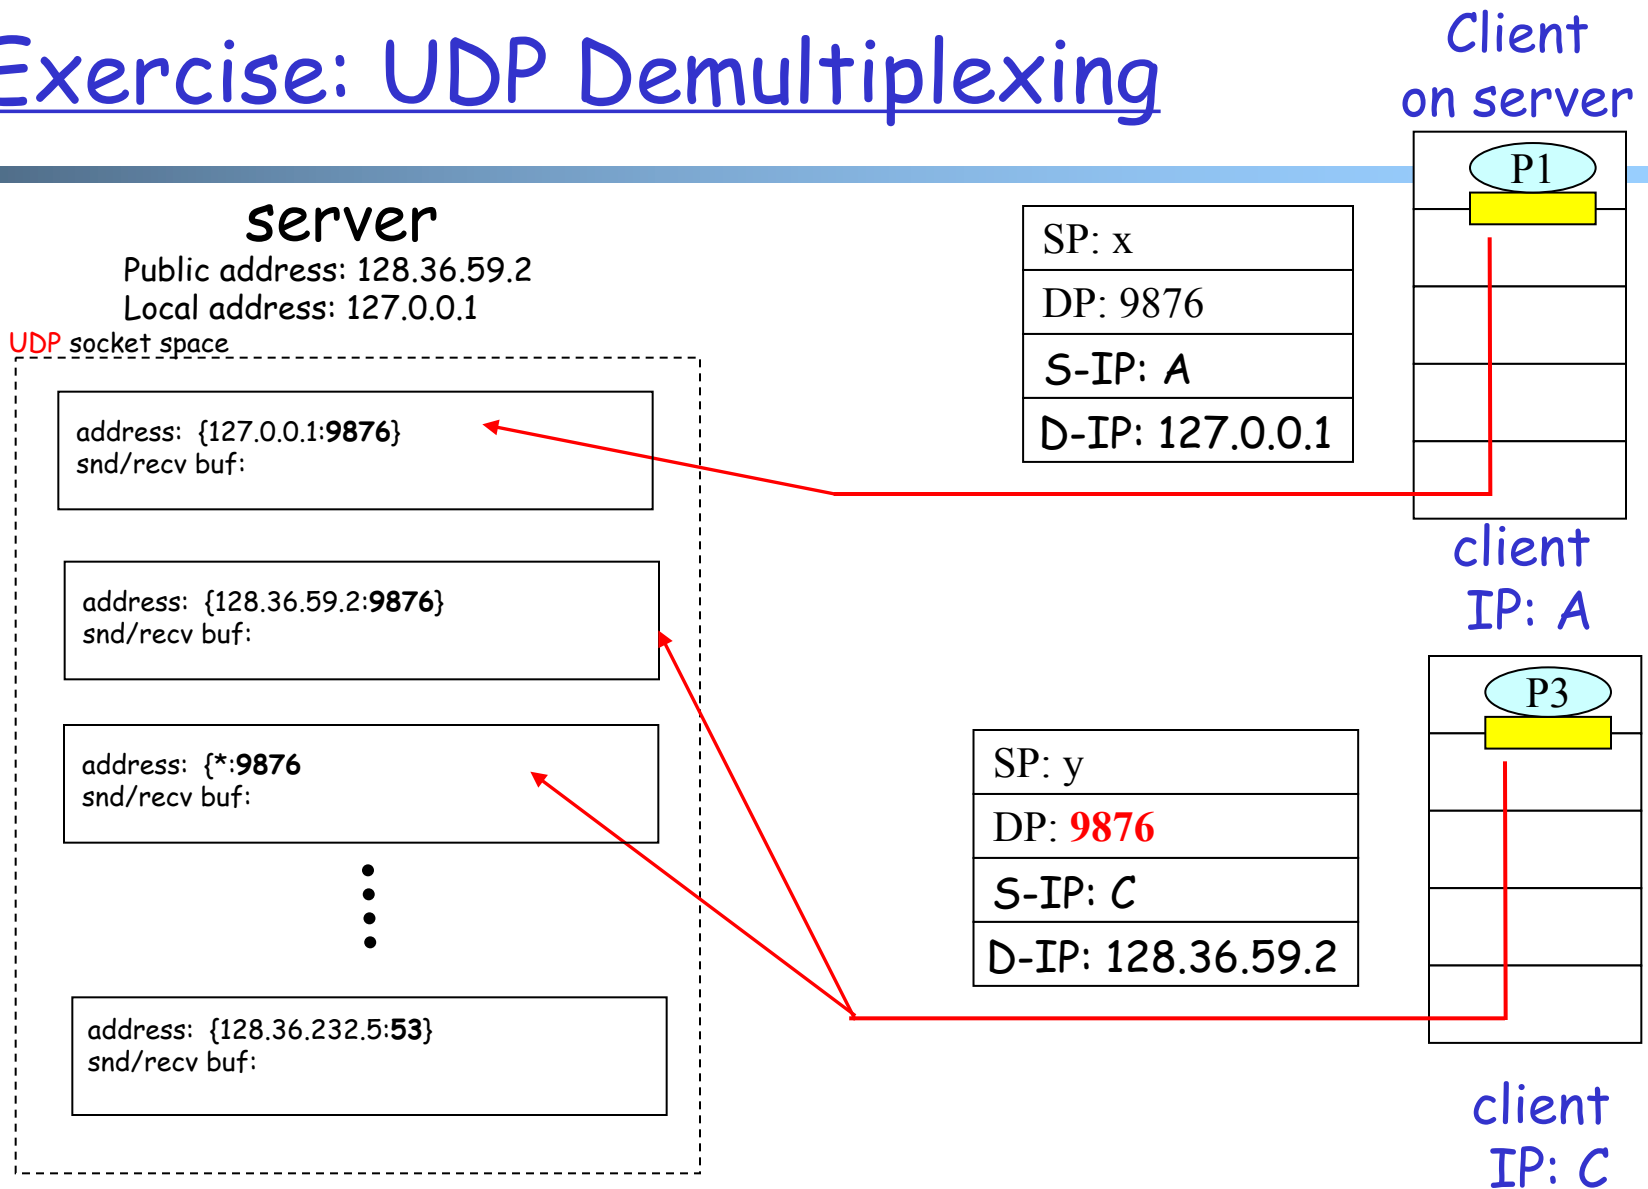
- ❑ Try to test all UDP bindings
- ❑ `[sudo] lsof -i4UDP -n -P`

Exercise: UDP Demultiplexing



UDP demultiplexing is based on matching state

Exercise: UDP Demultiplexing



UDP demultiplexing is based on matching state

Per Socket State

- ❑ Each Datagram socket has a set of states:
 - local address
 - send buffer size
 - receive buffer size
 - timeout
 - traffic class

See

<http://download.java.net/jdk7/archive/b123/docs/api/java/net/DatagramSocket.html>

Example: socket state after clients sent msgs to the server

Exercise: UDPClient

- ❑ Send messages to UDPServer from local, from a zoo machine
- ❑ Use wireshark to capture traffic

Java Server (UDP): Receiving

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception  
    {
```

```
        DatagramSocket serverSocket = new DatagramSocket(9876);
```

```
        byte[] receiveData = new byte[1024];  
        byte[] sendData = null;
```

```
        while(true)  
        {
```

Create space for
received datagram



```
            DatagramPacket receivePacket =  
                new DatagramPacket(receiveData, receiveData.length);
```

Receive
datagram



```
            serverSocket.receive(receivePacket);
```

DatagramPacket

❑ Receiving

- **DatagramPacket(byte[] buf, int length)**
constructs a DatagramPacket for receiving packets of length length.
- **DatagramPacket(byte[] buf, int offset, int length)**
constructs a DatagramPacket for receiving packets starting at offset, length length.

❑ Sending

- **DatagramPacket(byte[] buf, int length, InetAddress address, int port)**
constructs a datagram packet for sending packets of length length to the specified port number on the specified host.
- **DatagramPacket(byte[] buf, int offset, int length, InetAddress address, int port)**

Java Server (UDP): Processing

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception {
```

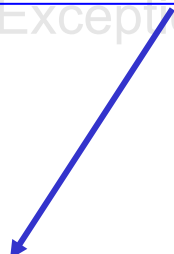
```
        ...
```

```
        // process data
```

```
        String sentence = new String(receivePacket.getData(),  
                                     0, receivePacket.getLength());
```

```
        String capitalizedSentence = sentence.toUpperCase();  
        sendData = capitalizedSentence.getBytes();
```

getData() returns a pointer to
an underlying buffer array



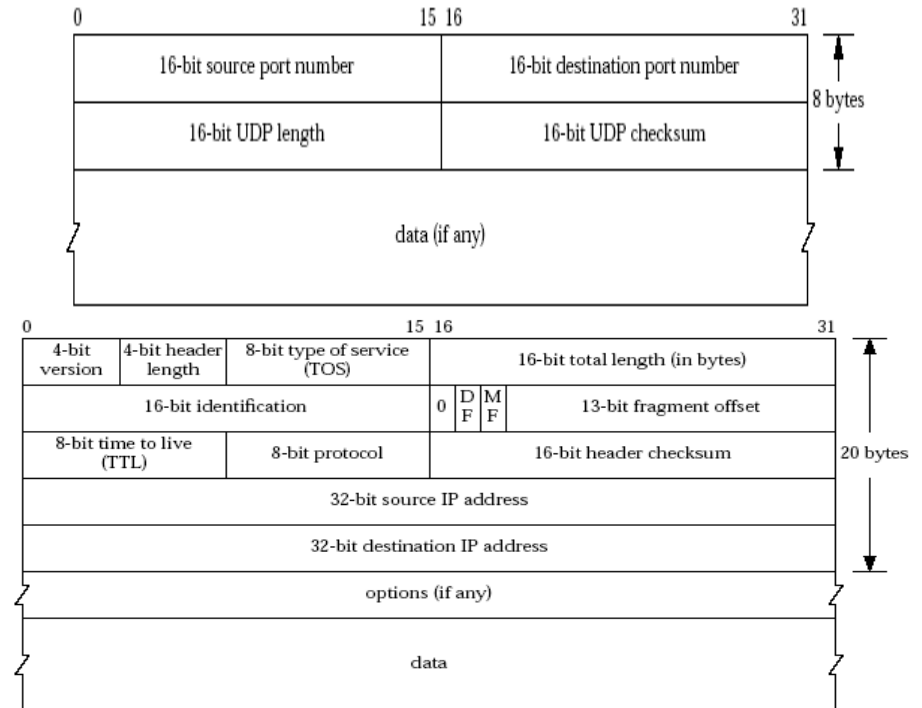
getLength() returns how much
data is valid.



Java Server (UDP): Response

❑ Java DatagramPacket:

- getAddress() / getPort()
() returns the **source** address/port



Java server (UDP): Reply

Get IP addr
port #, of
sender

```
InetAddress IPAddress = receivePacket.getAddress();  
int port = receivePacket.getPort();
```

Create datagram
to send to client

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length,  
        IPAddress, port);
```


Write out
datagram
to socket

```
serverSocket.send(sendPacket);  
}  
}
```

End of while loop,
loop back and wait for
another datagram

Example: UDPClient.java

- A simple UDP client which reads input from keyboard, sends the input to server, and reads the reply back from the server.



Example: Java client (UDP)

```
import java.io.*;  
import java.net.*;
```

```
class UDPClient {  
    public static void main(String args[]) throws Exception  
    {
```

Create
input stream

```
        BufferedReader inFromUser =  
            new BufferedReader(new InputStreamReader(System.in));  
        String sentence = inFromUser.readLine();  
        byte[] sendData = sentence.getBytes();
```

Create
client socket

```
        DatagramSocket clientSocket = new DatagramSocket();
```

Translate
hostname to IP
address using DNS

```
        InetAddress sIPAddress = InetAddress.getByName("servname");
```

Example: Java client (UDP), cont.

Create datagram
with data-to-send,
length, IP addr, port

Send datagram
to server

Read datagram
from server

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, sIPAddress, 9876);  
  
clientSocket.send(sendPacket);  
  
byte[] receiveData = new byte[1024];  
DatagramPacket receivePacket =  
    new DatagramPacket(receiveData, receiveData.length);  
  
clientSocket.receive(receivePacket);  
  
String modifiedSentence =  
    new String(receivePacket.getData());  
  
System.out.println("FROM SERVER:" + modifiedSentence);  
clientSocket.close();  
}  
}
```

Java Server (UDP): Processing

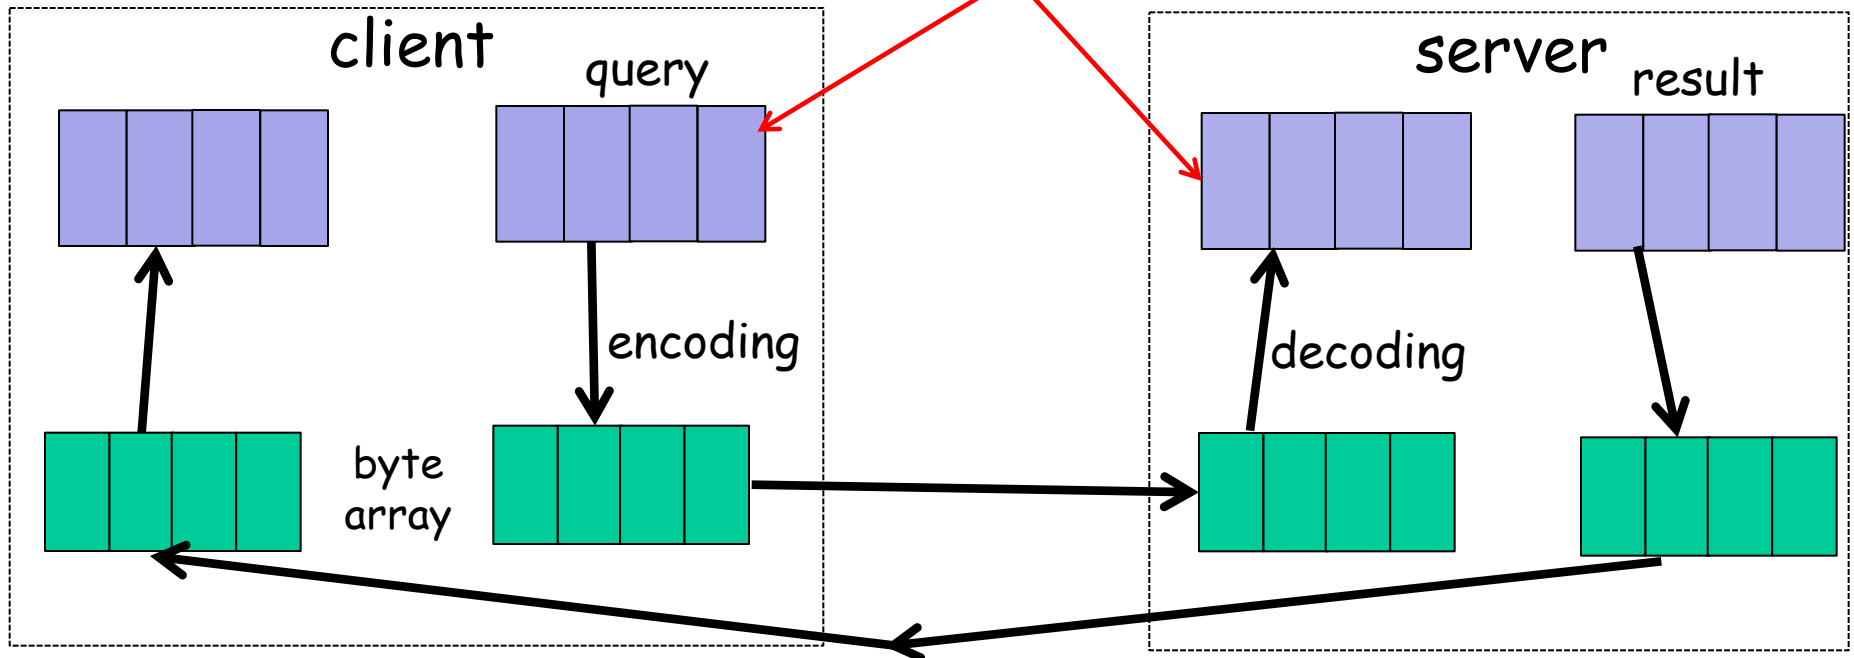
A simple upper-case UDP echo service is among the simplest network service. Are there any problems with the processing?

```
class UDPServer {  
    public static void main(String args[]) throws Exception {  
  
        ...  
        DatagramPacket receivePacket = new DatagramPacket(receiveData, receiveData.length);  
        serverSocket.receive(receivePacket);  
  
        // process  
        String sentence = new String(receivePacket.getData(),  
                                     0, receivePacket.getLength());  
        String capitalizedSentence = sentence.toUpperCase();  
        sendData = capitalizedSentence.getBytes();  
  
        // send  
        DatagramPacket sendPacket = new DatagramPacket(sendData, sendData.length,  
                                                         IPAddress, port);  
        serverSocket.send(sendPacket);  
    }  
}
```

Data Encoding/Decoding

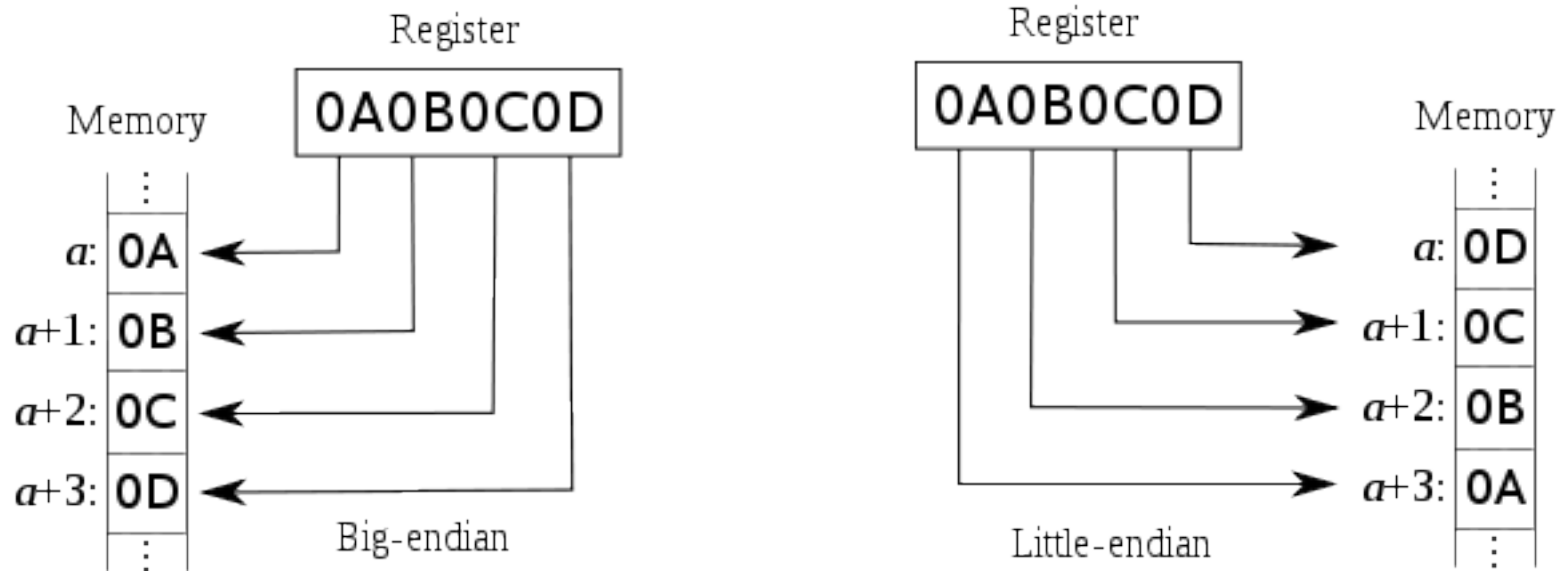
- ❑ Rule: **ALWAYS** pay attention to encoding/decoding of data

if not careful, query sent != query received (how?)



Example: Endianness of Numbers

□ `int var = 0x0A0B0C0D`



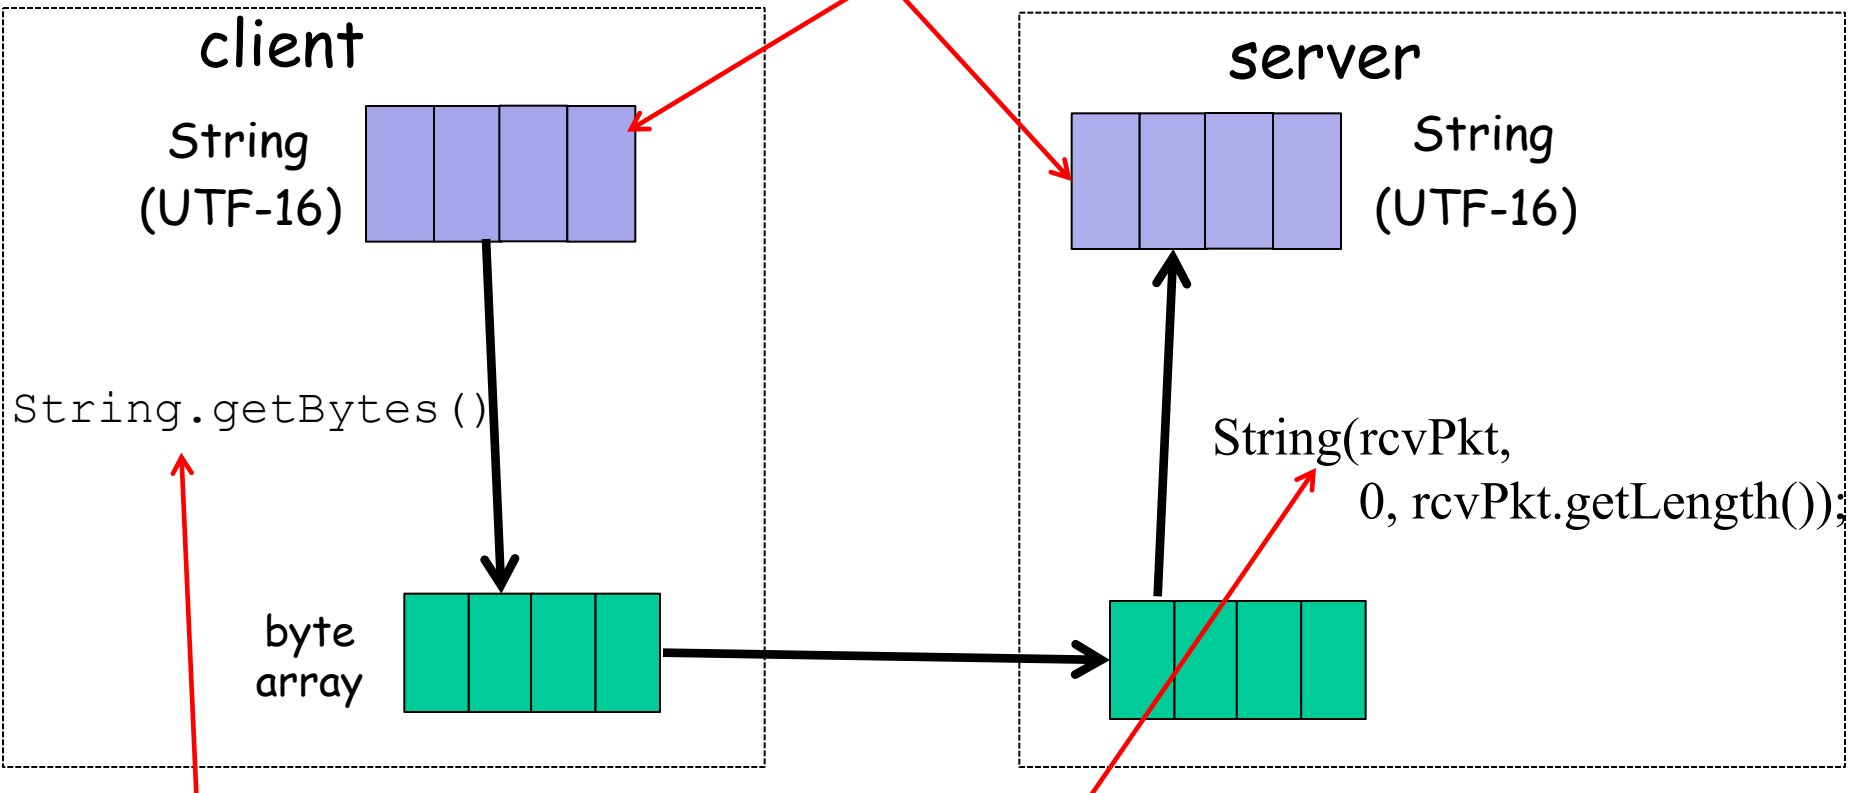
ARM, Power PC, Motorola 68k, IA-64

Intel x86

□ `sent != received`: take an int on a big-endian machine and send a little-endian machine

Example: String and Chars

Will we always get back the same string?



Depends on default local platform char set :
`java.nio.charset.Charset.defaultCharset()`

Example: Charset Troubles

- ❑ Try
 - java EncodingDecoding US-ASCII UTF-8

Encoding/Decoding as a Common Source of Errors

- ❑ Please read chapter 2 (Streams) of Java Network Programming for more details
 - Java stream, reader/writer can always be confusing, but it is good to finally understand

- ❑ Common mistake even in many (textbook) examples:
 - <http://www.java2s.com/Code/Java/Network-Protocol/UseDatagramSockettosendoutandreceiveDatagramPacket.htm>

Offline Exercise: UDP/DNS Server Pseudocode

- ❑ Modify the example UDP server code to implement a local DNS server.

Identification	Flags	
Number of questions	Number of answer RRs	12 bytes
Number of authority RRs	Number of additional RRs	
Questions (variable number of questions)		Name, type fields for a query
Answers (variable number of resource records)		RRs in response to query
Authority (variable number of resource records)		Records for authoritative servers
Additional information (variable number of resource records)		Additional "helpful" info that may be used

