
Network Applications:
High-Performance Server Design
(Thread, Async)

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

10/4/2018

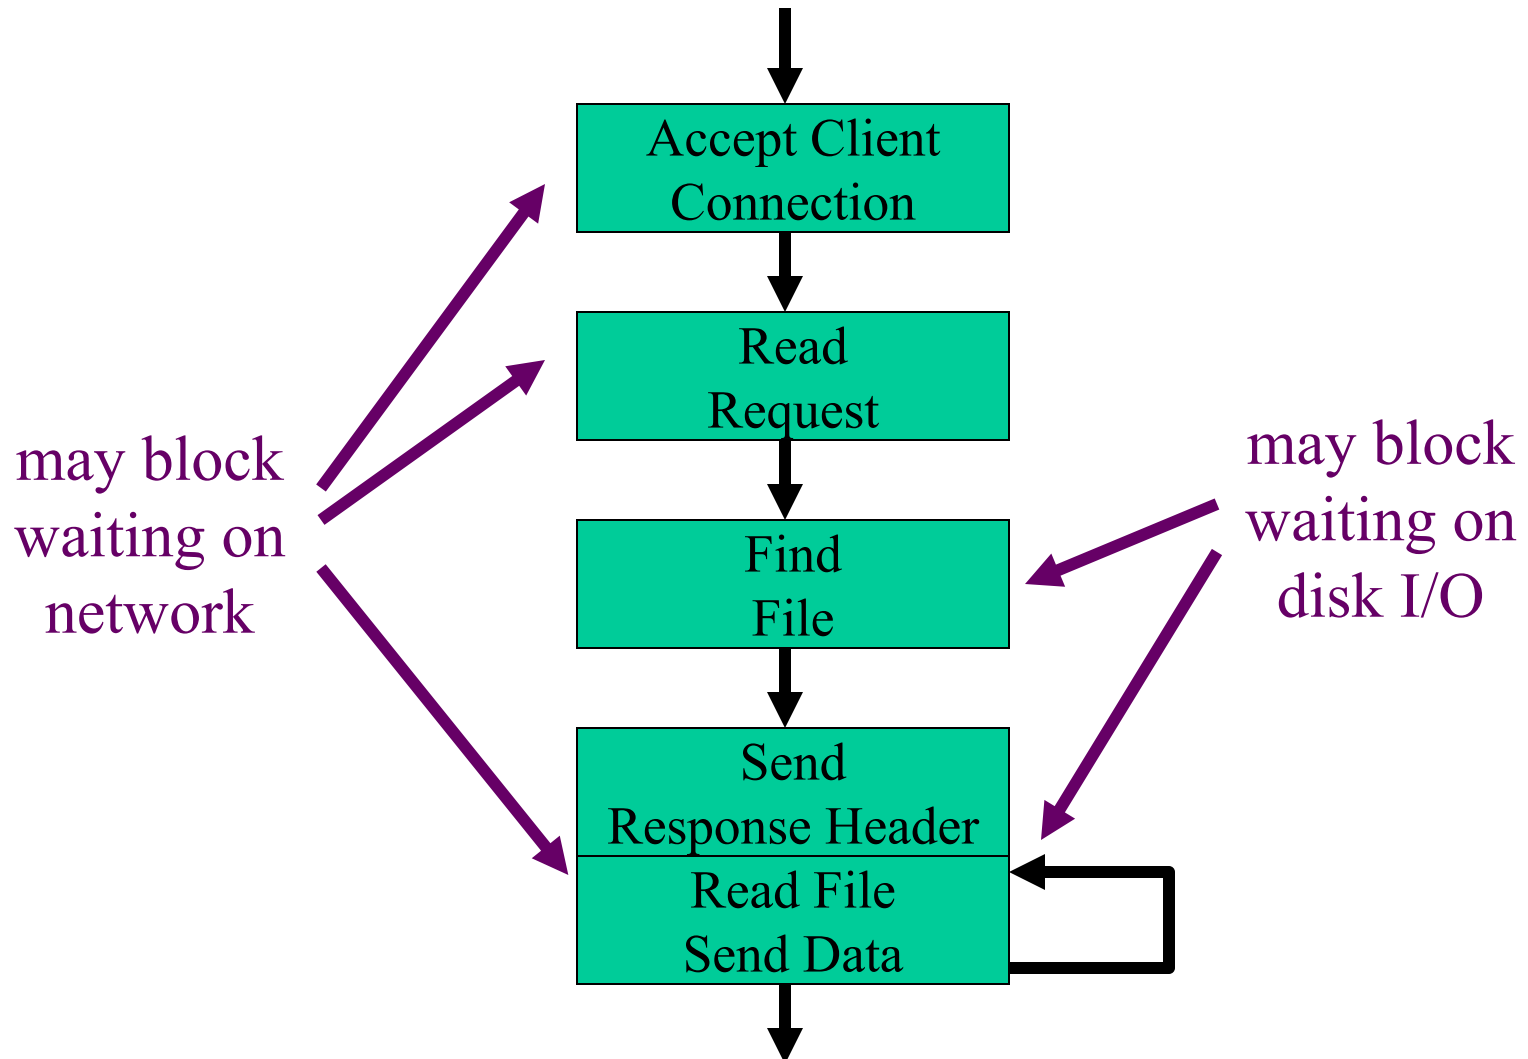
Outline

- ❑ Admin and recap
- ❑ High-performance network server design
 - Overview
 - Threaded servers



Admin

- ❑ Assignment Three (HTTP server)
Assignment Part 1 posted

Recap: Server Processing Steps




Recap: Best Server Design Limited Only by Resource Bottleneck


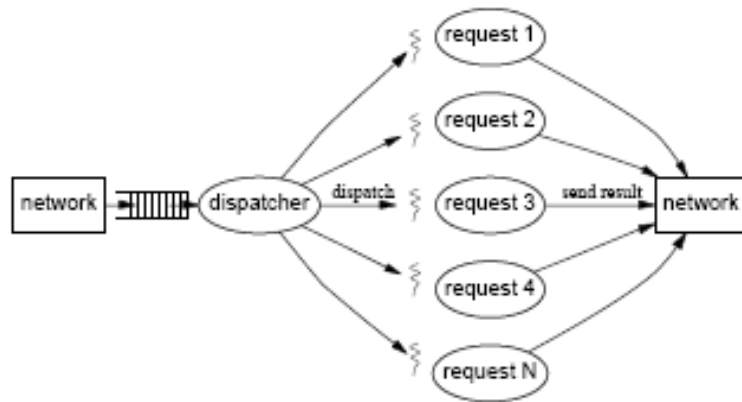


Recap: Per-Request Thread Server

```
main() {  
    ServerSocket s = new ServerSocket(port);  
    while (true) {  
        Socket conSocket = s.accept();  
        RequestHandler rh  
            = new RequestHandler(conSocket);  
        Thread t = new Thread (rh);  
        t.start();  
    }  
}
```



Recap: Problem of Per-Request Thread




(937 MHz x86, Linux 2.2.14, each thread reading 8KB file)

- ❑ High thread creation/deletion overhead
- ❑ Too many threads → resource overuse → throughput meltdown → response time explosion
- ❑ Handy tool: Little' Law

Using a Fixed Set of Threads (Thread Pool)


- ❑ Design issue: how to distribute the requests from the welcome socket to the thread workers



Design 1: Threads Share Access to the welcomeSocket

```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = welcomeSocket.accept();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```

sketch; not
working code

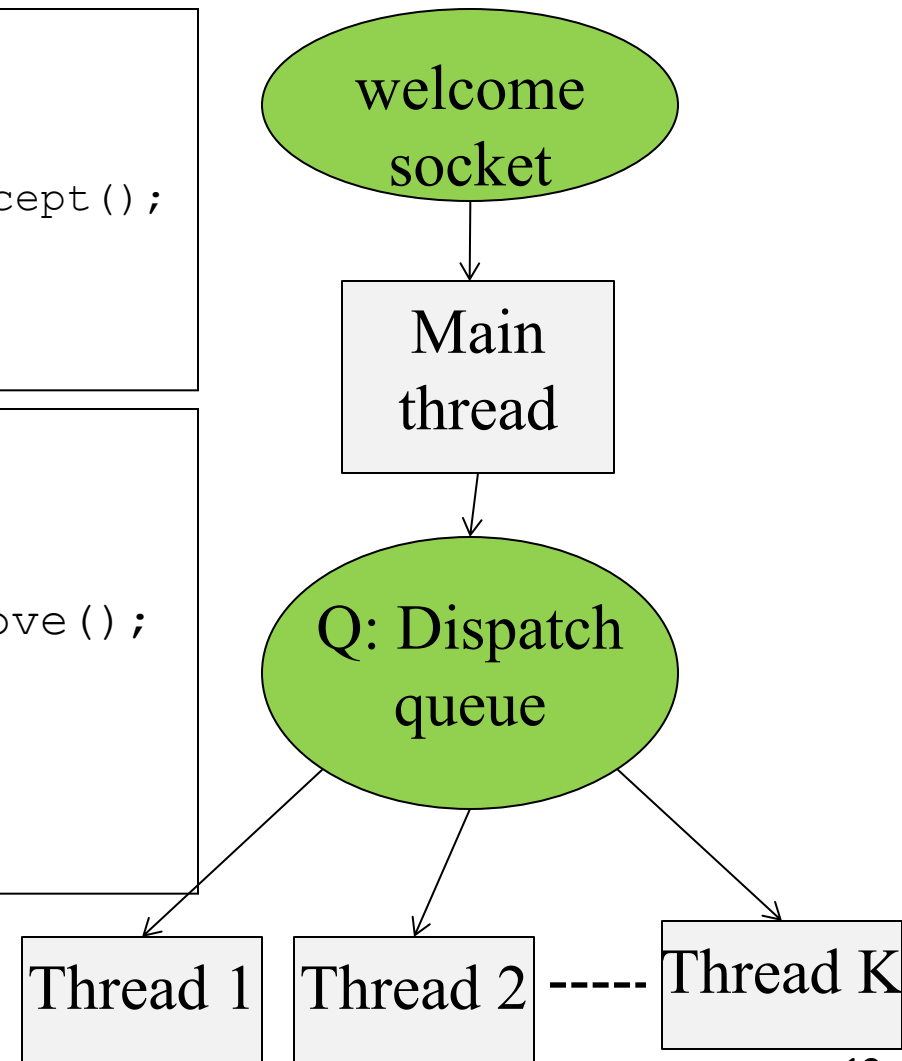


Design 2: Producer/Consumer

```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            Q.add(con);  
        } // end of while  
    }  
}
```

```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = Q.remove();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```

sketch; not
working code



Common Issues Facing Designs 1 and 2

- ❑ Both designs involve multiple threads modifying the same data concurrently
 - Design 1: `welcomeSocket`
 - Design 2: `Q`
- ❑ In our original `WebServerMT`, do we have multiple threads modifying the same data concurrently?

Concurrency and Shared Data

- ❑ Concurrency is easy if threads don't interact
 - Each thread does its own thing, ignoring other threads
 - Typically, however, threads need to communicate/coordinate with each other
 - Communication/coordination among threads is often done by *shared data*

Simple Example

```
public class ShareExample extends Thread {  
    private static int cnt = 10; // shared state, count  
                                // total decrease  
  
    public void run() {  
        int y = cnt;  
        cnt = y - 1;  
    }  
  
    public static void main(String args[]) {  
        Thread t1 = new ShareExample();  
        Thread t2 = new ShareExample();  
        t1.start();  
        t2.start();  
        Thread.sleep(1000);  
        System.out.println("cnt = " + cnt);  
    }  
}
```

Q: What is the result of the program?

Simple Example

What if we add a println:

```
int y = cnt;  
System.out.println("Calculating...");  
cnt = y - 1;
```

What Happened?

- ❑ A thread was preempted in the middle of an operation
- ❑ The operations from reading to writing `cnt` should be *atomic* with no interference access to `cnt` from other threads
- ❑ But the scheduler interleaves threads and caused a *race condition*
- ❑ Such bugs can be extremely hard to reproduce, and also hard to debug

Synchronization

- ❑ Refers to mechanisms allowing a programmer to control the execution order of some operations across different threads in a concurrent program.
- ❑ We use Java as an example to see synchronization mechanisms
- ❑ We'll look at locks first.

Java Lock (1.5)

```
interface Lock {  
    void lock();  
    void unlock();  
    ... /* Some more stuff, also */  
}  
class ReentrantLock implements Lock { ... }
```

- ❑ Only one thread can hold a lock at once
- ❑ Other threads that try to acquire it *block (or become suspended)* until the lock becomes available
- ❑ *Reentrant lock can be reacquired by same thread*
 - As many times as desired
 - No other thread may acquire a lock until it has been released the same number of times that it has been acquired
 - Do not worry about the reentrant perspective, consider it a lock

Java Lock

❑ Fixing the ShareExample.java problem

```
import java.util.concurrent.locks.*;
public class ShareExample extends Thread {
    private static int cnt = 10;
    static Lock lock = new ReentrantLock();

    public void run() {
        lock.lock();
        int y = cnt;
        cnt = y - 1;
        lock.unlock();
    }
    ...
}
```

Java Lock


- ❑ It is recommended to use the following pattern (why?)

```
...
lock.lock();
try {
    // processing body
} finally {
    lock.unlock();
}
```

Background: Java synchronized

- ❑ This pattern is really common
 - Acquire lock, do something, release lock after we are done, **under any circumstances, even if exception was raised, the method returned in the middle, etc.**
- ❑ Java has a language construct for this
 - `synchronized (obj) { body }`
 - ❑ Utilize the design that every Java object has its own **implicitly lock** object, also called the **intrinsic lock**, **monitor lock** or simply **monitor**
 - Obtains the lock associated with **obj**
 - Executes **body**
 - Release lock when scope is exited
 - Even in cases of exception or method return

Background: Java Object and its Associated Lock



- ❑ An object and its associated lock are different !
- ❑ Holding the lock on an object does not affect what you can do with that object in any way
- ❑ Examples:
 - `synchronized(o) { ... } // acquires lock named o`
 - `o.f (); // someone else can call o's methods`
 - `o.x = 3; // someone else can read and write o's fields`

Synchronization on this

```
class C {  
    int cnt;  
    void dec() {  
        synchronized (this) {  
            cnt--;  
        } // end of sync  
    } // end of dec  
}
```

```
C c = new C();
```

```
Thread 1  
c.dec();
```

```
Thread 2  
c.dec();
```

- ❑ A program can often use `this` as the object to lock
- ❑ Does the program above have a data race?
 - No, both threads acquire locks on the same object before they access shared data

Synchronization on this

```
class C {  
    static int cnt;  
    void inc() {  
        synchronized (this) {  
            cnt++;  
        } // end of sync  
    } // end of inc  
  
    void dec() {  
        synchronized (this) {  
            cnt--;  
        } // end of sync  
    } // end of dec  
}
```

```
C c = new C();
```

```
Thread 1  
c.inc();
```

```
Thread 2  
c.dec();
```

- ❑ Does the program above have a data race?
 - No, both threads acquire locks on the same object before they access shared data

Example

- ❑ See

- ShareWelcome/Server.java
- ShareWelcome/ServiceThread.java

- ❑ See `jstack -l` to see lock

Discussion

- ❑ You would not need the lock before `accept` if Java were to label the call as thread safe (synchronized)
- ❑ One reason Java does not specify `accept` as thread safe is that one could register your own socket implementation with [`ServerSocket.setSocketFactory`](#)
- ❑ Always consider thread safety in your design
 - If a resource is shared through concurrent read/write, write/write), consider thread-safe issues.

Why not Synchronization

- ❑ Synchronized method invocations generally are going to be slower than non-synchronized method invocations
- ❑ Synchronization gives rise to the possibility of deadlock, a severe performance problem in which your program appears to hang

Synchronization Overhead

- ❑ Try SyncOverhead.java

Synchronization Overhead


□ Try SyncOverhead.java

Method	Time (ms; 5,000,000 exec)
no sync	9 ms
synchronized method	116 ms
synchronized on this	110 ms
lock	117 ms
lock and finally	113 ms

Design 2: Producer/Consumer

```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            Q.add(con);  
        } // end of while  
    }  
}
```


```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = Q.remove();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```



How to turn it into working code (no race condition)?

Main

```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            Q.add(con);  
        } // end of while  
    }  
}
```



```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            synchronized(Q) {  
                Q.add(con);  
            }  
        } // end of while  
    }  
}
```

Worker

```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = Q.remove();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```



```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        synchronize(Q) {  
            if (!Q.isEmpty())  
                myConn = (Socket) Q.remove();  
        }  
    } // end of while  
    // process myConn  
}
```

Example

- ❑ try
 - ShareQ/Server.java
 - ShareQ/ServiceThread.java

Problem of ShareQ Design

- ❑ Worker thread continually spins (**busy wait**) until cond

```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        Q.lock()  
        if (!Q.isEmpty()) // cond !empty() met  
            myConn = (Socket) Q.remove();  
        } else {}  
        Q.unlock();  
    } // end of while  
    // process myConn  
}
```

- ❑ Can lead to high utilization and slow response time
- ❑ Q: Does the shared welcomeSock have busy-wait?
 - what did jstack -l show?

Outline

- ❑ Admin and recap
- ❑ High-performance network server design
 - Overview
 - Threaded servers
 - Per-request thread
 - problem: large # of threads and their creations/deletions may let overhead grow out of control
 - Thread pool
 - Design 1: Service threads compete on the welcome socket
 - Design 2: Service threads and the main thread coordinate on the shared queue
 - » polling (busy wait)
 - » suspension: wait/notify

Solution: Suspension

- ❑ Put thread to sleep to avoid busy spin
- ❑ Thread life cycle: while a thread executes, it goes through a number of different phases
 - New: created but not yet started
 - Runnable: is running, or can run on a free CPU
 - Blocked: waiting for socket/I/O, a lock, or **suspend** (after call wait())
 - Sleeping: paused for a user-specified interval
 - Terminated: completed

Solution: Suspension

```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        lock Q;  
        if (!Q.isEmpty()) // {  
            // get myConn from Q  
        } else {  
            // stop and wait ← Hold lock?  
        }  
        unlock Q;  
    }  
    // get the next request; process  
}
```

Solution: Suspension




```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        lock Q;  
        if (!Q.isEmpty()) // {  
            // get myConn from Q  
        } else {  
            // stop and wait  
        }  
        unlock Q;  
    }  
    // get the next request; process  
}
```

Design pattern:

- Need to release lock to avoid deadlock (to allow write into Q)
- Typically need to reacquire lock after waking up


Wait-sets and Notification

- ❑ Every Java Object has an associated wait-set (called wait list) in addition to a lock object



Wait-sets and Notification

- ❑ Wait list object can be **manipulated only while the object lock is held**
 - Otherwise, `IllegalMonitorStateException` is thrown



Wait-sets: Wait (suspend)

- ❑ Thread enters the wait-set of an object by invoking the object's `wait()` method
 - `wait()` releases the lock
 - No other held locks are released
 - then the thread is suspended
- ❑ Can add optional time `wait(long millis)`
 - `wait()` is equivalent to `wait(0)` - wait forever
 - for robust programs, it is typically a good idea to add a timer

Wait-set: Wakeup (release)

- ❑ Threads are released from the wait-set when:
 - `notifyAll()` is invoked on the object
 - All threads released (typically recommended)
 - `notify()` is invoked on the object
 - One thread selected at 'random' for release
 - The specified time-out elapses
 - The thread has its `interrupt()` method invoked
 - `InterruptedException` thrown
 - A spurious wakeup occurs
 - Not (yet!) spec'd but an inherited property of underlying synchronization mechanisms e.g., POSIX condition variables

Wait-set: Notify


- ❑ Caller of `notify()` must hold lock associated with the object
- ❑ Those threads awoken must reacquire lock before continuing
 - (This is part of the function; you don't need to do it explicitly)
 - Can't be acquired until notifying thread releases it
 - A released thread contends with all other threads for the lock

Exercise: Dispatcher

```
main {  
  void run {  
    while (true) {  
      Socket con = welcomeSocket.accept();  
      synchronized(Q) {  
        Q.add(con);  
      }  
    } // end of while  
  }  
}
```





```
main {  
  void run {  
    while (true) {  
      Socket con = welcomeSocket.accept();  
      synchronize(Q) {  
        Q.add(con);  
        Q.notifyAll();  
      }  
    } // end of while  
  }  
}
```



Worker

```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        synchronize(Q) {  
            if (! Q.isEmpty()) // {  
                myConn = Q.remove();  
            }  
        }  
    } // end of while  
    // process myConn  
}
```


Busy wait



```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        synchronize(Q) {  
            if (! Q.isEmpty()) // {  
                myConn = Q.remove();  
            } else {  
                Q.wait();  
            }  
        }  
    } // end of while  
    // process myConn  
}
```



Main
thread




Thread 1

Thread K

Worker: Another Format

```
while (true) {  
    // get next request  
    Socket myConn = null;  
    synchronized(Q) {  
        while (Q.isEmpty()) {  
            Q.wait();  
        }  
        myConn = Q.remove();  
    } // end of sync  
    // process request in myConn  
} // end of while
```

Note the while
loop; no guarantee
that Q is not empty
when wake up



Example

□ See

- WaitNotify/Server.java
- WaitNotify/ServiceThread.java

Summary: Guardian via Suspension: Waiting

```
synchronized (obj) {  
    while (!condition) {  
        try { obj.wait(); }  
        catch (InterruptedException ex)  
        { ... }  
    } // end while  
    // make use of condition  
} // end of sync
```

- ❑ **Golden rule:** Always test a condition in a loop
 - Change of state may not be what you need
 - Condition may have changed again
- ❑ Break the rule only after you are sure that it is safe to do so

Summary: Guarding via Suspension: Changing a Condition

```
synchronized (obj) {  
    condition = true;  
    obj.notifyAll(); // or obj.notify()  
}
```

- ❑ Typically use `notifyAll()`
- ❑ There are subtle issues using `notify()`, in particular when there is interrupt

Blackbox: Use Java ThreadPoolExecutor

```
public class TimeServerHandlerExecutePool {  
  
    private ExecutorService executor;  
  
    public TimeServerHandlerExecutePool(int maxPoolSize, int queueSize) {  
        executor = new ThreadPoolExecutor(  
            Runtime.getRuntime().availableProcessors(),  
            maxPoolSize,  
            120L, TimeUnit.SECONDS,  
            new ArrayBlockingQueue<java.lang.Runnable>(queueSize)  
        );  
    }  
  
    public void execute(java.lang.Runnable task) {  
        executor.execute(task);  
    }  
}
```

For Java ThreadPoolExecutor scheduling algorithm, see:

<https://docs.oracle.com/javase/10/docs/api/java/util/concurrent/ThreadPoolExecutor.html>

Blackbox: Use Java ThreadPoolExecutor

```
server = new ServerSocket(port);  
System.out.println("Time server listens at port: " + port);
```

```
// Create Java Executor Pool  
TimeServerHandlerExecutePool myExecutor  
    = new TimeServerHandlerExecutePool(50, 10000);
```

```
Socket socket = null;  
while (true) {  
    socket = server.accept();  
    myExecutor.execute(new TimeServerHandler(socket));  
} // end of while
```

Note: Beyond Java

- ❑ Use of wait(), notifyAll() and notify() similar to
 - Condition queues of classic Monitors
 - Condition variables of POSIX PThreads API
 - In C# it is called Monitor (<http://msdn.microsoft.com/en-us/library/ms173179.aspx>)

- ❑ Python Thread module in its Standard Library is based on Java Thread model
(<https://docs.python.org/3/library/threading.html>)
 - "The design of this module is loosely based on Java's threading model. However, where Java makes locks and condition variables basic behavior of every object, they are separate objects in Python."

Whitebox: Java (1.5)

```
interface Lock { Condition newCondition(); ... }  
interface Condition {  
    void await();  
    void signalAll(); ...  
}
```


❑ Condition created from a Lock

❑ `await` called with lock held

- Releases the lock
 - But not any other locks held by this thread
- Adds this thread to wait set for lock
- Blocks the thread

❑ `signalAll` called with lock held

- Resumes all threads on lock's wait set
- Those threads must reacquire lock before continuing
 - (This is part of the function; you don't need to do it explicitly)



Beyond Class: Producer/Consumer Example

```
Lock lock = new ReentrantLock();
Condition ready = lock.newCondition();
boolean valueReady = false;
Object value;
```

```
void produce(Object o) {
    lock.lock();
    while (valueReady)
        ready.await();
    value = o;
    valueReady = true;
    ready.signalAll();
    lock.unlock();
}

Object consume() {
    lock.lock();
    while (!valueReady)
        ready.await();
    Object o = value;
    valueReady = false;
    ready.signalAll();
    lock.unlock();
}
```

Beyond Class: Blocking Queues in Java

- ❑ Design Pattern for producer/consumer pattern with blocking, e.g.,
 - put/take
- ❑ Two handy implementations
 - `LinkedBlockingQueue` (FIFO, may be bounded)
 - `ArrayBlockingQueue` (FIFO, bounded)
 - (plus a couple more)

<https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/BlockingQueue.html>

Beyond Class: Complete Java Concurrency Framework

Executors

- **Executor**
- **ExecutorService**
- **ScheduledExecutorService**
- Callable
- **Future**
- ScheduledFuture
- Delayed
- CompletionService
- **ThreadPoolExecutor**
- **ScheduledThreadPoolExecutor**
- AbstractExecutorService
- Executors
- FutureTask
- ExecutorCompletionService

Queues

- BlockingQueue
- ConcurrentLinkedQueue
- LinkedBlockingQueue
- ArrayBlockingQueue
- SynchronousQueue
- PriorityBlockingQueue
- DelayQueue

Concurrent Collections

- ConcurrentMap
- ConcurrentHashMap
- CopyOnWriteArray{List,Set}

Synchronizers

- CountdownLatch
- Semaphore
- Exchanger
- CyclicBarrier

Locks: `java.util.concurrent.locks`

- Lock
- Condition
- ReadWriteLock
- AbstractQueuedSynchronizer
- LockSupport
- ReentrantLock
- ReentrantReadWriteLock

Atomics: `java.util.concurrent.atomic`

- Atomic{Type}
- Atomic{Type}Array
- Atomic{Type}FieldUpdater
- Atomic{Markable,Stampable}Reference

See jcf slides for a tutorial.

Outline

- ❑ Admin and recap
- ❑ High-performance network server design
 - Overview
 - Threaded servers
 - Per-request thread
 - problem: large # of threads and their creations/deletions may let overhead grow out of control
 - Thread pool
 - Design 1: Service threads compete on the welcome socket
 - Design 2: Service threads and the main thread coordinate on the shared queue
 - » polling (busy wait)
 - » suspension: wait/notify
 - Basic correctness analysis

Correctness

❑ Threaded programs are typically more complex.

❑ What types of properties do you analyze concurrent server for correctness?

```
// worker
void run() {
    while (true) {
        // get next request
        Socket myConn = null;
        synchronized(Q) {
            while (Q.isEmpty()) {
                Q.wait();
            } // end of while
            myConn = Q.remove();
        } // end of sync
        // process request in myConn
    } // end of while
} // end of run()
```

```
// master
void run() {
    while (true) {
        Socket con = welcomeSocket.accept();
        synchronize(Q) {
            Q.add(con);
            Q.notifyAll();
        } // end of sync
    } // end of while
} // end of run()
```

Key Correctness Properties

- Safety

- Liveness (progress)

- Fairness

- For example, in some settings, a designer may want the threads to share load equally

Safety Properties

- ❑ What safety properties?
 - No read/write; write/write conflicts
 - holding lock `Q` before reading or modifying shared data `Q` and `Q.wait_list`
 - `Q.remove()` is not on an empty queue
- ❑ There are formal techniques to model server programs and analyze their properties, but we will use basic analysis
 - This is enough in many cases

Make Program Explicit

```
// dispatcher
void run() {
    while (true) {
        Socket con = welcomeSocket.accept();
        synchronize(Q) {
            Q.add(con);
            Q.notifyAll();
        } // end of sync
    } // end of while
} // end of run()
```

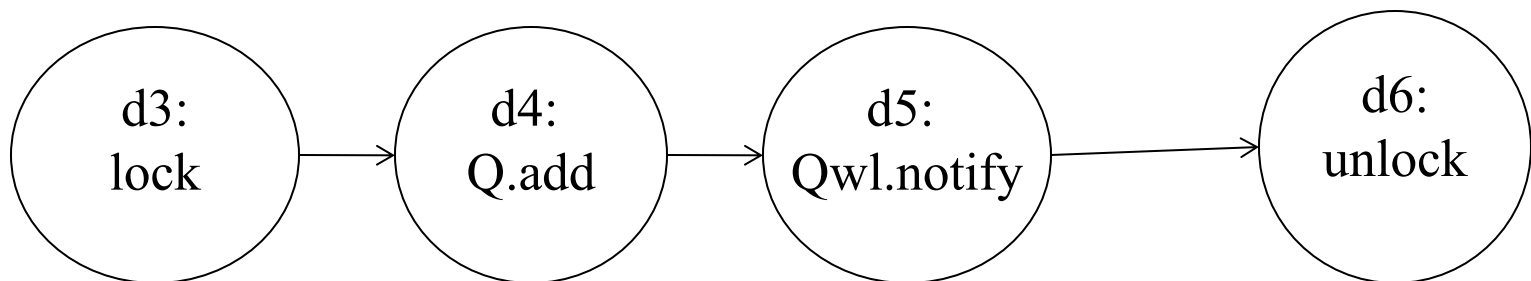
```
// dispatcher
void run() {
1.   while (true) {
2.       Socket con = welcomeSocket.accept();
3.       lock(Q) {
4.           Q.add(con);
5.           notify Q.wait_list; // Q.notifyAll();
6.           unlock(Q);
           } // end of while
       } // end of run()
```

```
// service thread
void run() {
    while (true) {
        // get next request
        Socket myConn = null;
        synchronized(Q) {
            while (Q.isEmpty()) {
                Q.wait();
            } // end of while
            myConn = Q.remove();
        } // end of sync
        // process request in myConn
    } // end of while
}
```

```
// service thread
void run() {
1. while (true) {
    // get next request
2.     Socket myConn = null;
3.     lock(Q);
4.     while (Q.isEmpty()) {
5.         unlock(Q)
6.         add to Q.wait_list;
7.         yield until marked to wake; //wait
8.         lock(Q);
9.     } // end of while
10.    myConn = Q.remove();
11.    unlock(Q);
    // process request in myConn
    } // end of while
}
```

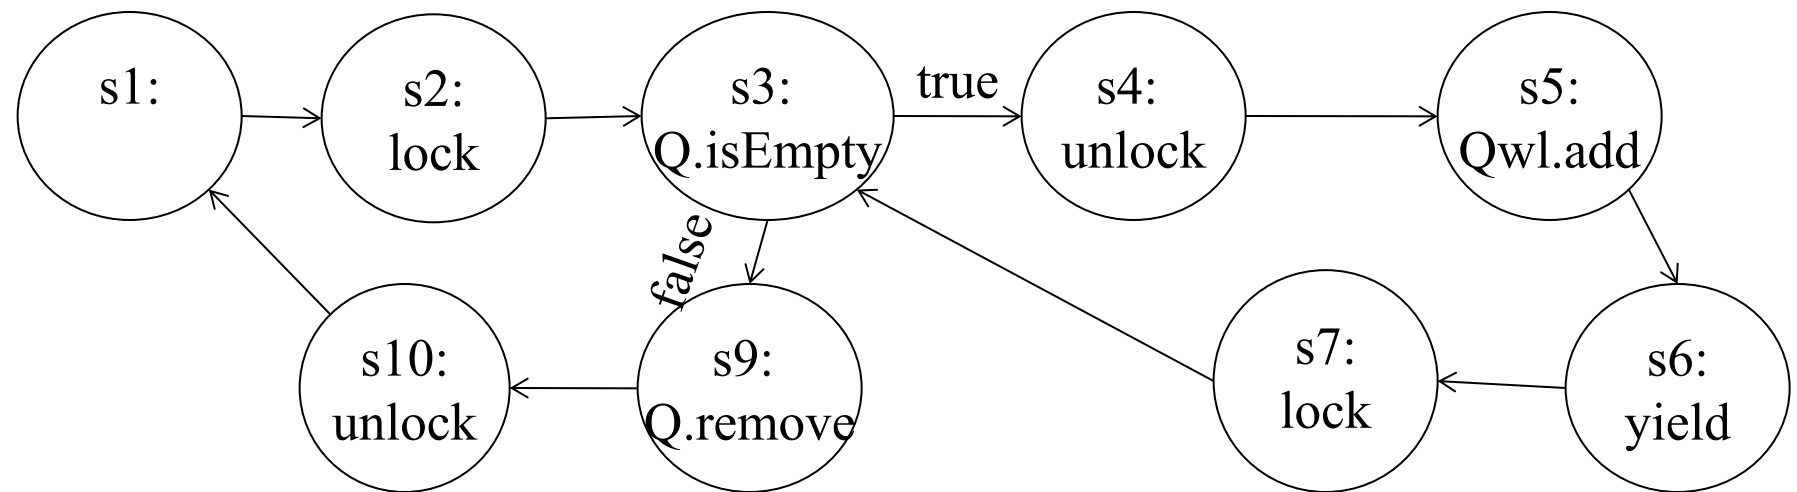
Statements to States (Dispatcher)

```
// dispatcher
void run() {
1.     while (true) {
2.         Socket con = welcomeSocket.accept();
3.         lock(Q) {
4.             Q.add(con);
5.             notify Q.wait_list; // Q.notifyAll();
6.             unlock(Q);
           } // end of while
} // end of run()
```

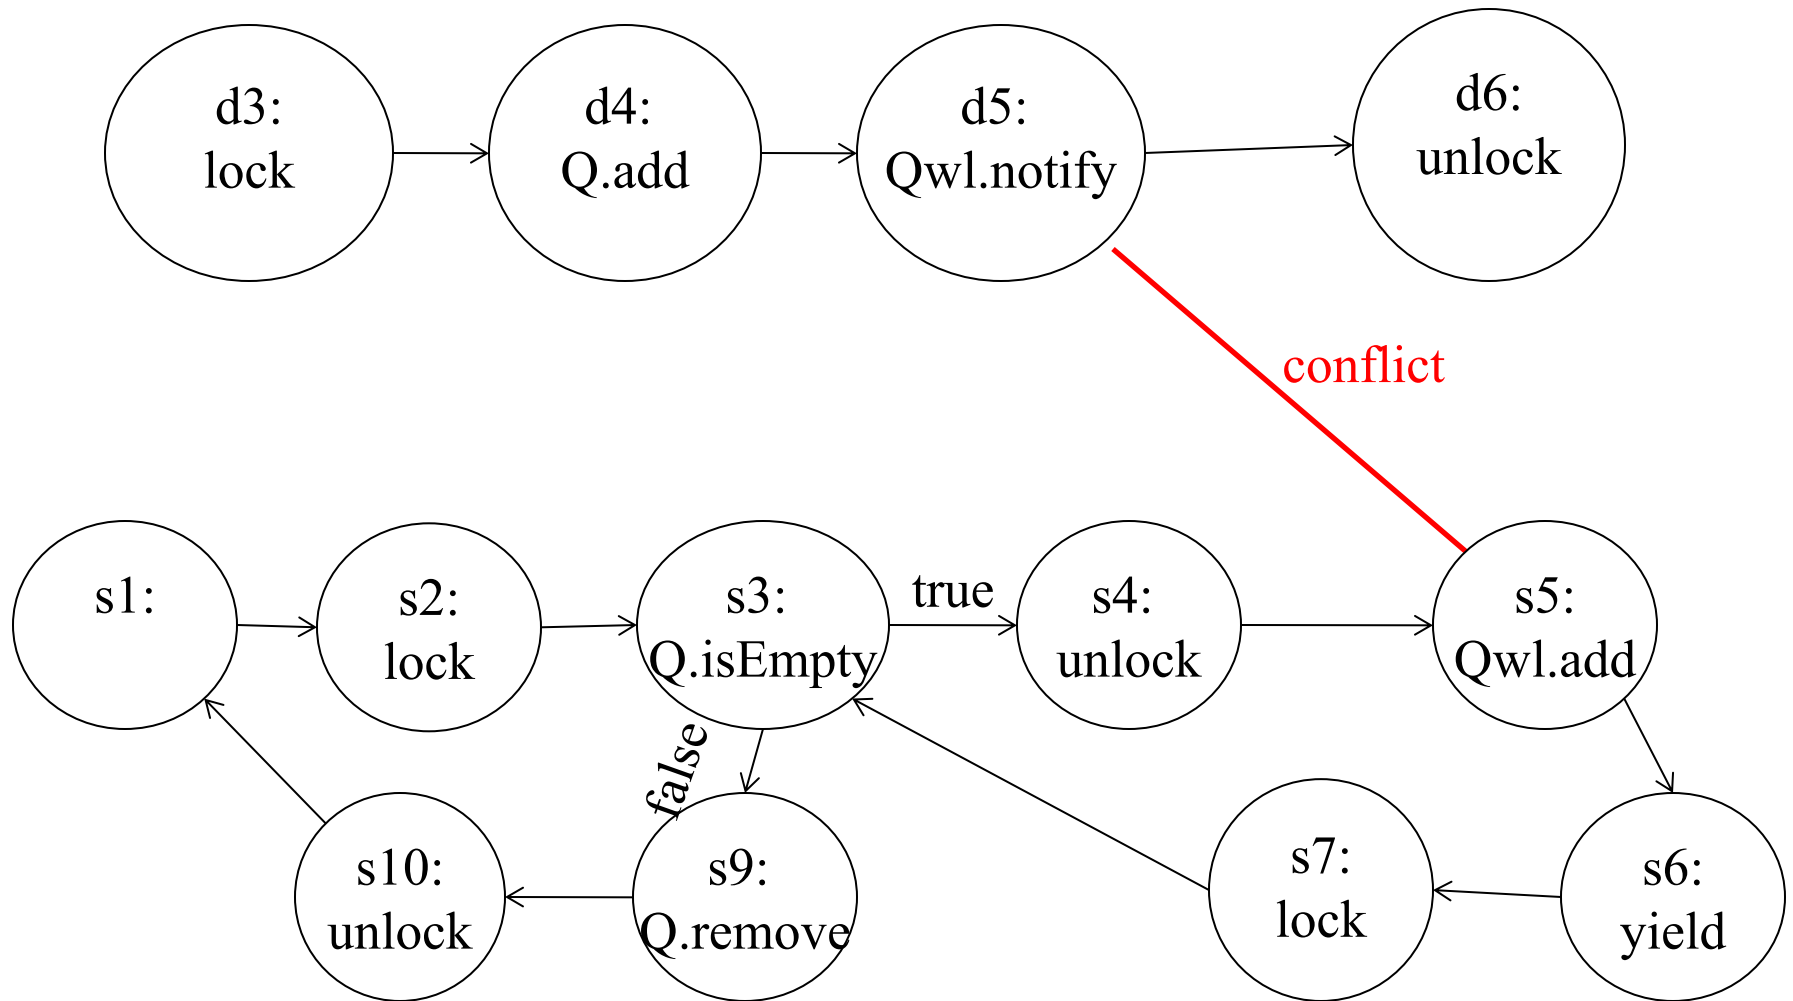


Statements to States (Service)

```
while (true) {  
    // get next request  
1.   Socket myConn = null;  
2.   lock(Q);  
3.   while (Q.isEmpty()) {  
4.       unlock(Q)  
5.       add to Q.wait_list;  
6.       yield; //wait  
7.       lock(Q);  
8.   } // end of while isEmpty  
9.   myConn = Q.remove();  
10.  unlock(Q);  
    // process request in myConn  
} // end of while
```



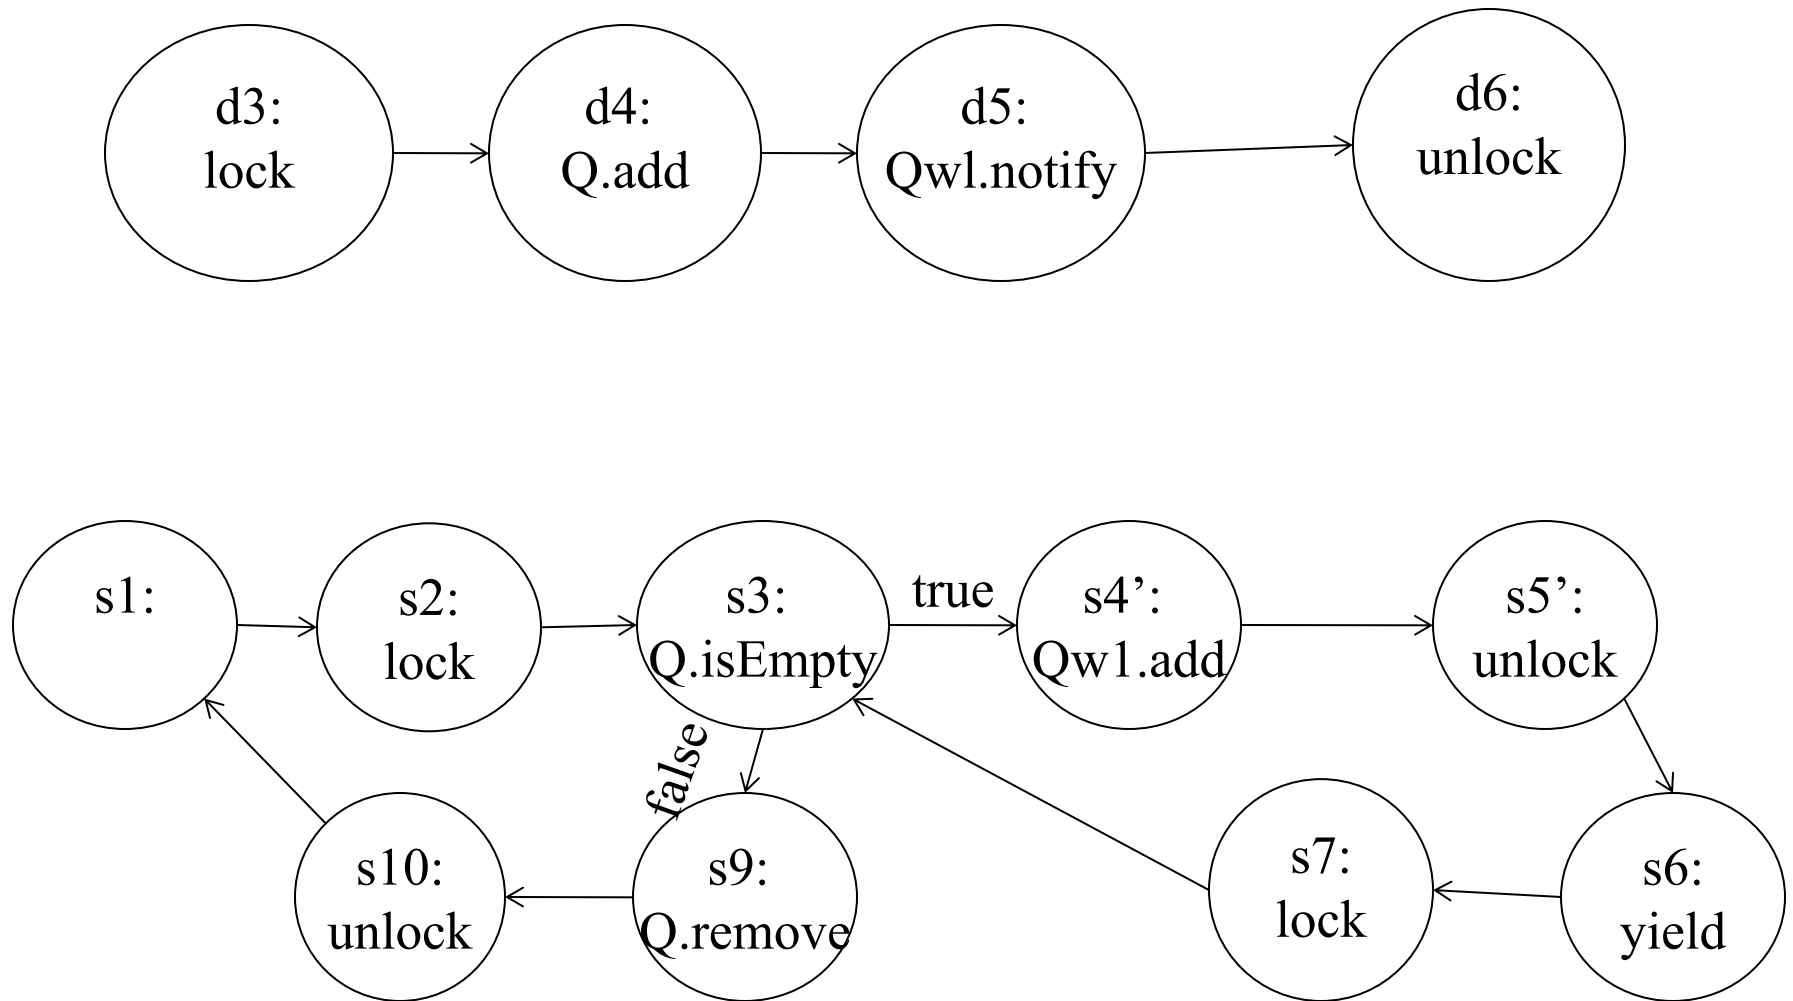
Check Safety



Real Implementation of wait

```
while (true) {  
    // get next request  
1.   Socket myConn = null;  
2.   lock(Q);  
3.   while (Q.isEmpty()) {  
4.       add to Q.wait_list;  
5.       unlock(Q); after add to wait list  
6.       yield; //wait  
7.       lock(Q);  
8.   }  
9.   myConn = Q.remove();  
10.  unlock(Q);  
    // process request in myConn  
} // end of while
```

Check Safety

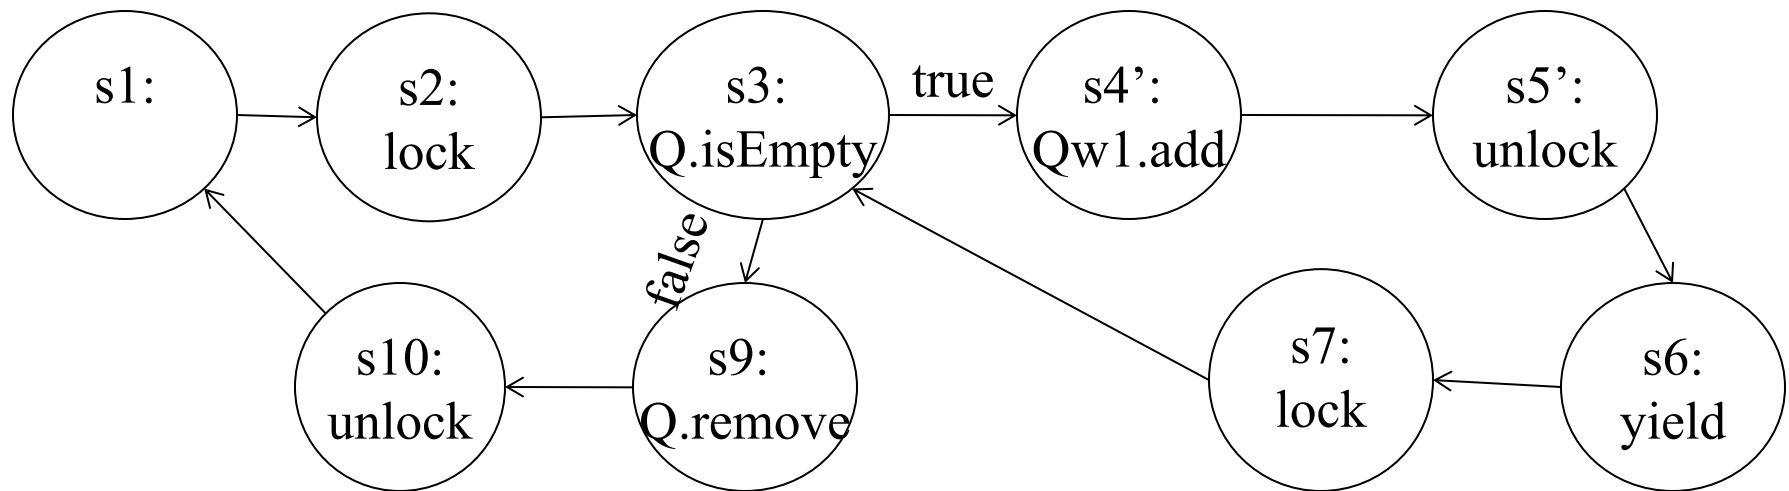


Liveness Properties

- ❑ What liveness (progress) properties?
 - dispatcher thread can always add to Q
 - every connection in Q will be processed

Dispatcher Thread Can Always Add to Q

- ❑ Assume dispatcher thread is blocked
- ❑ Suppose Q is not empty, then each iteration removes one element from Q
- ❑ In finite number of iterations, all elements in Q are removed and all service threads unlock and block
 - Need to assume each service takes finite amount of time (bound by a fixed T_0)



Each Connection in Q is Processed

- ❑ Cannot be guaranteed unless
 - there is fairness in the thread scheduler, or
 - put a limit on Q size to block the dispatcher thread

Summary: Server Correctness Analysis

❑ Safety

- No read/write; write/write conflicts
 - holding lock Q before reading or modifying shared data Q and Q.wait_list
- Q.remove() is not on an empty queue


❑ Liveness (progress)

- dispatcher thread can always add to Q
- every connection in Q will be processed

❑ Fairness

- For example, in some settings, a designer may want the threads to share load equally

Summary: Using Threads



Advantages

- Intuitive (sequential) programming model
- Shared address space simplifies optimizations

Disadvantages

- Overhead: thread stacks, synchronization
- Thread pool parameter (how many threads) difficult to tune

Should You Use Threads?

❑ Typically avoid threads for io

- Use event-driven, not threads, for GUIs, servers, distributed systems.

❑ Use threads where true CPU concurrency is needed.

- Where threads needed, isolate usage in threaded application kernel: keep most of code single-threaded.

