
Network Applications:
High-Performance Server Design
(Async Select NonBlocking Servers)

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

10/9/2018

Admin

- ❑ Assignment Three (HTTP server) Part 1 check point
- ❑ Assignment Part 2 to be posted on Wednesday

Recap: Thread-Based Network Servers

- ❑ Why: blocking operations; threads (execution sequences) so that only one thread is blocked
- ❑ How:
 - Per-request thread
 - problem: large # of threads and their creations/deletions may let overhead grow out of control
 - Thread pool
 - Design 1: Service threads compete on the welcome socket
 - Design 2: Service threads and the dispatcher thread coordinate on a shared queue
 - polling (busy wait)
 - suspension: wait/notify
 - An example control see <http://httpd.apache.org/docs/2.4/mod/worker.html>

Recap: Program Correctness Analysis

❑ Safety

- Consistency (exclusive access)
- app requirement, e.g., `Q.remove()` is not on an empty queue

❑ Liveness (progress)

- main thread can always add to `Q`
- every connection in `Q` will be processed

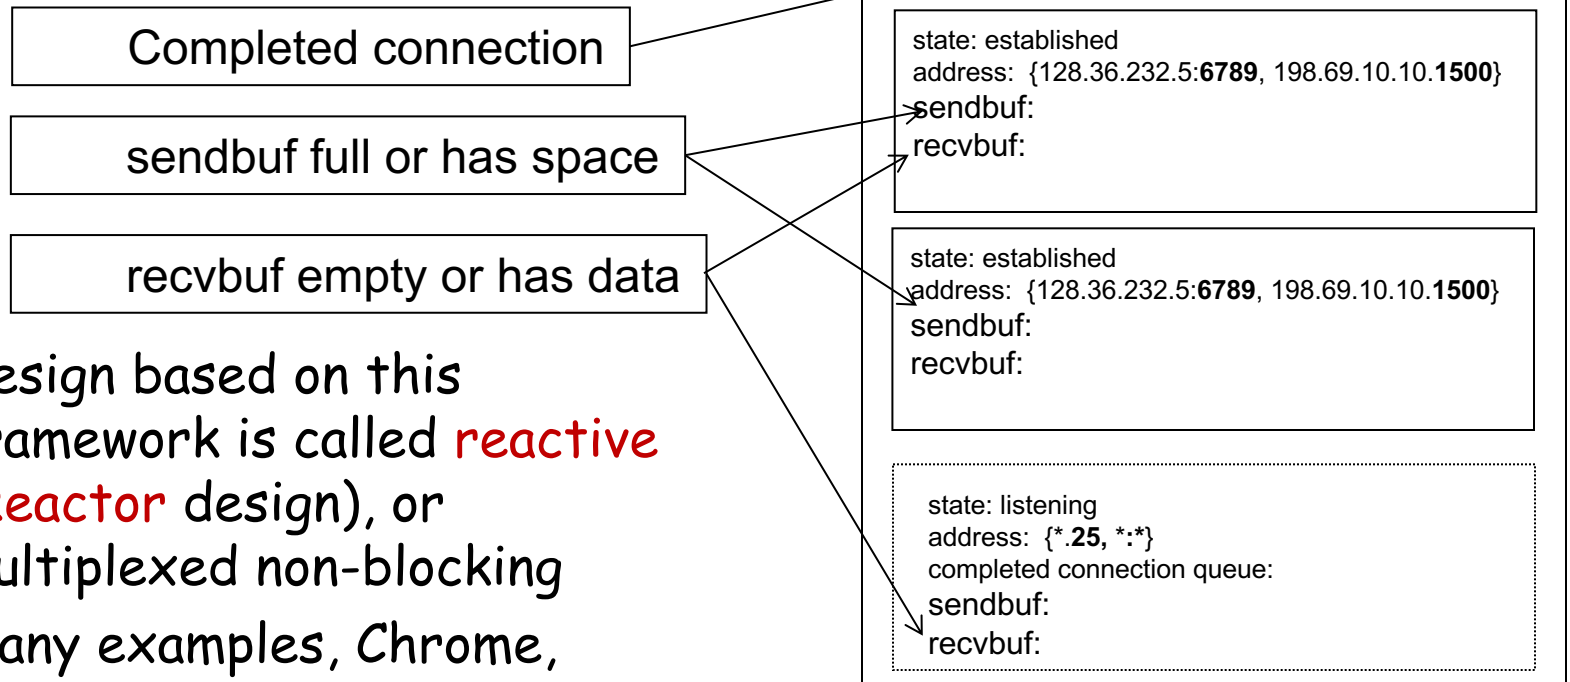
❑ Fairness

- For example, in some settings, a designer may want the threads to share load equally

Recap: Multiplexed, Reactive I/O

- A different approach for avoiding blocking: **peek** system state, issue function calls only for those that are **ready**

- Linux: select, epoll (2.6)
- Mac/FreeBSD: kqueue




- Design based on this framework is called **reactive** (**Reactor** design), or multiplexed non-blocking
- Many examples, Chrome, Dropbox, nginx

Outline


- ❑ Admin and recap
- ❑ High performance servers
 - Thread design
 - Per-request thread
 - Thread pool
 - Busy wait
 - Wait/notify
 - Asynchronous design
 - Overview
 - Multiplexed (selected), reactive programming

Multiplexed, Reactive Server Architecture



- ❑ Program registers events (e.g., acceptable, readable, writable) on channels (sources) to be monitored
- ❑ An infinite dispatcher loop:
 - Dispatcher asks OS to check if any ready event
 - Dispatcher calls (**multiplexes**) the handler of each ready event of each source
 - **Handler should be non-blocking**, to avoid blocking the event loop

Multiplexed, Non-Blocking Network Server



```
// clients register interests/handlers
on events/sources
while (true) {
    - ready events = select()
      /* or selectNow(),
        or select(int timeout) to
        check ready events from the
        registered interests */

    - foreach ready event {
        switch event type:
        accept: call accept handler
        readable: call read handler
        writable: call write handler
    }

    - handle other events
}
```


Main Abstractions

- ❑ Main abstractions of multiplexed IO for network servers:
 - Channel (source): represents a connection to an entity capable of performing I/O operations;
 - Selection facilities;
 - Protocol control block (PCB): container to keep state/handler for each event/channel.
- ❑ Java abstractions see:
<https://docs.oracle.com/javase/8/docs/api/java/nio/package-summary.html>

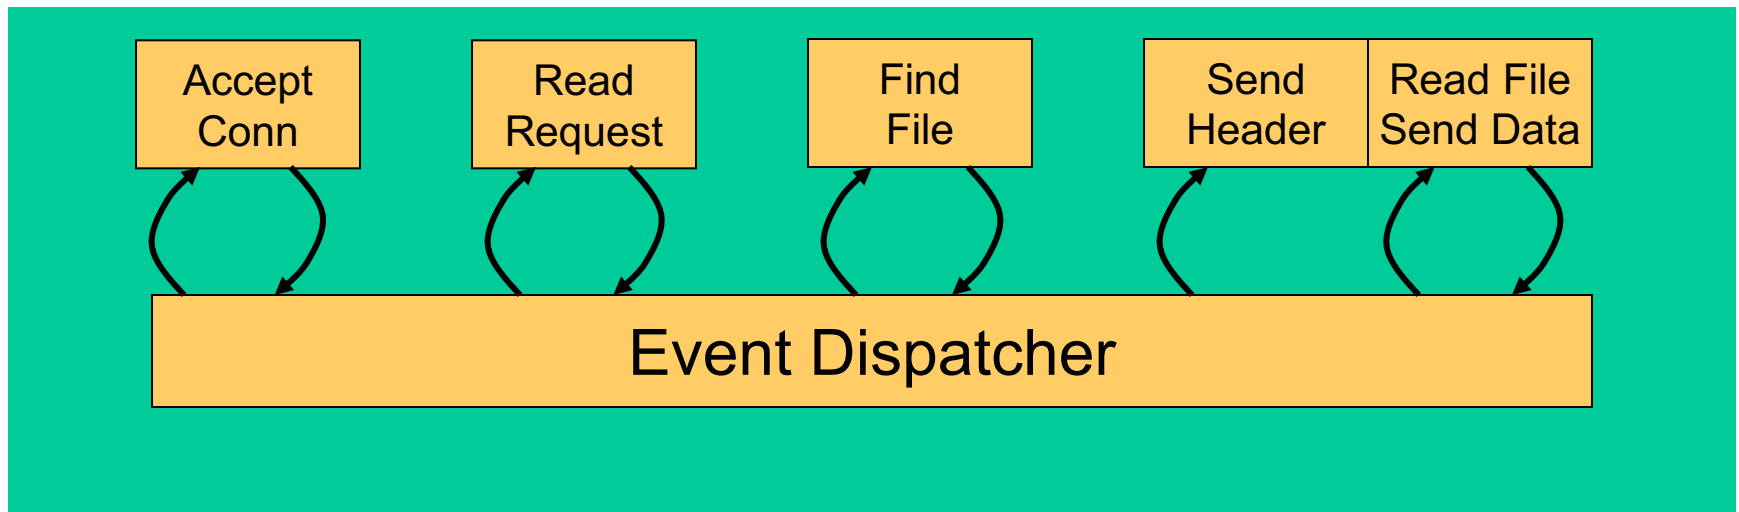
Multiplexed (Selectable), Non-Blocking Channels

<code>SelectableChannel</code>	A channel that can be multiplexed
<code>DatagramChannel</code>	A channel to a datagram-oriented socket
<code>Pipe.SinkChannel</code>	The write end of a pipe
<code>Pipe.SourceChannel</code>	The read end of a pipe
<code>ServerSocketChannel</code>	A channel to a stream-oriented listening socket
<code>SocketChannel</code>	A channel for a stream-oriented connecting socket

- Use `configureBlocking(false)` to make a channel non-blocking
- Note: Java `SelectableChannel` does not include file I/O

Selector

- ❑ The class `Selector` is the base of the multiplexer/dispatcher
- ❑ Constructor of `Selector` is protected; create by invoking the `open` method to get a selector (which design pattern?)




Selector and Registration

- ❑ A selectable channel registers events to be monitored with a `selector` with the `register` method
- ❑ The registration returns an object called a `SelectionKey`:

```
SelectionKey key =  
    channel.register(selector, ops);
```

Java Selection I/O Structure

- ❑ A `SelectionKey` object stores:
 - **interest set**: events to check:
`key.interestOps (ops)`
 - **ready set**: after calling `select`, it contains the events that are ready, e.g.
`key.isReadable ()`
 - **an attachment** that you can store anything you want, typically `PCB`
`key.attach (myObj)`



Checking Events

- ❑ A program calls `select` (or `selectNow()`, or `select(int timeout)`) to check for ready events from the registered `SelectableChannels`
 - Ready events are called the selected key set

```
selector.select();  
Set readyKeys = selector.selectedKeys();
```
- ❑ The program iterates over the selected key set to process all ready events


I/O in Java: ByteBuffer

- ❑ Java SelectableChannels typically use ByteBuffer for read and write
 - `channel.read(byteBuffer);`
 - `channel.write(byteBuffer);`
- ❑ ByteBuffer is a powerful class that can be used for both read and write
- ❑ It is derived from the class Buffer
- ❑ Please be sure to read these data structures

Java ByteBuffer Hierarchy


Buffers	Description
<u>Buffer</u>	Position, limit, and capacity; clear, flip, rewind, and mark/reset
<u>ByteBuffer</u>	Get/put, compact, views; allocate, wrap
<u>MappedByteBuffer</u>	A byte buffer mapped to a file
<u>CharBuffer</u>	Get/put, compact; allocate, wrap
<u>DoubleBuffer</u>	' '
<u>FloatBuffer</u>	' '
<u>IntBuffer</u>	' '
<u>LongBuffer</u>	' '
<u>ShortBuffer</u>	' '

Buffer (relative index)




- ❑ Each Buffer has three numbers: position, limit, and capacity
 - Invariant: $0 \leq \text{position} \leq \text{limit} \leq \text{capacity}$
- ❑ `Buffer.clear()`: `position = 0; limit=capacity`

channel.read(Buffer)




- ❑ Put data into Buffer, starting at `position`, not to reach `limit`

channel.write(Buffer)




- ❑ Move data from Buffer to channel, starting at `position`, not to reach `limit`

Buffer.flip()



- ❑ `Buffer.flip()`: `limit=position; position=0`
- ❑ Why flip: used to switch from preparing data to output, e.g.,
 - `buf.put(header); // add header data to buf`
 - `in.read(buf); // read in data and add to buf`
 - `buf.flip(); // prepare for write`
 - `out.write(buf);`
- ❑ Typical pattern: read, flip, write

Buffer.compact()



- ❑ Move [position , limit) to 0
- ❑ Set position to limit-position, limit to capacity

// typical design pattern

```
buf.clear(); // Prepare buffer for use
for (;;) {
    if (in.read(buf) < 0 && !buf.hasRemaining())
        break; // No more bytes to transfer
    buf.flip();
    out.write(buf);
    buf.compact(); // In case of partial write
}
```

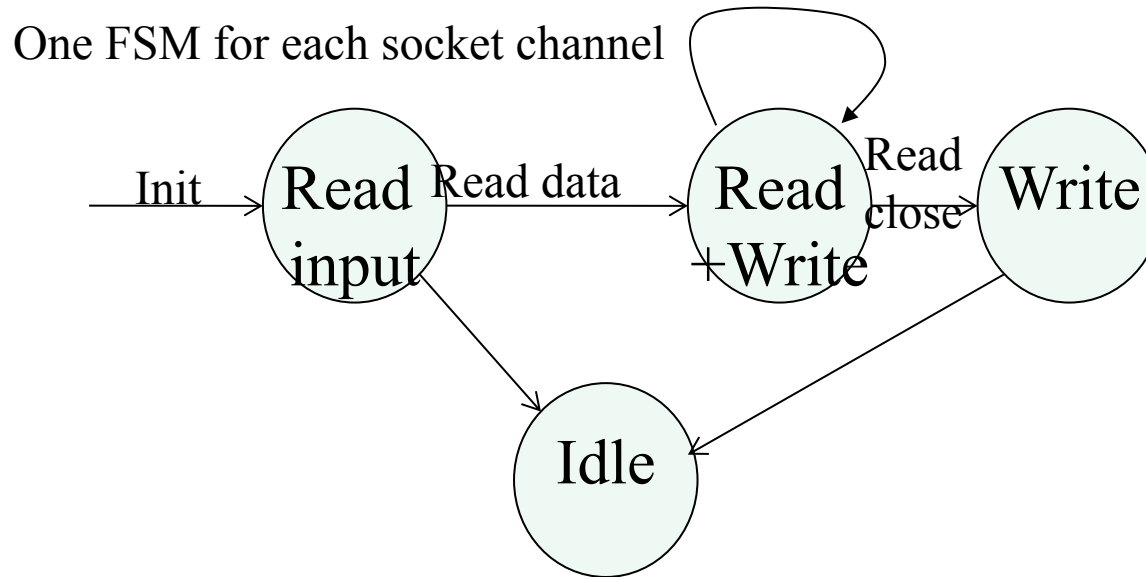
Example and Design Exercise

- See
SelectEchoServer/v1/SelectEchoServer.java

Summary: Steps We Took to Refine the Echo Server

- ❑ Register READ for newly accepted connection
 - otherwise, no read events
- ❑ Register only READ, not WRITE
 - otherwise empty write
 - Imagine empty write with 10,000 sockets
- ❑ After read data, turn on write to enable echo output
 - otherwise no output
- ❑ After write, check if there is data remaining to write, if no, turn off write
 - otherwise, empty write calls
- ❑ After reading end of stream (read returns -1), turn off read interest (or better deregister)
- ❑ All above are state management!

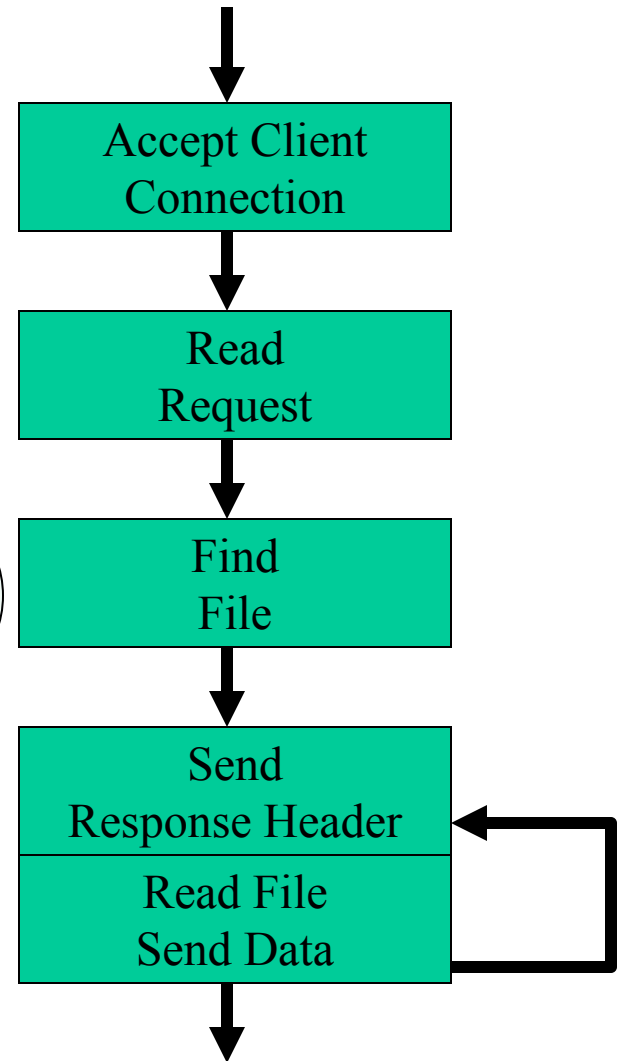

Finite-State Machine and Async Server



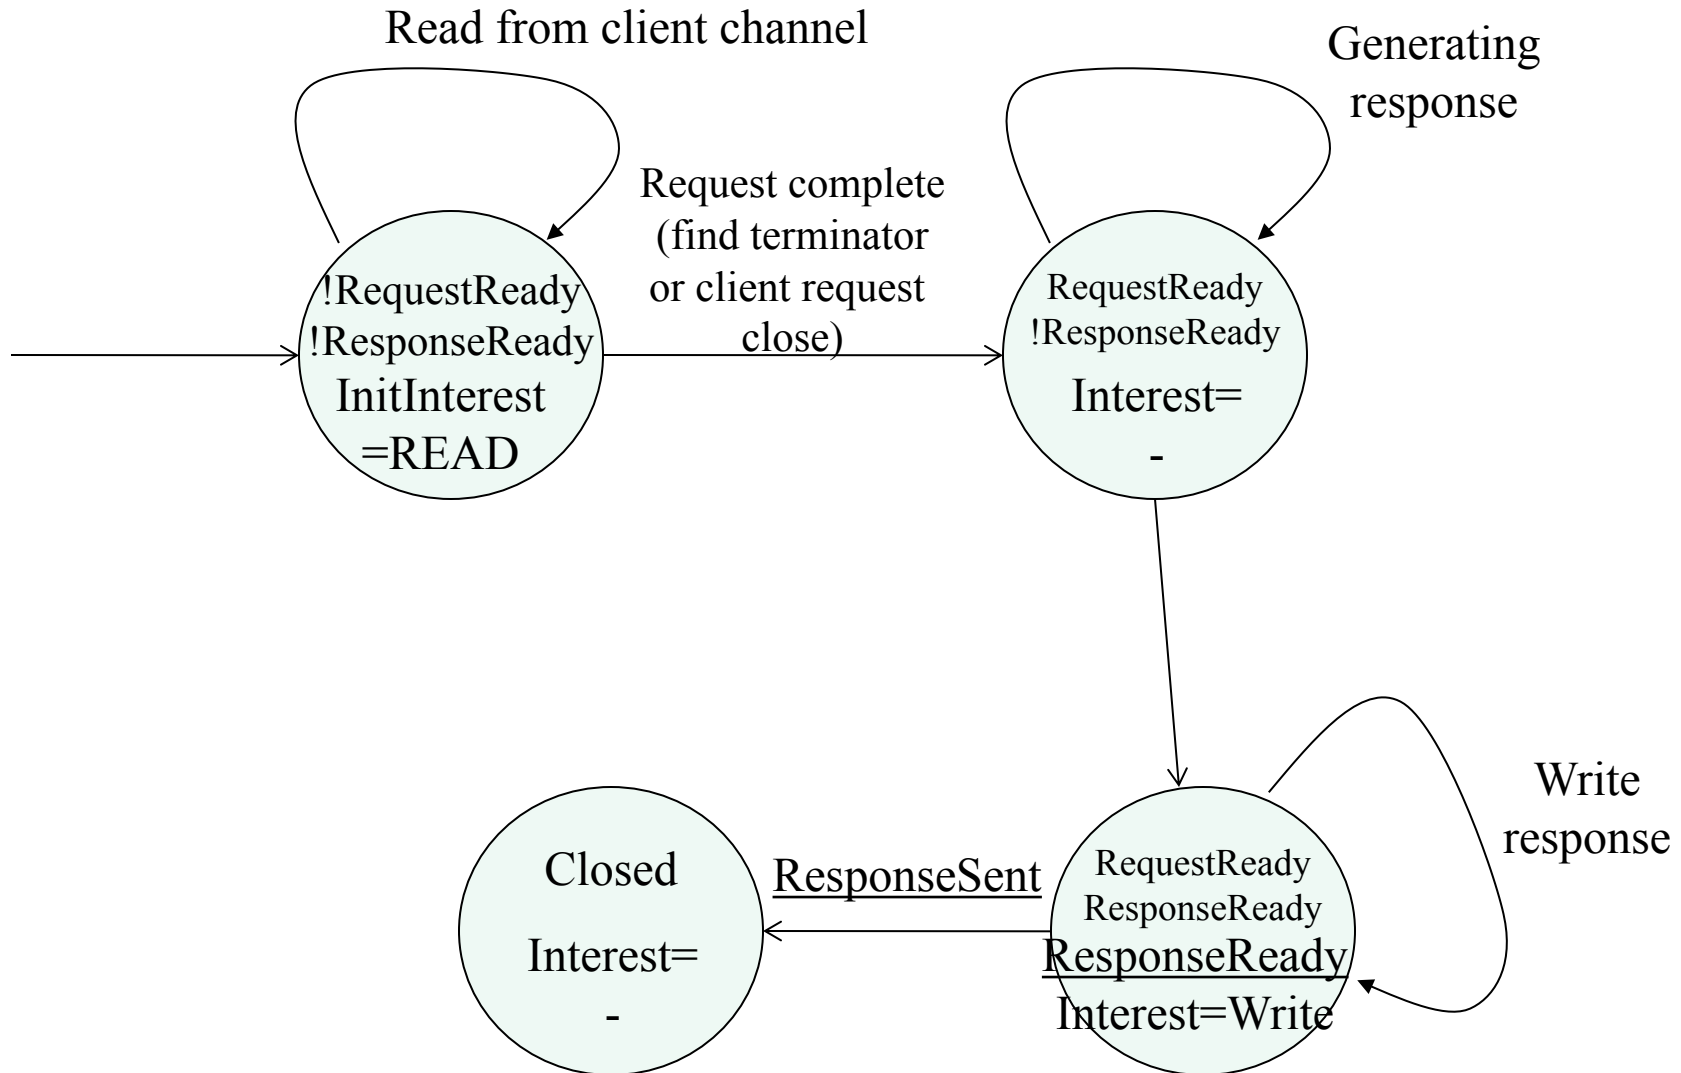
Not the most effective FSM, but an example.

Finite-State Machine and Thread

- Why no need to introduce FSM for a thread version?



A More Typical Finite State Machine



FSM and Reactive Programming

- ❑ Designing the FSM is key to non-blocking servers, and there can be multiple types of FSMs, to handle protocols correctly
 - Staged: first read request and then write response
 - Mixed: read and write mixed
- ❑ Choice depends on protocol and tolerance of complexity, e.g.,
 - HTTP/1.0 channel may use staged
 - HTTP/1.1/2/Chat channel may use mixed

Toward More Generic Select Server Software Framework

- ❑ Non-blocking, select programming framework is among the more complex software systems, and we want to reuse the software as much as possible
 - E.g., consider a setting where a single server monitors multiple ports, with each port may run a different protocol
- ❑ Question: Which design of the EchoServer is not generic (i.e., reusable for different protocols)?

EchoServer Design Issues

- ❑ Fixed accept/read/write functions (handlers) are not general design
- ❑ PCB is customized for echo servers only

A More Extensible Dispatcher Design

□ Attachment stores generic event handler

○ Define interfaces

- IAcceptHandler and
- IReadWriteHandler

○ Retrieve handlers at run time

```
if (key.isAcceptable()) { // a new connection is ready
    IAcceptHandler aH = (IAcceptHandler) key.attachment();
    aH.handleAccept(key);
}

if (key.isReadable() || key.isWritable()) {
    IReadWriteHandler rwH = (IReadWriteHandler) key.attachment();
    if (key.isReadable()) rwH.handleRead(key);
    if (key.isWritable()) rwH.handleWrite(key);
}
```

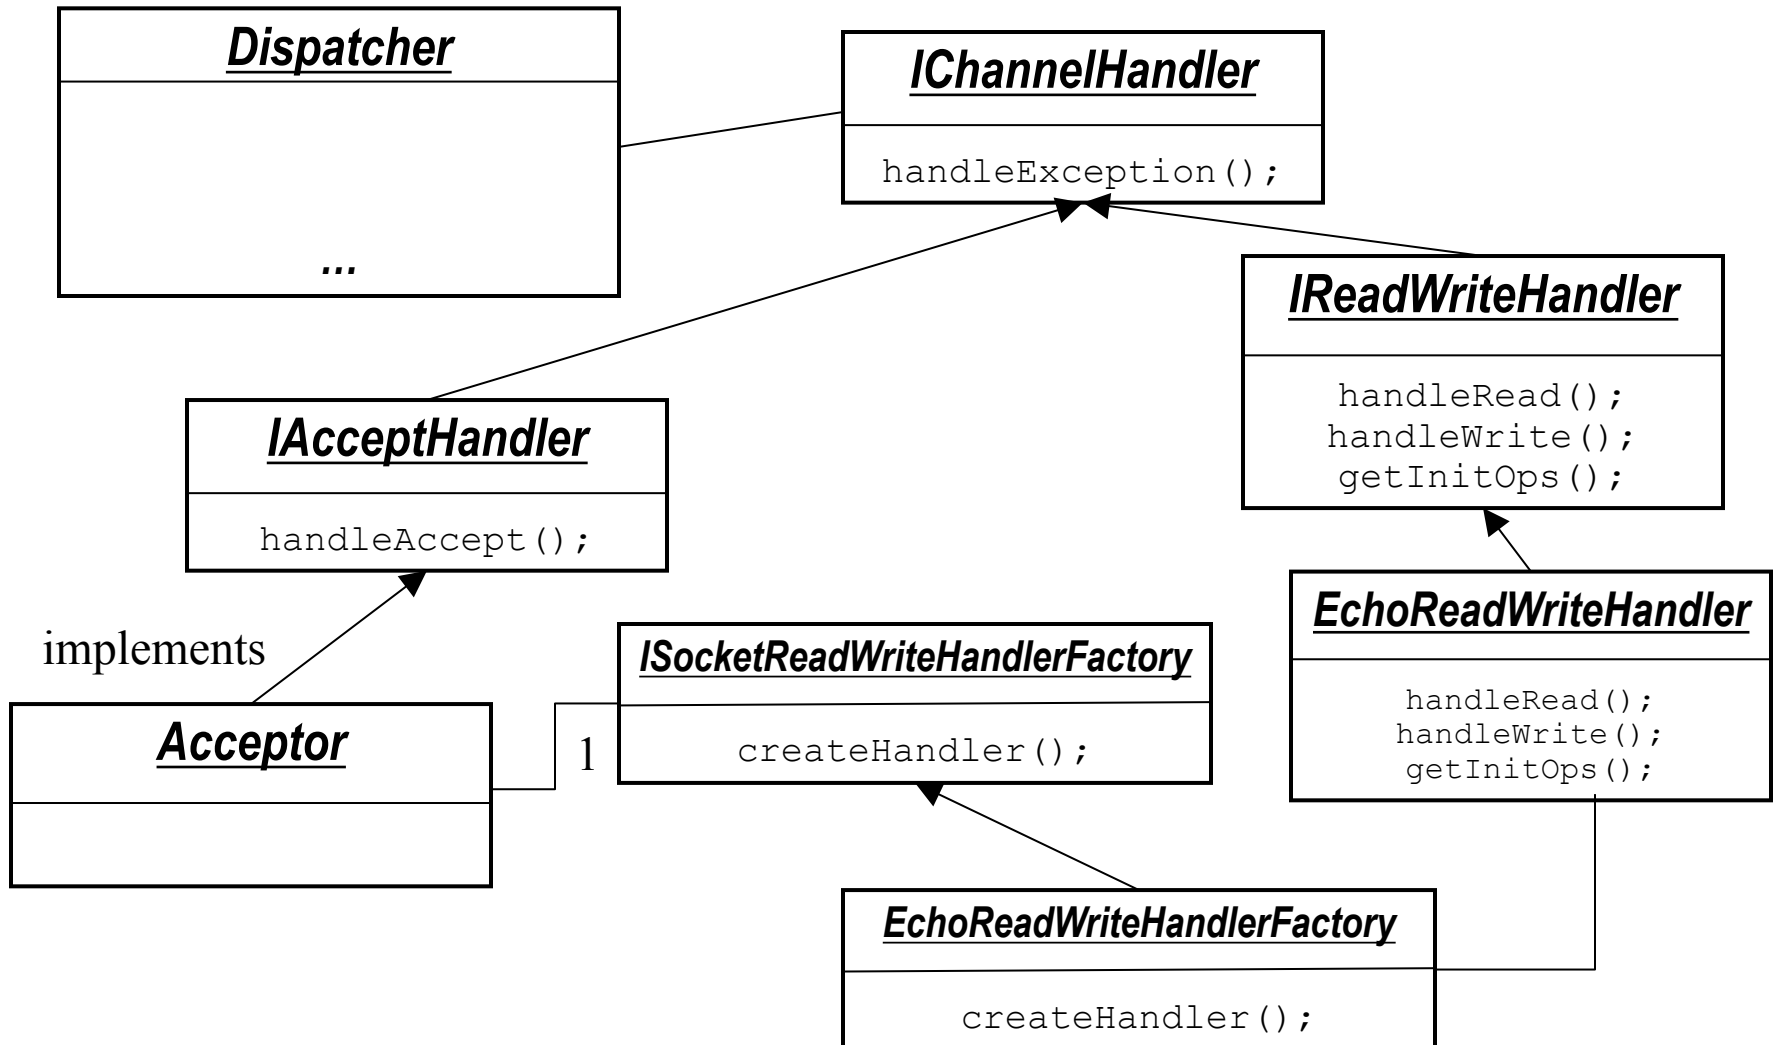
Handler Design: Acceptor

- ❑ What should an accept handler object know?
 - ServerSocketChannel (so that it can call accept)
 - Can be derived from SelectionKey in the call back
 - Selector (so that it can register new connections)
 - Can be derived from SelectionKey in the call back
 - What ReadWrite object to create (different protocols may use different ones)?
 - Pass a Factory object: SocketReadWriteHandlerFactory

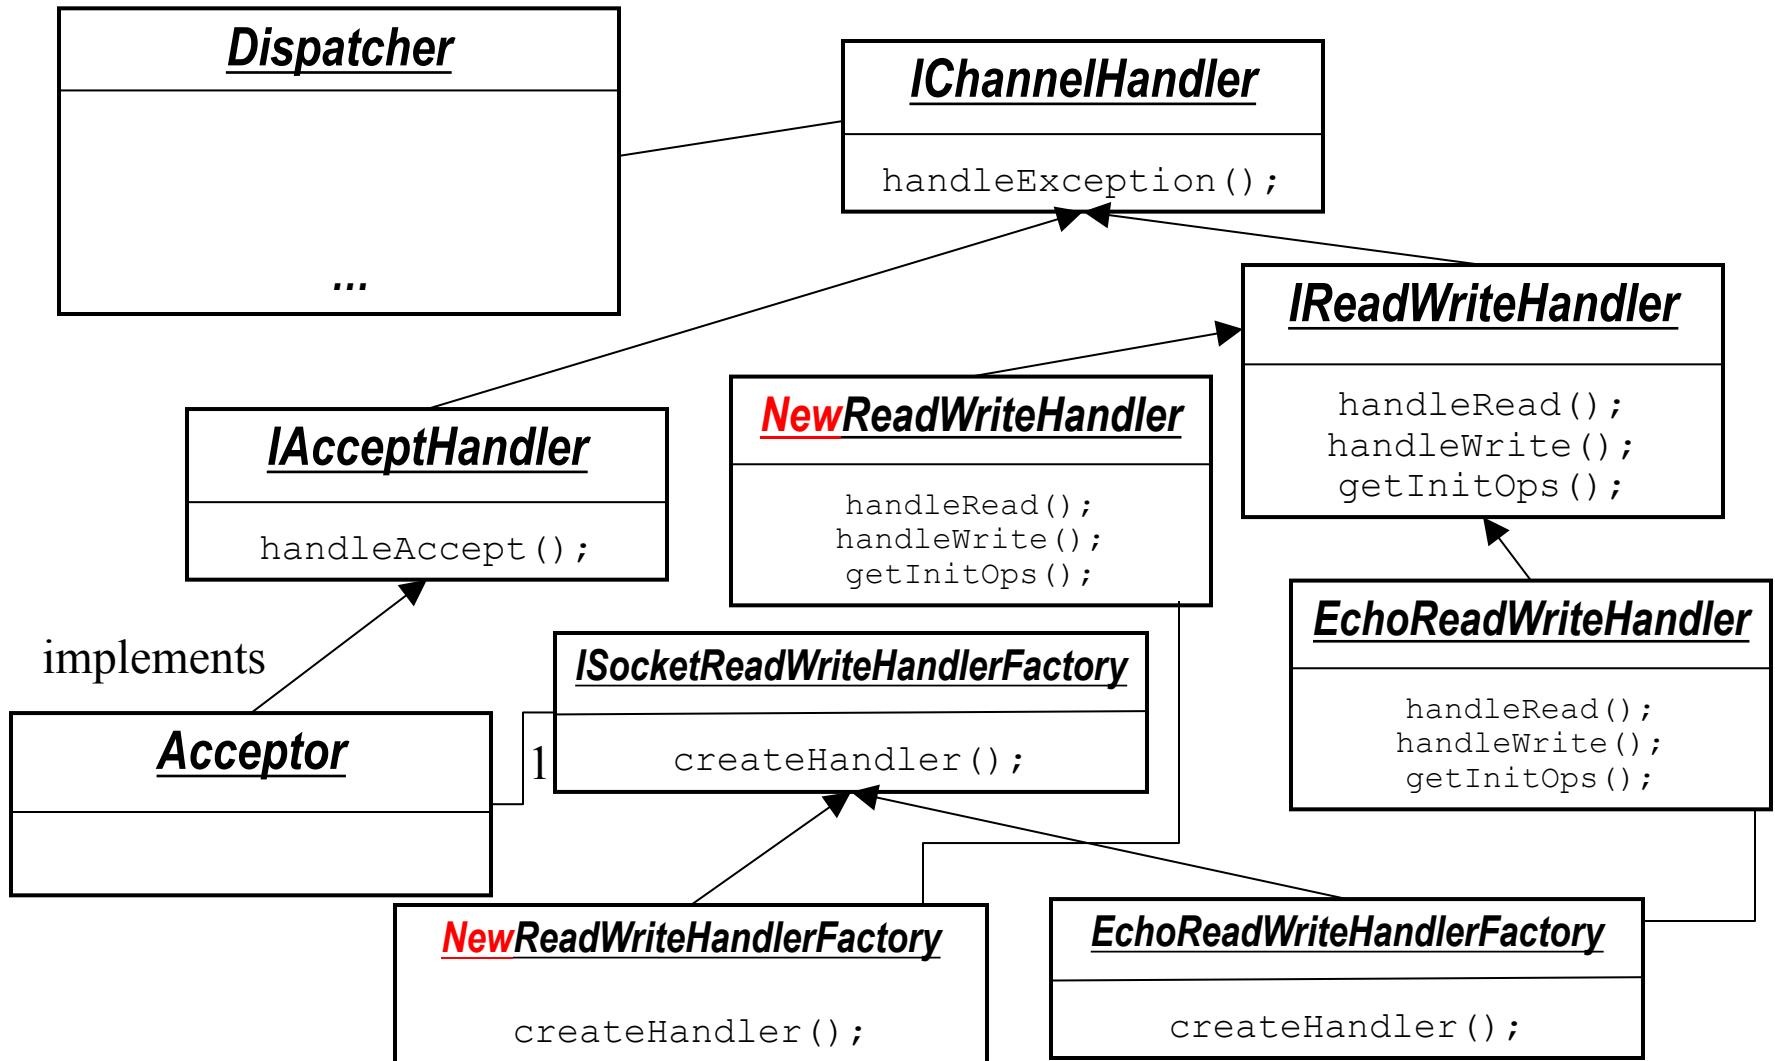
Handler Design: ReadWriteHandler

- ❑ What should a ReadWrite handler object know?
 - SocketChannel (so that it can read/write data)
 - Can be derived from SelectionKey in the call back
 - Selector (so that it can change state)
 - Can be derived from SelectionKey in the call back

Class Diagram of SimpleNAIO



Class Diagram of SimpleNAIO



SimpleNAIO

□ See `SelectEchoServer/v2/*.java`

Design Exercise

❑ In our current implementation (Server.java)

1. Create dispatcher
2. Create server socket channel
3. Register server socket channel to dispatcher
4. Start dispatcher thread

Can we simply switch 3 and 4?

Design Exercise to Understand Server Structure

- ❑ A production network server often closes a connection if it does not receive a complete request in TIMEOUT
 - ❑ One way to implement time out is that
 - the read handler registers a timeout event with a timeout watcher thread with a call back
 - the watcher thread invokes the call back upon TIMEOUT
 - the callback closes the connection
- Any problem?

Extending Dispatcher Interface

- ❑ Interacting from another thread to the dispatcher thread can be tricky
- ❑ Typical solution: async command queue

```
while (true) {  
    - process async. command queue  
    - ready events = select (or selectNow(), or  
      select(int timeout)) to check for ready events  
      from the registered interest events of  
      SelectableChannels  
  
    - foreach ready event  
      call handler  
}
```

Execute Commands by Dispatcher

```
public void invokeLater(Runnable run) {  
    synchronized (pendingInvocations) {  
        pendingInvocations.add(run);  
    }  
    selector.wakeup();  
}
```

Design Exercise to Understand Server Structure

- ❑ What if another thread wants to wait until a command is finished by the dispatcher thread?
 - AKA: How to block another thread until its command is executed by the dispatcher thread


```

public void invokeAndWait(final Runnable task)
    throws InterruptedException
{
    if (Thread.currentThread() == selectorThread) {
        // We are in the selector's thread. No need to schedule
        // execution
        task.run();
    } else {
        // Used to deliver the notification that the task is executed
        final Object latch = new Object();
        synchronized (latch) {
            // Uses the invokeLater method with a newly created task
            this.invokeLater(new Runnable() {
                public void run() {
                    task.run();
                    // Notifies
                    synchronized(latch) { latch.notify(); }
                }
            });
            // Wait for the task to complete.
            latch.wait();
        }
        // Ok, we are done, the task was executed. Proceed.
    }
}

```


Backup Slides

Another view

- Events obscure control flow
 - For programmers *and* tools

<i>Threads</i>	<i>Events</i>
<pre>thread_main(int sock) { struct session s; accept_conn(sock, &s); read_request(&s); pin_cache(&s); write_response(&s); unpin(&s); } pin_cache(struct session *s) { pin(&s); if(!in_cache(&s)) read_file(&s); }</pre>	<pre>AcceptHandler(event e) { struct session *s = new_session(e); RequestHandler.enqueue(s); } RequestHandler(struct session *s) { ...; CacheHandler.enqueue(s); } CacheHandler(struct session *s) { pin(s); if(!in_cache(s)) ReadFileHandler.enqueue(s); else ResponseHandler.enqueue(s); } ... ExitHandler(struct session *s) { ...; unpin(&s); free_session(s); }</pre>

Web Server



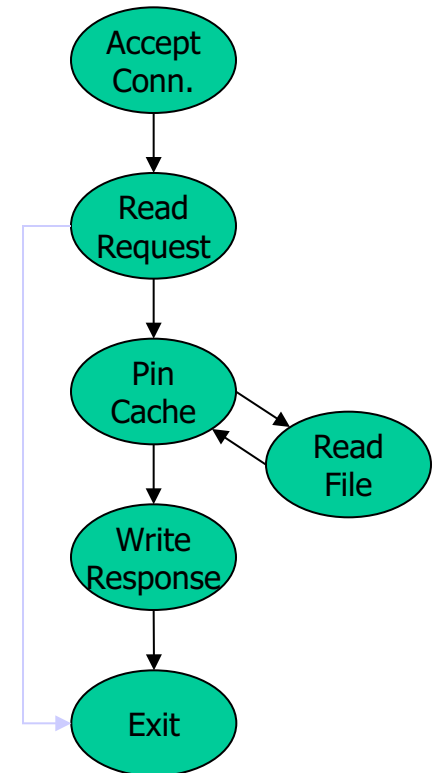
[von Behren]

State Management

- ❑ Events require manual state management
- ❑ Hard to know when to free
 - Use GC or risk bugs

<i>Threads</i>	<i>Events</i>
<pre>thread_main(int sock) { struct session s; accept_conn(sock, &s); if(!read_request(&s)) return; pin_cache(&s); write_response(&s); unpin(&s); }</pre> <pre>pin_cache(struct session *s) { pin(&s); if(!in_cache(&s)) read_file(&s); }</pre>	<pre>CacheHandler(struct session *s) { pin(s); if(!in_cache(s)) ReadFileHandler.enqueue(s); else ResponseHandler.enqueue(s); } RequestHandler(struct session *s) { ...; if(error) return; CacheHandler.enqueue(s); } ... ExitHandler(struct session *s) { ...; unpin(&s); free_session(s); } AcceptHandler(event e) { struct session *s = new_session(e); RequestHandler.enqueue(s); }</pre>

Web Server



[von Behren]