

---

Network Transport Layer:  
Transport Reliability:  
Sliding Windows; Connection Management

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

11/6/2018

# Admin.: PS4

## □ Part 1

- Discussion checkpoint: Nov. 11; code checkpoint Nov. 13

## □ Part 2

- Discussion checkpoint: Nov. 16; all due Nov. 27

proj-sol:

129 400 3045 FishThread.java  
**388** 1457 12873 Node.java  
51 167 1145 PingRequest.java  
**83** 250 2106 SimpleTCPSockSpace.java  
**181** 605 5248 TCPManager.java  
**889** 3088 26381 TCPSock.java  
**60** 149 1316 TCPSockID.java  
123 382 3866 TransferClient.java  
147 500 5059 TransferServer.java

**2051** 6998 61039 total

proj:

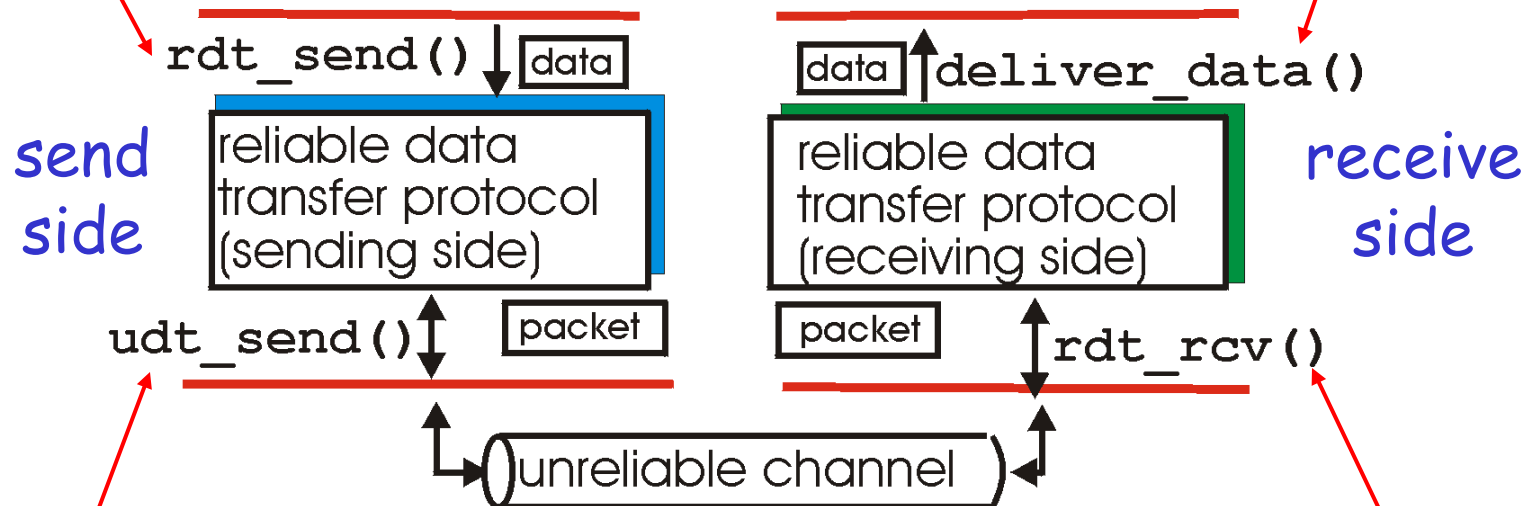
129 400 3045 FishThread.java  
**341** 1301 11313 Node.java  
51 167 1145 PingRequest.java  
**50** 128 909 TCPManager.java  
132 460 3146 TCPSock.java  
  
123 382 3866 TransferClient.java  
147 500 5059 TransferServer.java

973 3338 28483 total

# Recap: Reliable Data Transfer Context

**rdt\_send()** : called from above,  
(e.g., by app.)

**deliver\_data()** : called by  
rdt to deliver data to upper



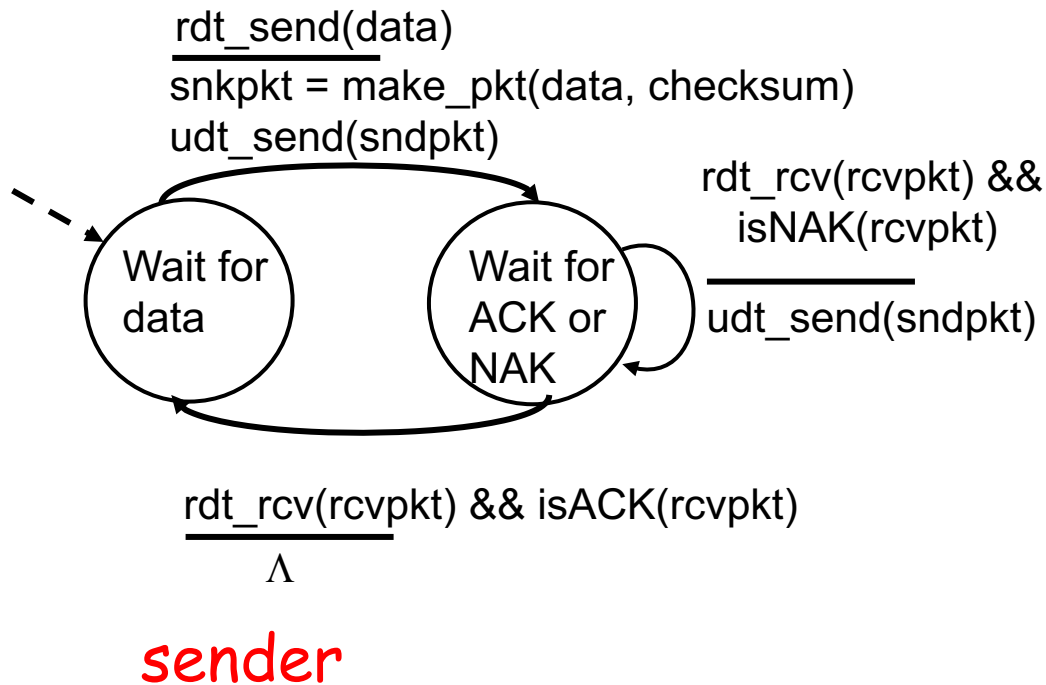
**udt\_send()** : called by rdt,  
to transfer packet over  
unreliable channel to receiver

**rdt\_rcv()** : called from below;  
when packet arrives on rcv-side of  
channel


# Recap: Potential Channel Errors

- ❑ Factors to pay attention when designing rdt
  - Types of channel errors: Characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt).
    - bit errors
    - loss (drop) of packets
    - reordering or duplication
  - Not only protocol but also analysis techniques

# Recap: rdt2.0: Reliability allowing only Data Msg Corruption

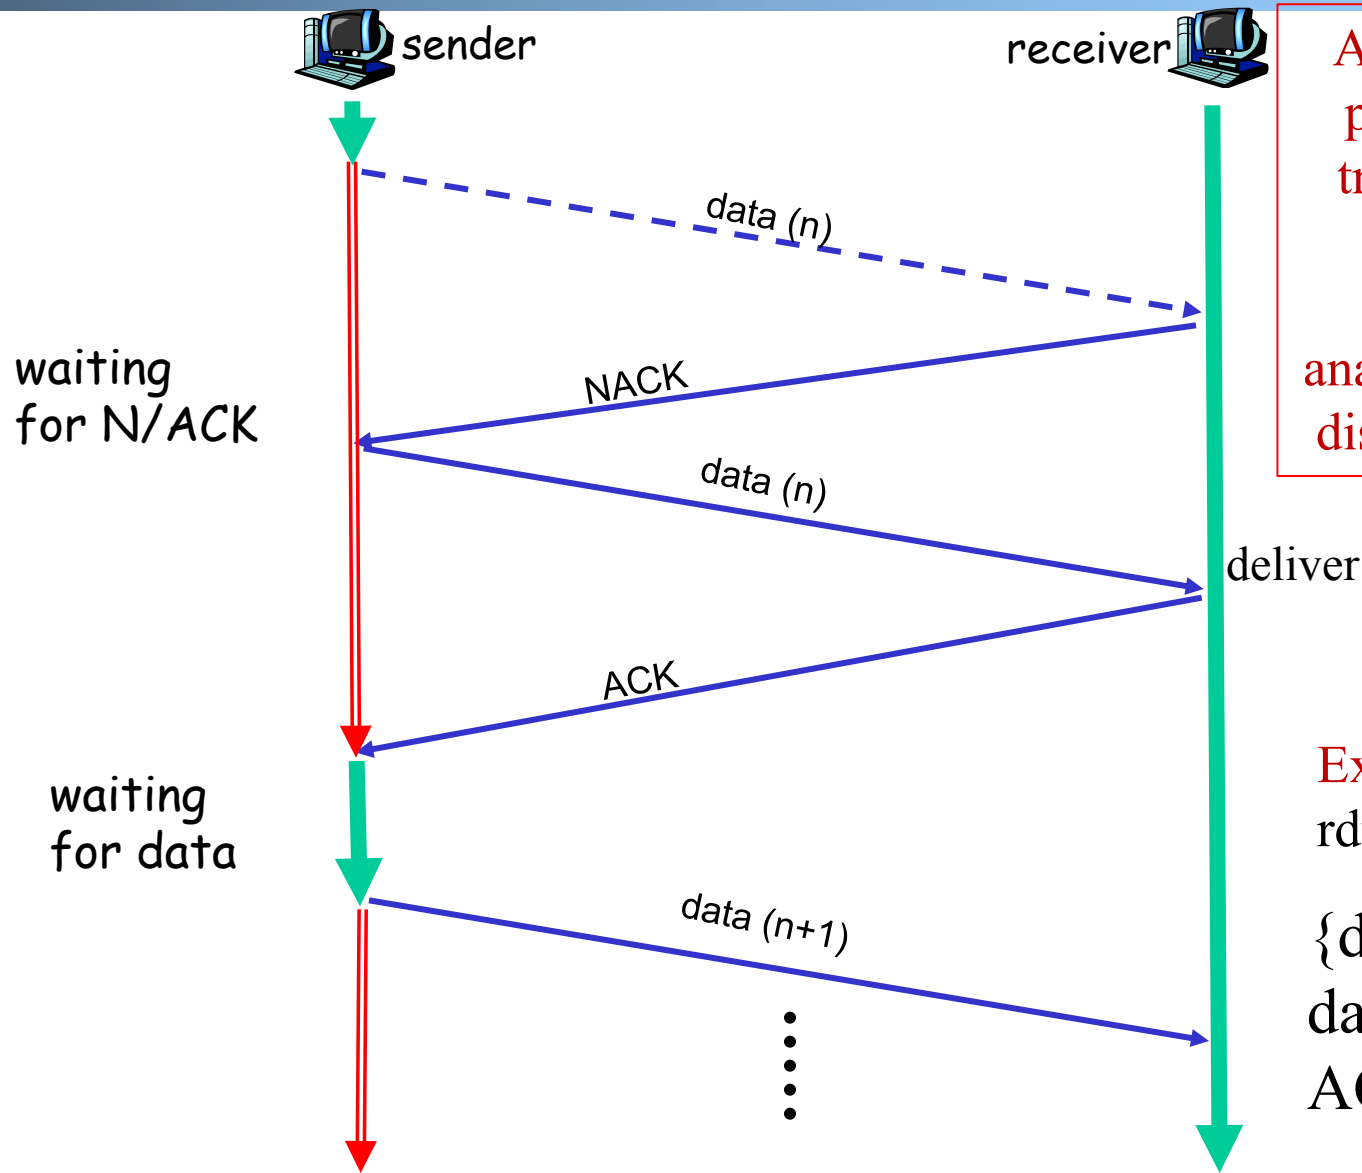


receiver



# Recap: Rdt2.0 Analysis

$\text{data}^\wedge: \langle S \text{ data} \rangle \langle R \text{ data} \rangle^*$   
 $\text{data}: \langle S \text{ data} \rangle \langle R \text{ data} \rangle$

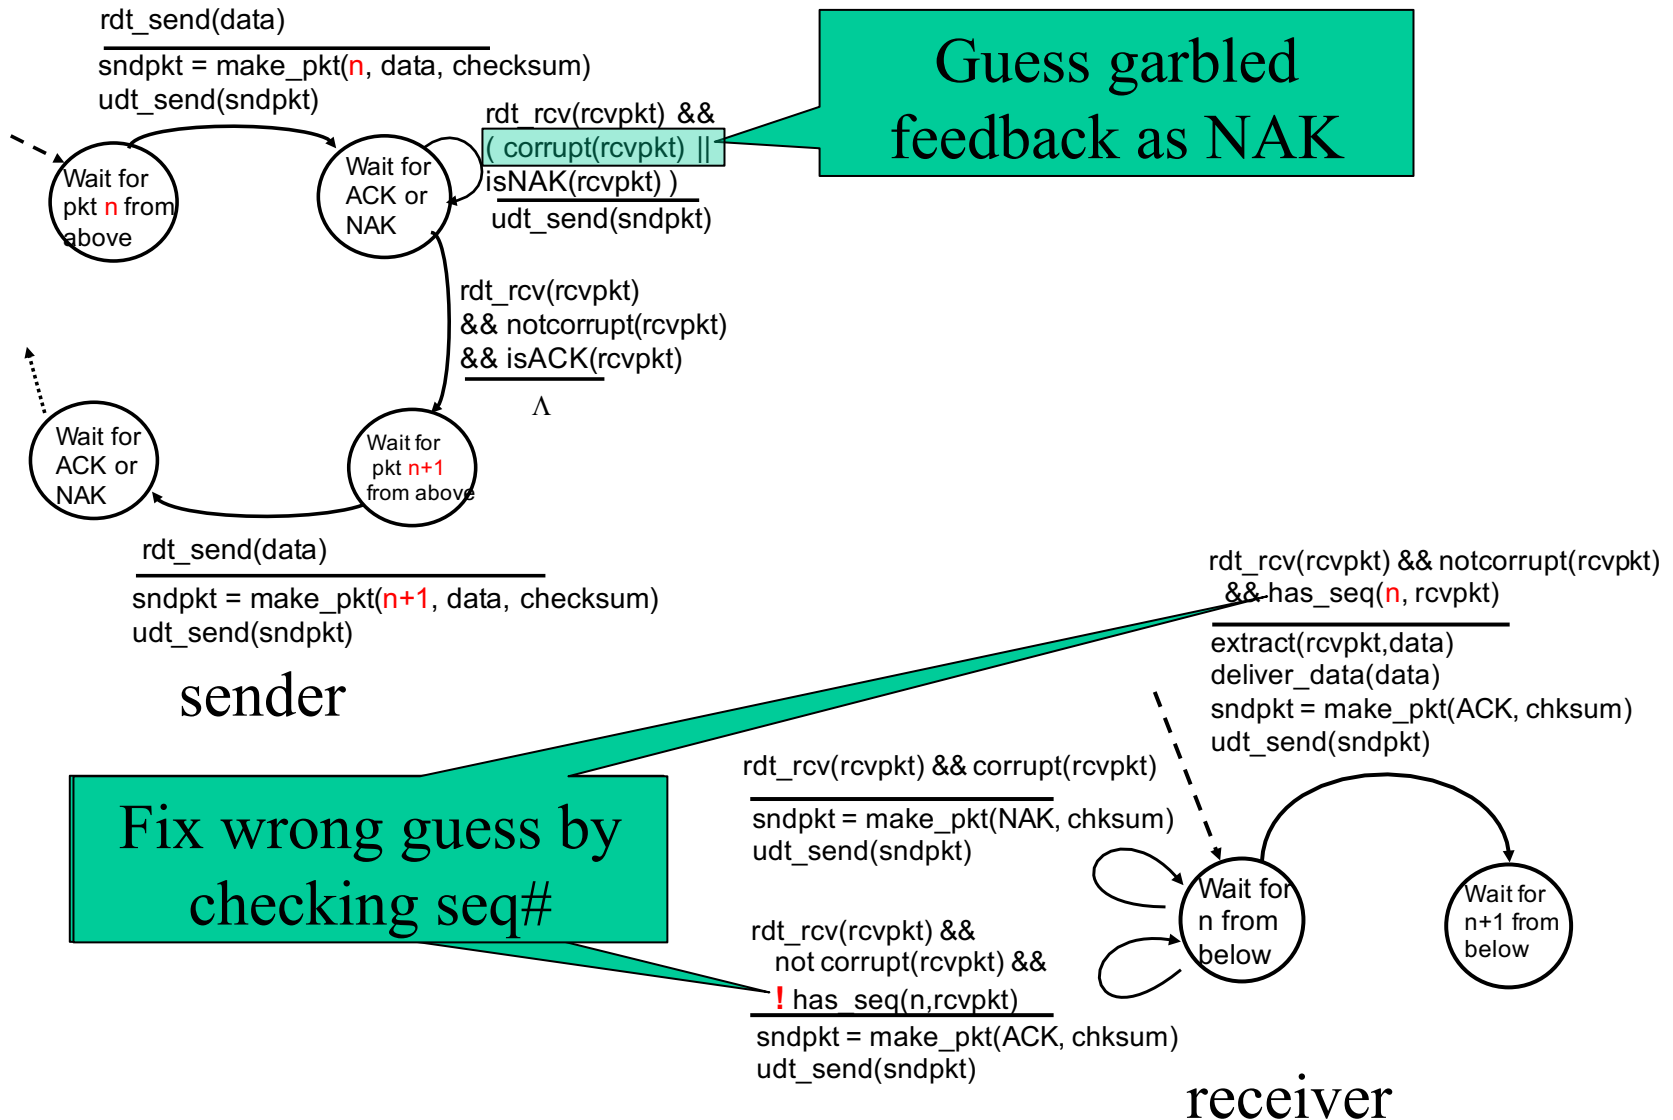


Analyzing set of all possible execution traces is a common technique to understand and analyze many types of distributed protocols.

Execution traces of rdt2.0:

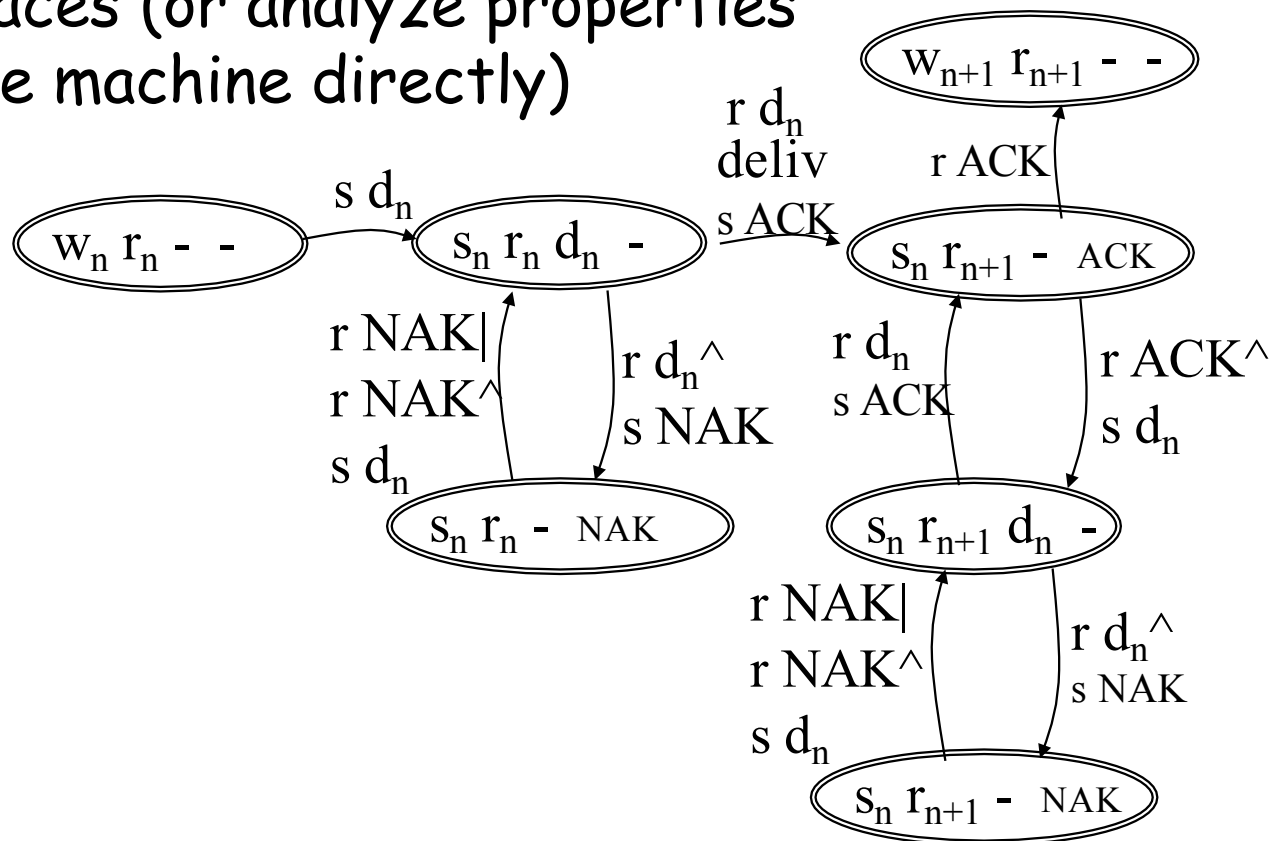
$\{\text{data}^\wedge \text{ NACK} \}^*$   
data deliver  
ACK

# Recap: rdt2.1b: Reliability allowing Data/**Control** Msg Corruption



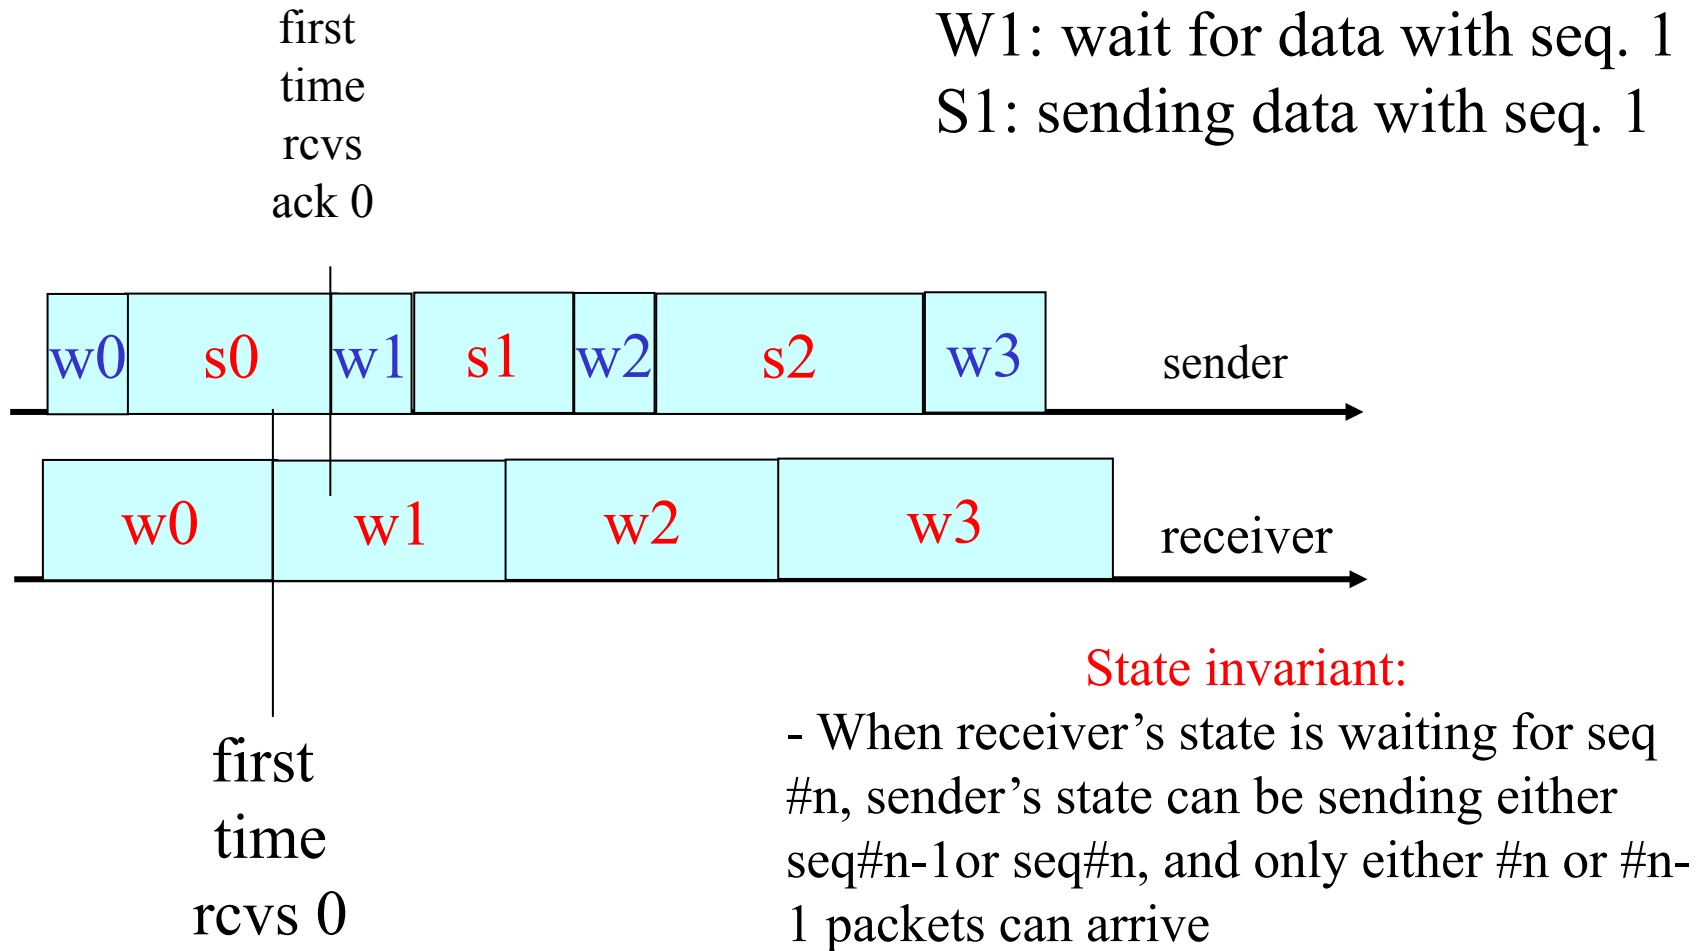
# Recap: Protocol Analysis using (Generic) Execution Traces Technique

- A systematic approach to enumerating execution traces is to compute **joint sender/receiver/channels state machine**, and then convert the state machine to traces (or analyze properties on the machine directly)



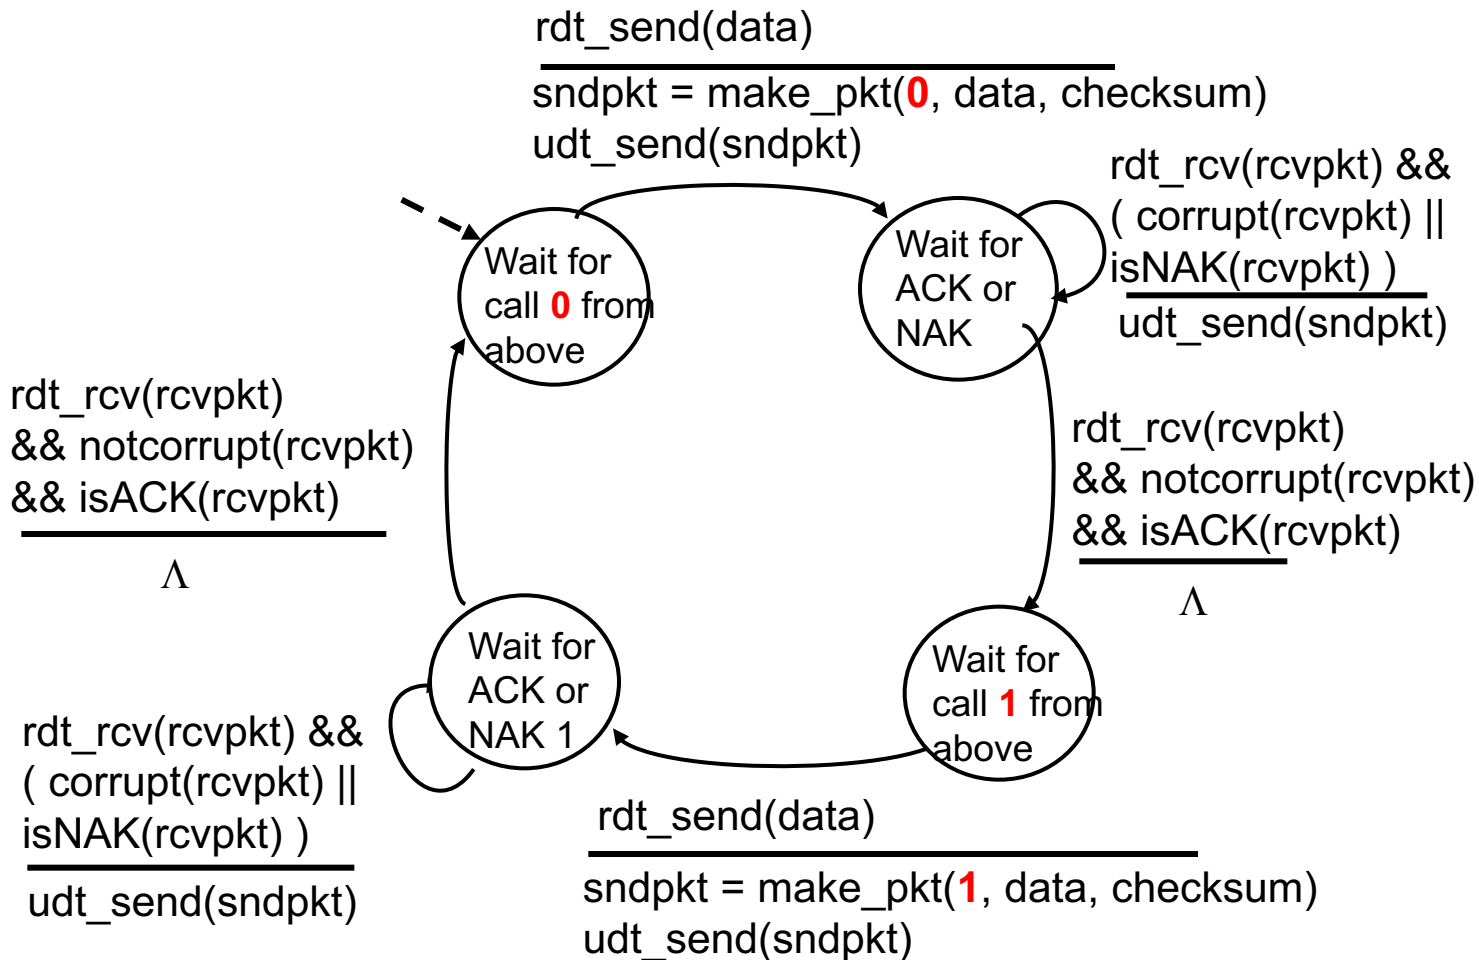


# Recap: Protocol Analysis using State Invariants



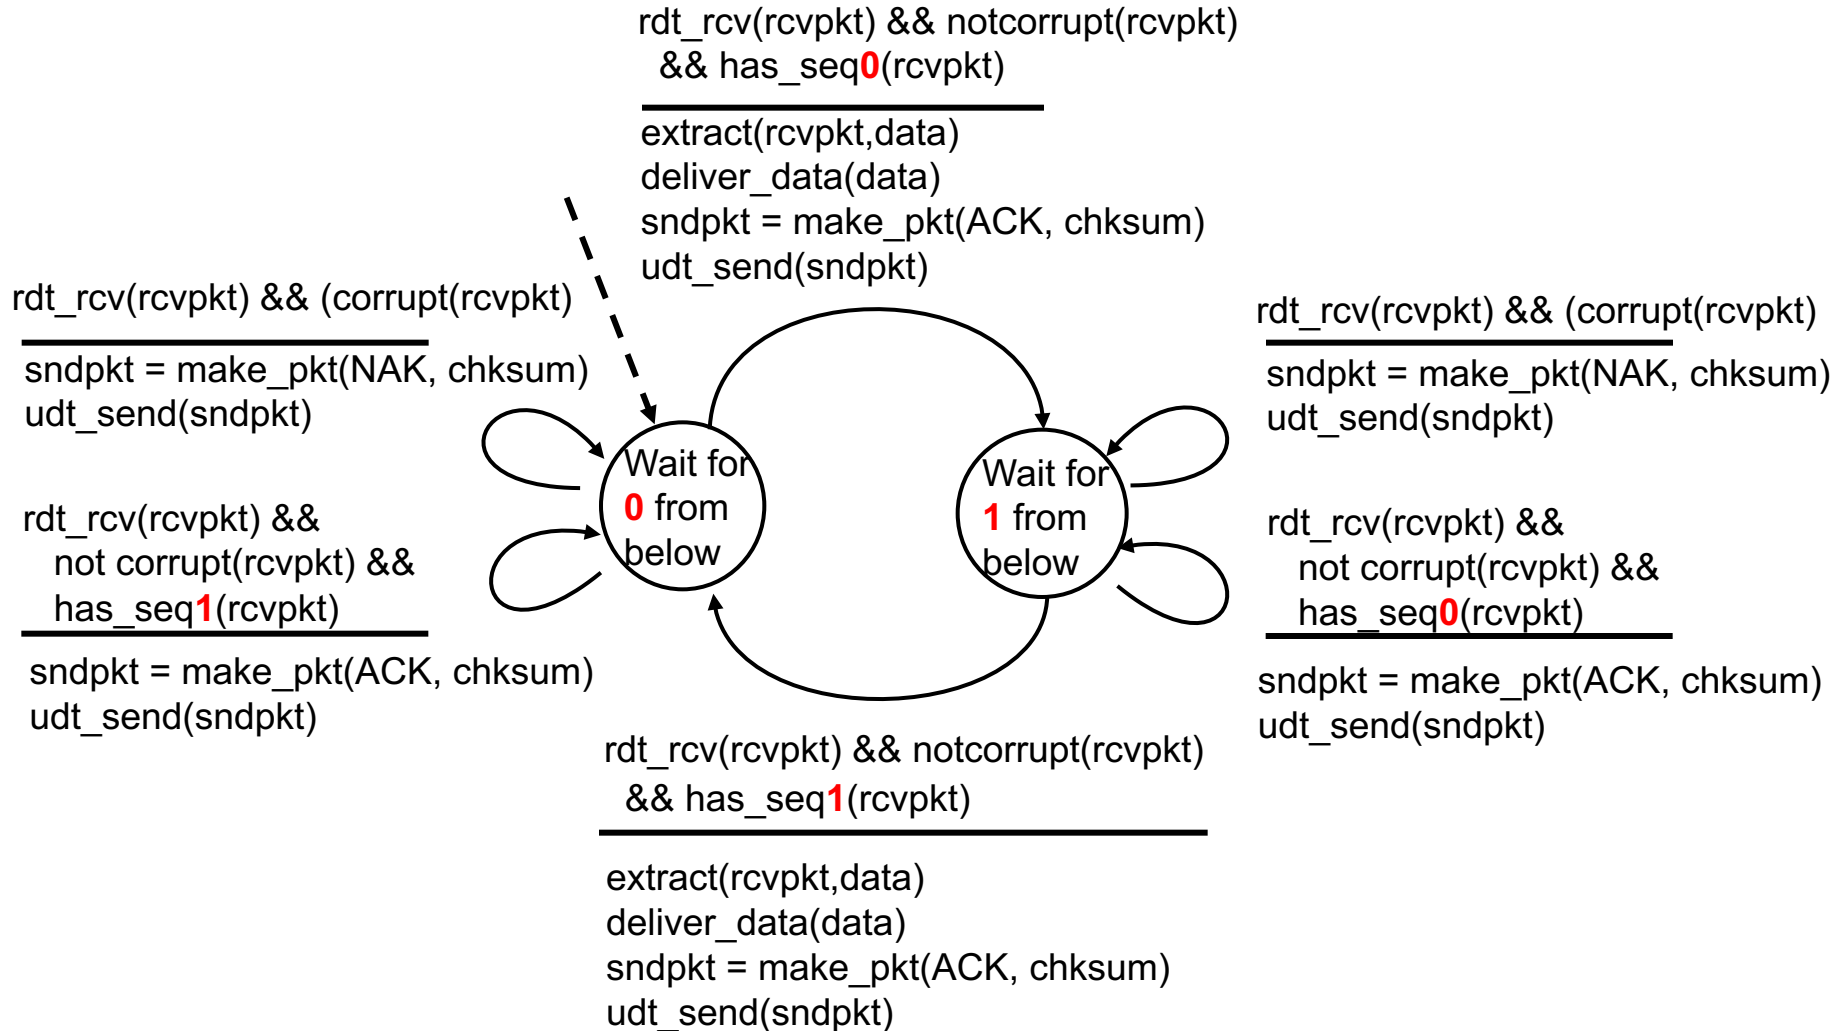
**Implication: One bit (the last bit) is enough to distinguish the two states**

## rdt2.1c: Sender, Handles Garbled ACK/NAKs: Using 1 bit (Alternating-Bit Protocol)



# rdt2.1c: Receiver, Handles Garbled

## ACK/NAKs: Using 1 bit



# rdt2.1c: Summary

---

## Sender:

- state must “remember” whether “current” pkt has 0 or 1 seq. #

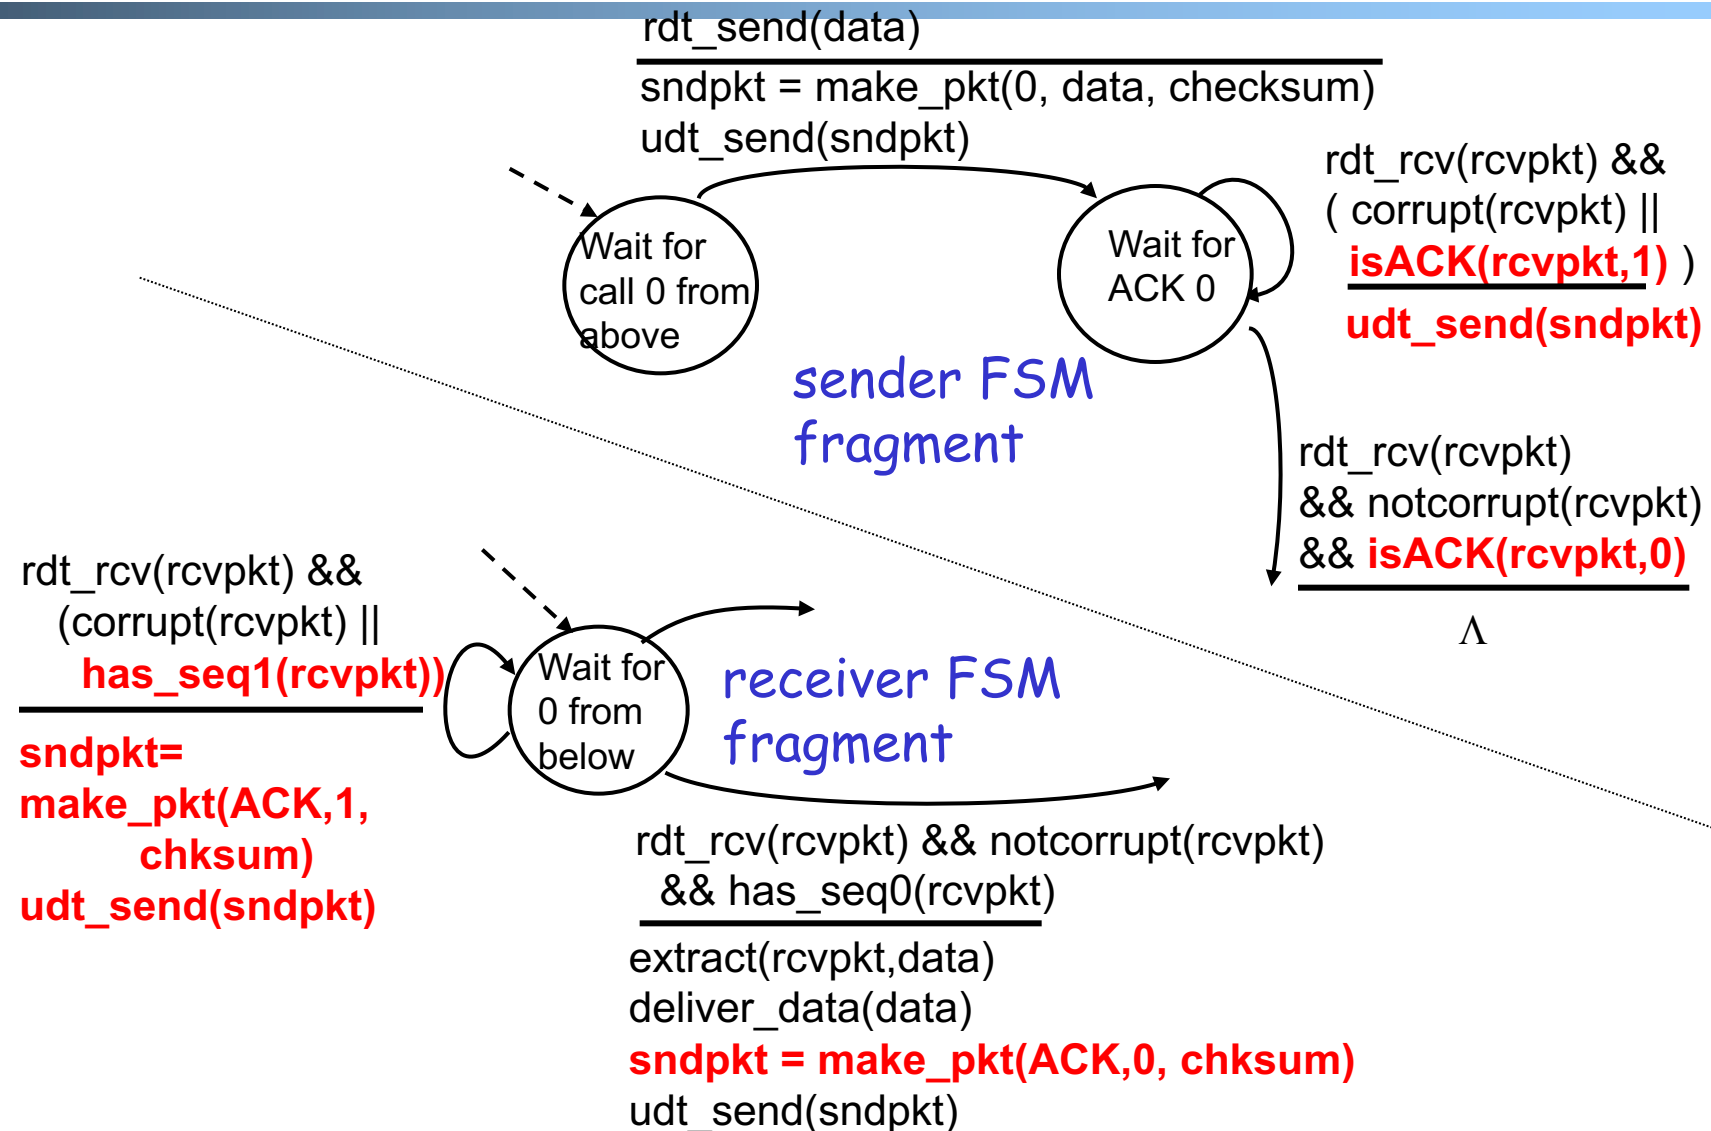
## Receiver:

- must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #

## rdt2.2: a NAK-free protocol

- ❑ Same functionality as rdt2.1c, using ACKs only
- ❑ Instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must *explicitly* include seq # of pkt being ACKed
- ❑ Duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

# rdt2.2: Sender, Receiver Fragments



# Outline

---

- ❑ Admin and review
- Reliable data transfer
  - perfect channel
  - channel with bit errors
  - channel with bit errors and losses

# rdt3.0: Channels with Errors and Loss

## New assumption:

underlying channel can also lose packets (data or ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help, but not enough

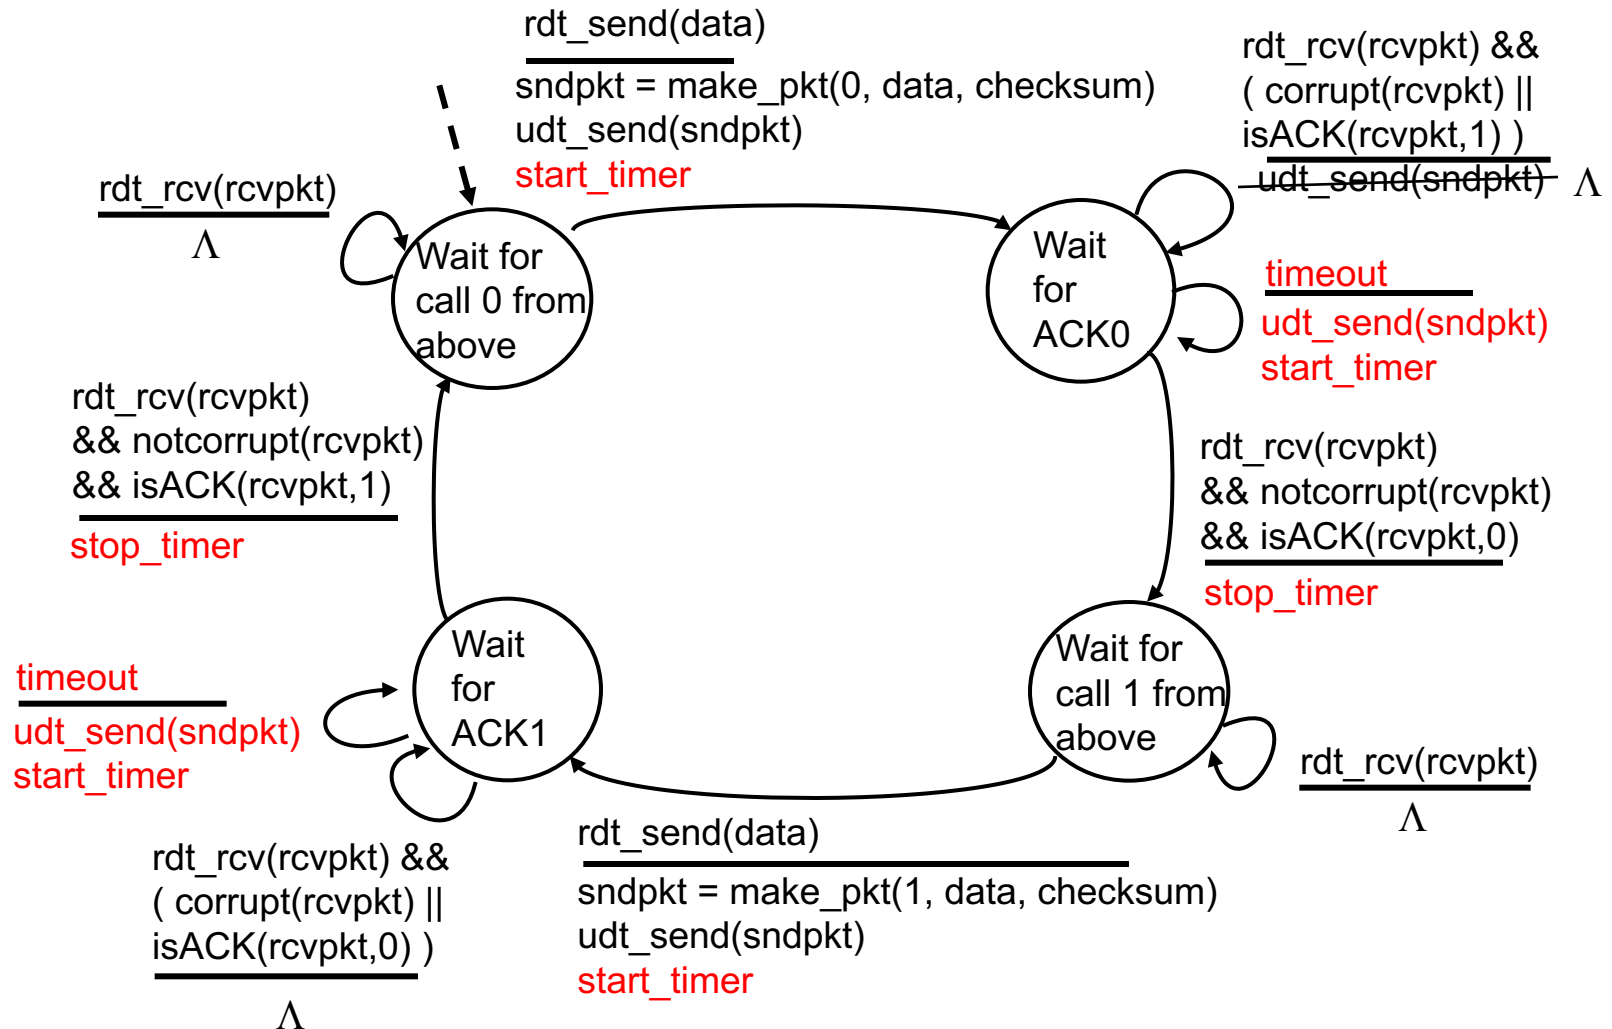
Q: What can rdt2.2 go wrong under losses?

Approach: sender waits “reasonable” amount of time for ACK

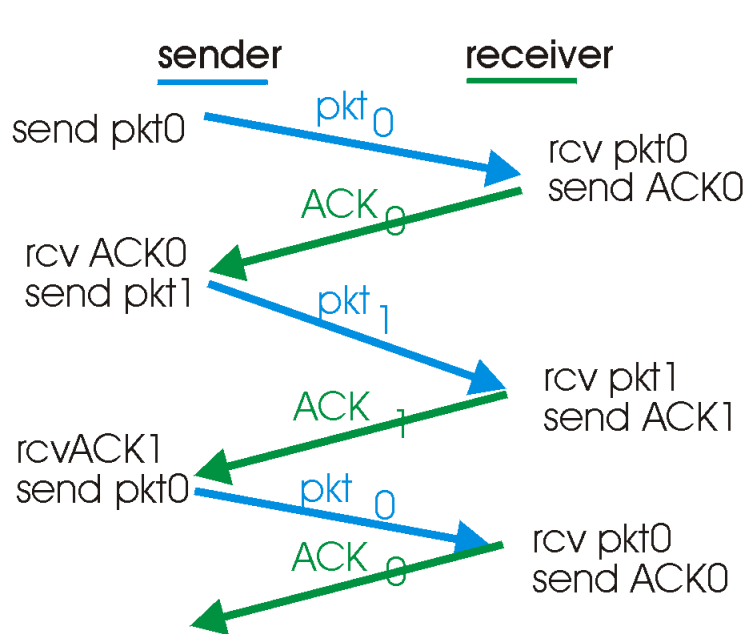
- requires countdown timer
- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but use of seq. #'s already handles this
  - receiver must specify seq # of pkt being ACKed



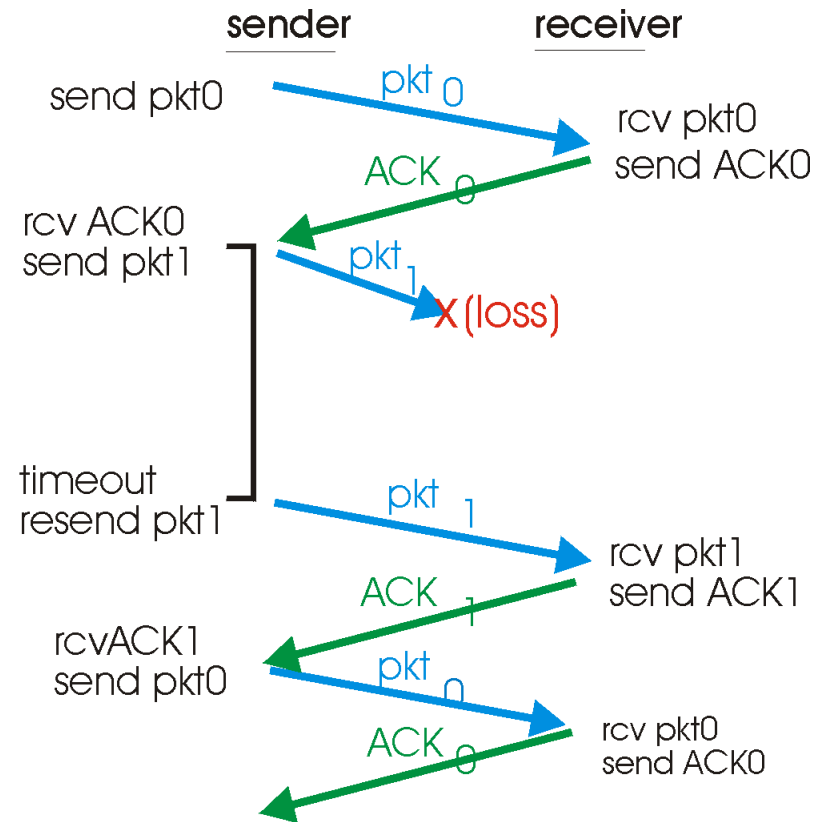
# rdt3.0 Sender



# rdt3.0 in Action




(a) operation with no loss

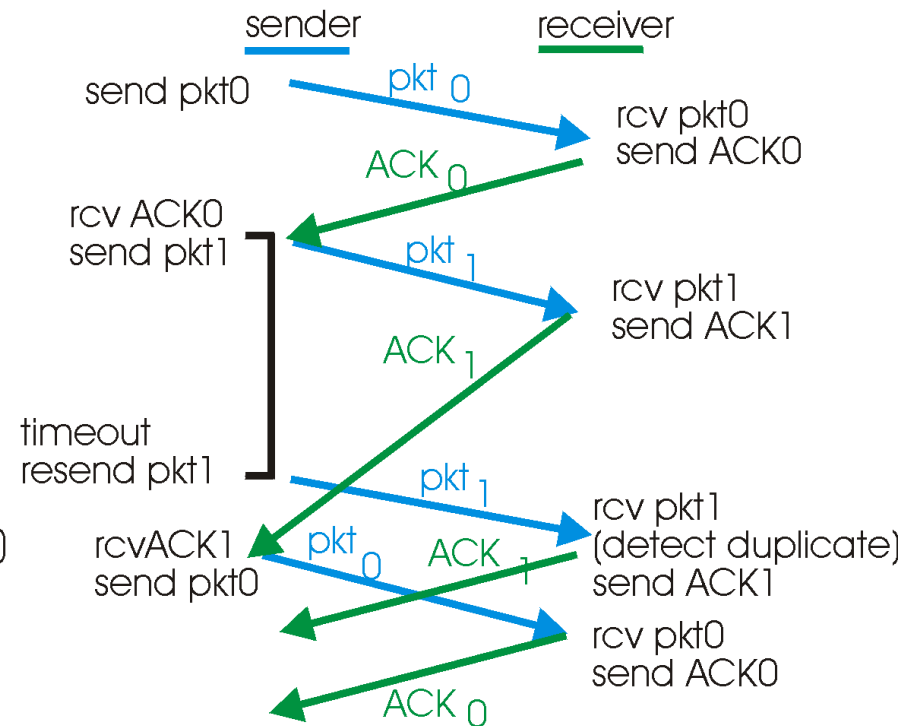


(b) lost packet

# rdt3.0 in Action



(c) lost ACK

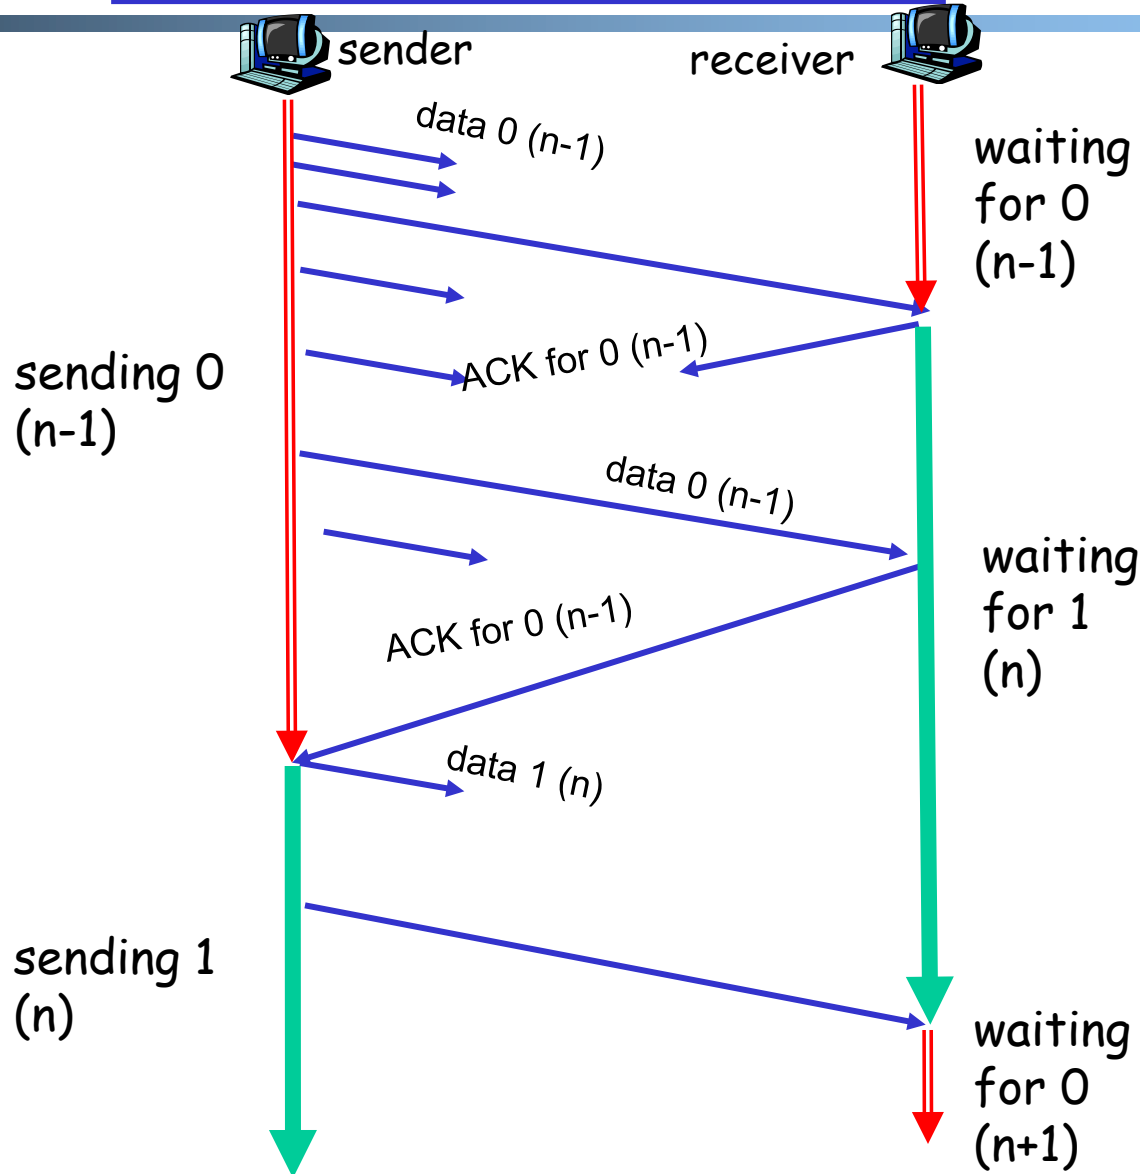


(d) premature timeout

Question to think about: How to determine a good timeout value?

Home exercises: (1) What are execution traces of rdt3.0? What are some state invariants of rdt3.0? (2) What are some state invariants?

# rdt3.0: Protocol Analysis using State Invariants




## State consistency:

When receiver's state is waiting  $n$ , the state of the sender is either sending for  $n-1$  or sending for  $n$

When sender's state is sending for  $n$ , receiver's state is waiting for  $n$  or  $n + 1$

# rdt3.0: Stop-and-Wait Performance



What is  $U_{\text{sender}}$ : **utilization** – fraction of time link busy sending?

Assume: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet

# Performance of rdt3.0

- ❑ rdt3.0 works, but performance stinks
- ❑ Example: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet:

$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^{**9} \text{ b/sec}} = 8 \text{ microsec}$$


$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 1KB pkt every 30 msec -> 33kB/sec thruput over 1 Gbps link
- network protocol limits use of physical resources !

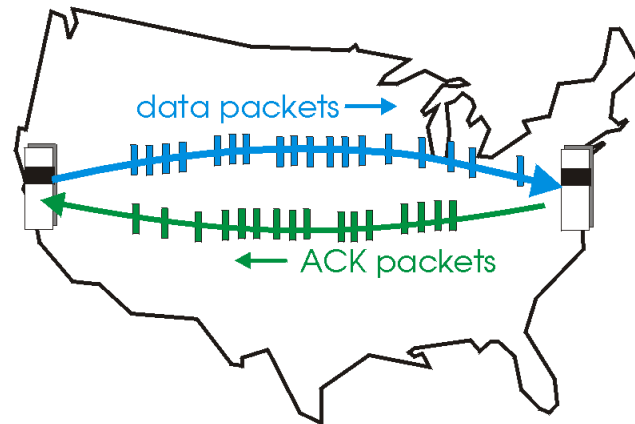
# Sliding Window Protocols: Pipelining

**Pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation

# Pipelining: Increased Utilization



$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$


Question: a rule-of-thumb window size?



# Realizing Sliding Window: Go-Back-n


## Sender:

- ❑ k-bit seq # in pkt header
- ❑ “window” of up to  $W$ , consecutive unack’ed pkts allowed

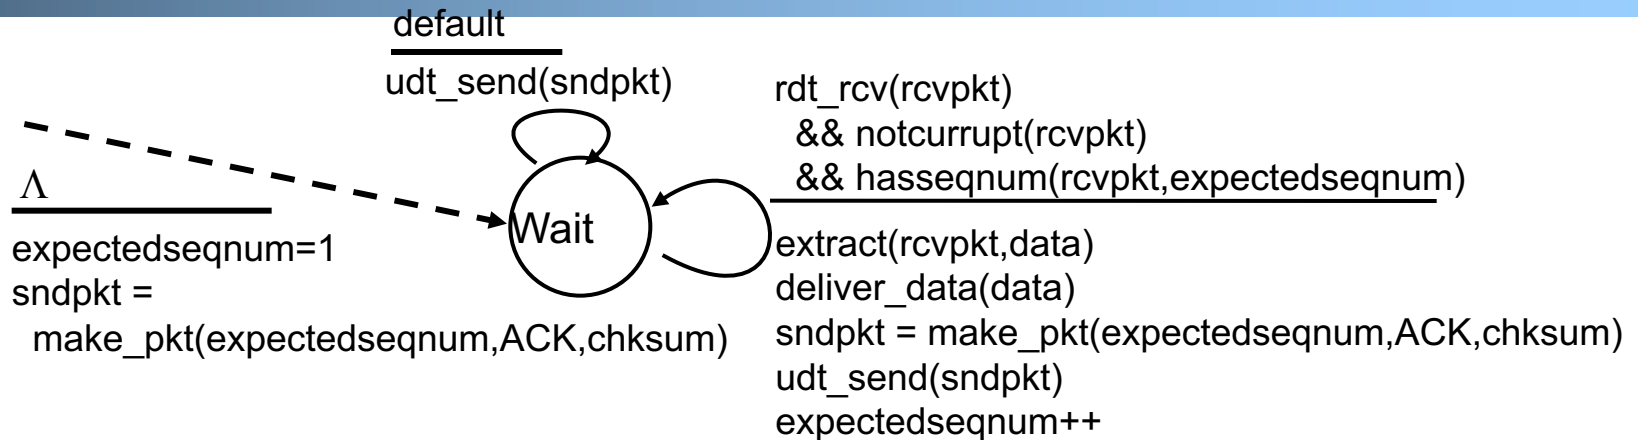


- ❑ **ACK(n): ACKs all pkts up to, including seq #  $n$  - “cumulative ACK”**
  - note: ACK(n) could mean two things: I have received **upto and include**  $n$ , or I am waiting for  $n$
- ❑ timer for the packet at base
- ❑ **timeout(n):** retransmit pkt  $n$  and all higher seq # pkts in window

# GBN: Sender FSM



# GBN: Receiver FSM



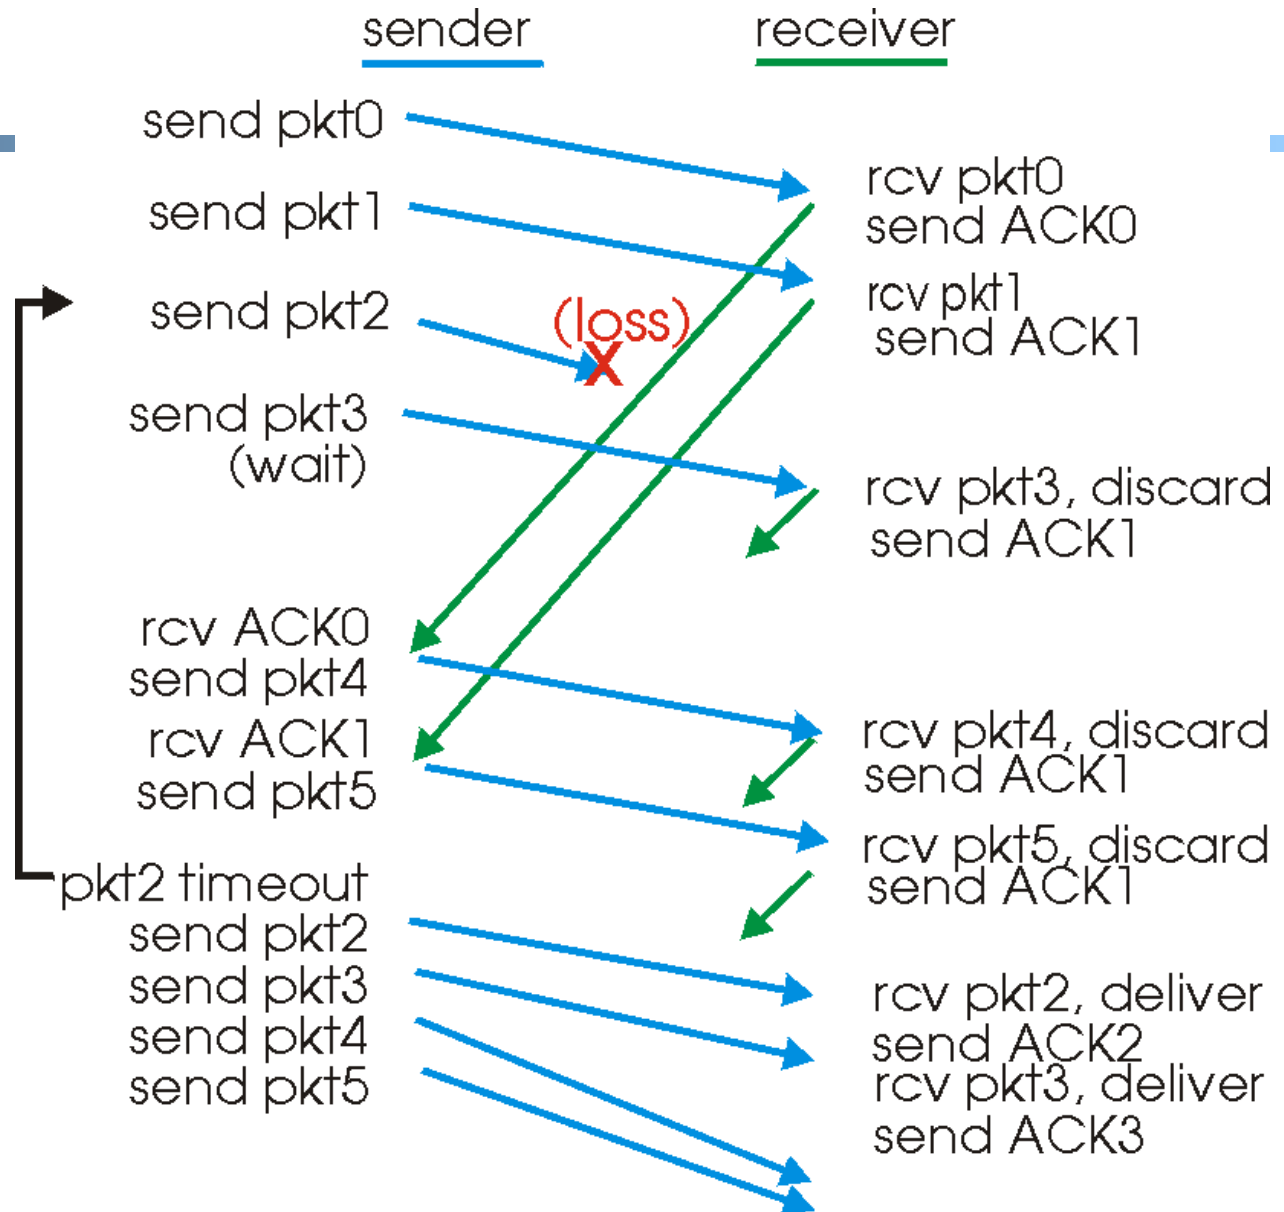
Only state: **expectedseqnum**

❑ out-of-order pkt:

- discard (don't buffer) -> **no receiver buffering!**
- re-ACK pkt with highest in-order seq #
- may generate duplicate ACKs

# GBN in Action

window  
size = 4



# Analysis: Efficiency of Go-Back-n

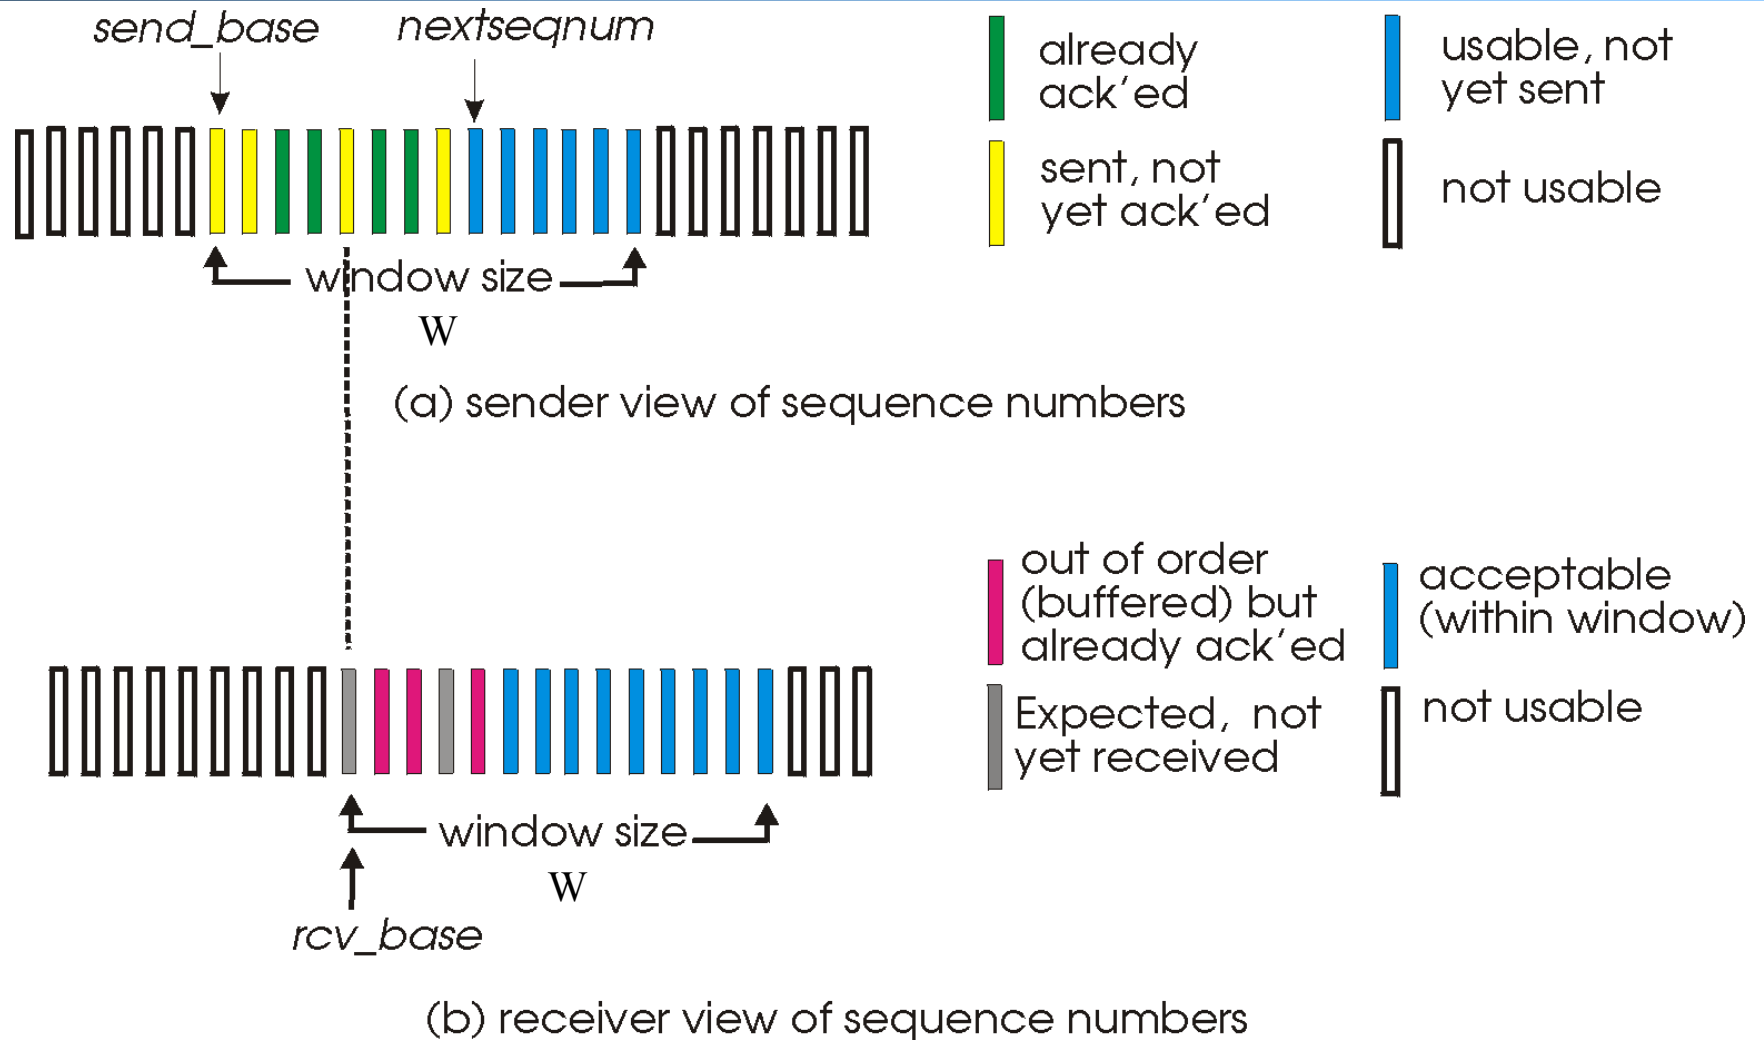
---

- ❑ Assume window size  $W$
- ❑ Assume each packet is lost with probability  $p$
- ❑ On average, how many packets do we send for each data packet received?

# Selective Repeat

- ❑ Sender window
  - Window size  $W$ :  $W$  consecutive unACKed seq #'s
- ❑ Receiver *individually* acknowledges correctly received pkts
  - buffers out-of-order pkts, for eventual in-order delivery to upper layer
  - ACK(n) means received packet with seq# n only
  - question: buffer size at receiver?
- ❑ Sender only resends pkts for which ACK not received
  - sender timer for each unACKed pkt

# Selective Repeat: Sender, Receiver Windows



# Selective Repeat

## sender

data from above :

- unACKed packets is less than window size  $W$ , send; otherwise block app.

timeout(n):

- resend pkt  $n$ , restart timer

ACK(n) in [sendbase, sendbase+ $W$ -1]:

- mark pkt  $n$  as received
- update sendbase to the first packet unACKed

## receiver

pkt  $n$  in [rcvbase, rcvbase+ $W$ -1]

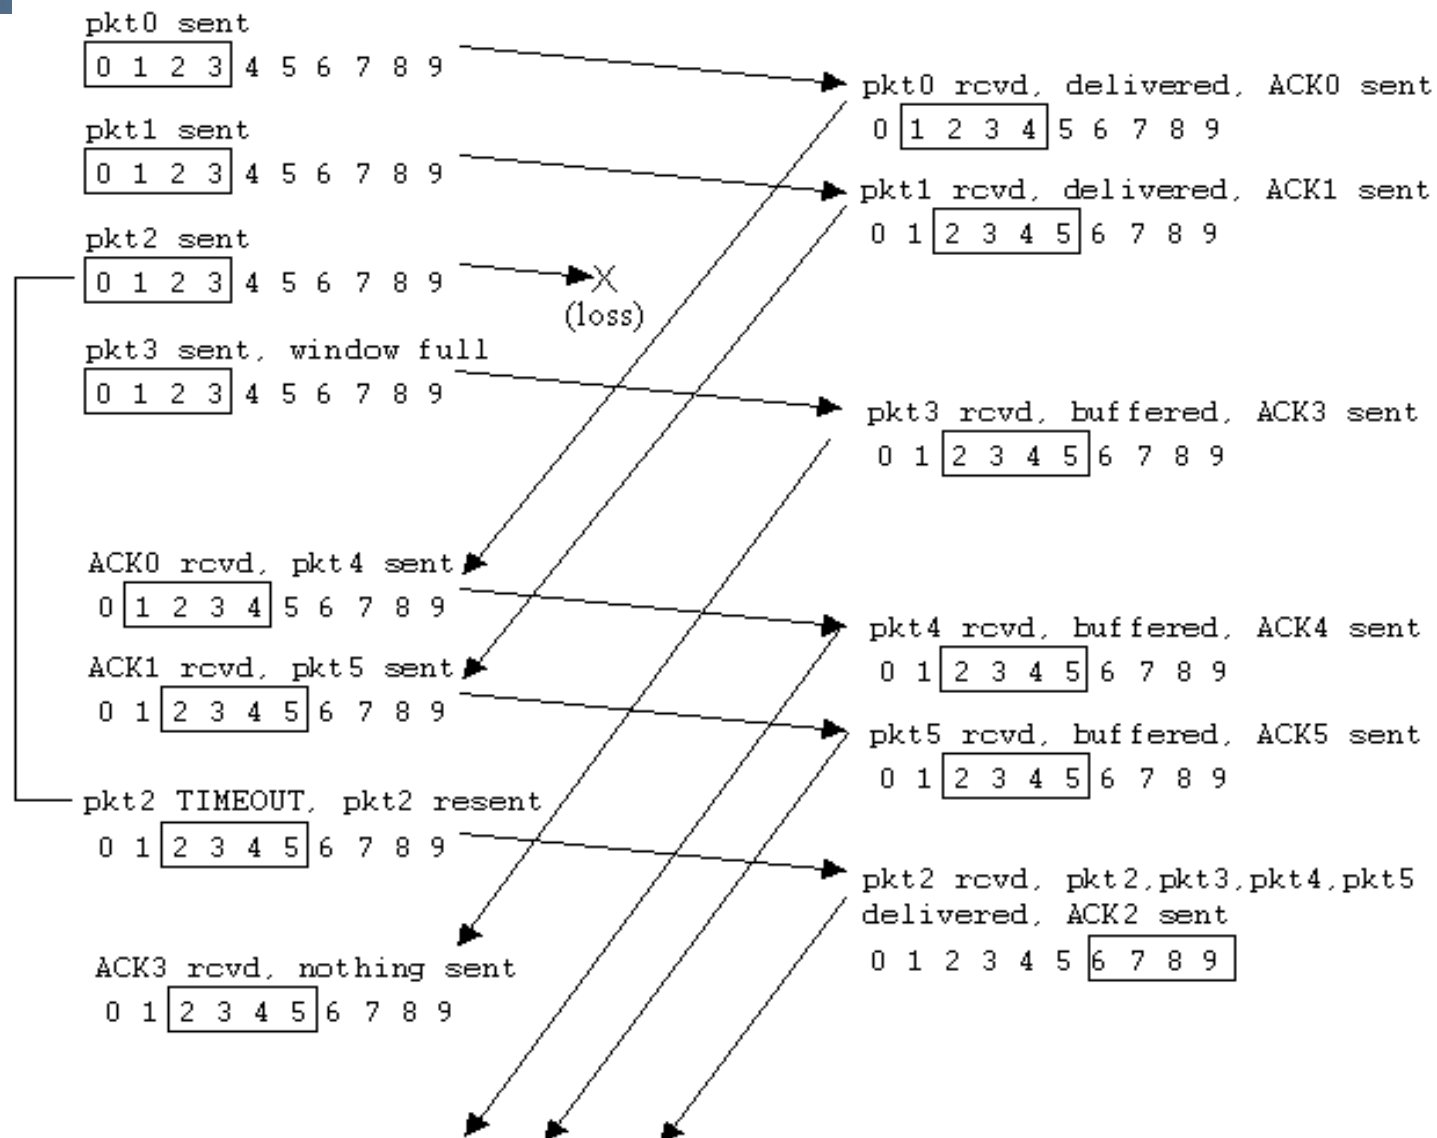
- send ACK(n)
- if (out-of-order)  
mark and buffer pkt  $n$   
else /\*in-order\*/  
deliver any in-order packets

otherwise:

- ignore



# Selective Repeat in Action



## Discussion: Efficiency of Selective Repeat


---

- ❑ Assume window size  $W$
- ❑ Assume each packet is lost with probability  $p$
- ❑ On average, how many packets do we send for each data packet received?

# Selective Repeat: Seq# Size and Window Size


Example:

- ❑ seq #'s (2 bits):  
0, 1, 2, 3
- ❑ window size=3
- ❑ Error: incorrectly passes duplicate data as new.




# State Invariant: Window Location

- Go-back-n (GBN)




- Selective repeat (SR)




# Window Location

Q: what relationship between seq # size and window size?

- Go-back-n (GBN)



- Selective repeat (SR)

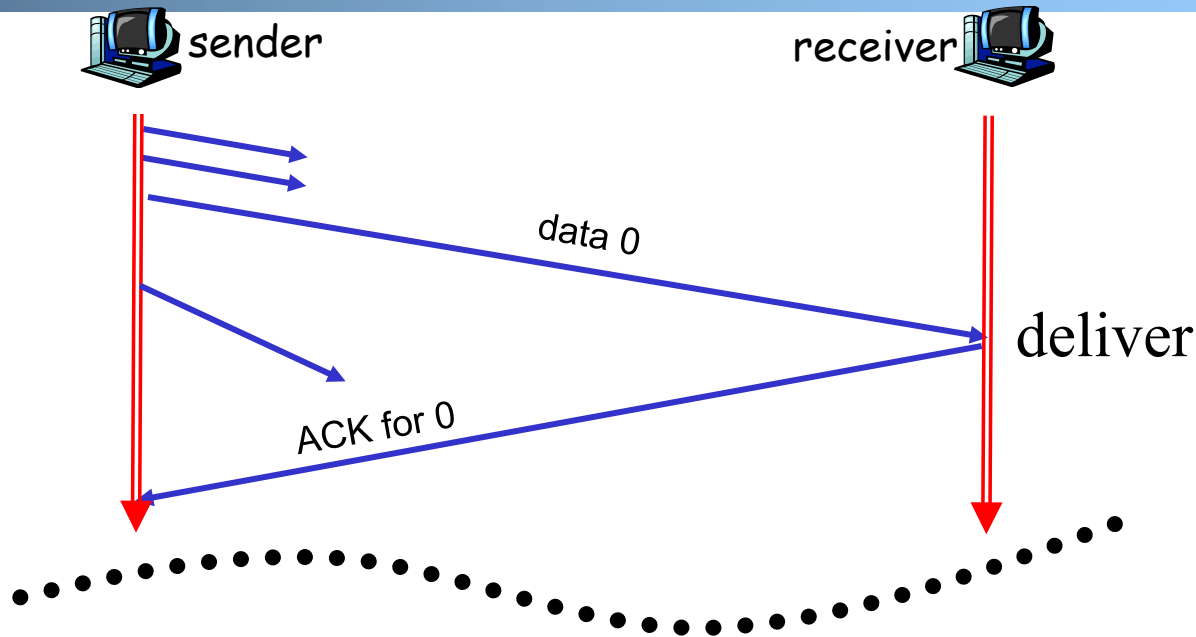


# Sliding Window Protocols: Go-back-n and Selective Repeat


	Go-back-n	Selective Repeat
data bandwidth: sender to receiver (avg. number of times a pkt is transmitted)	Less efficient $\frac{1-p+pw}{1-p}$	More efficient $\frac{1}{1-p}$
ACK bandwidth (receiver to sender)	More efficient	Less efficient
Relationship between M (the number of seq#) and W (window size)	$M > W$	$M \geq 2W$
Buffer size at receiver	1	W
Complexity	Simpler	More complex

$p$ : the loss rate of a packet; M: number of seq# (e.g., 3 bit M = 8); W: window size

# Question: What is Initial Seq# and When to Accept First Packet?



# Question: What is Initial Seq# and When to Accept First Packet?



Discussion: Condition for receiver to deliver first data?



# Outline

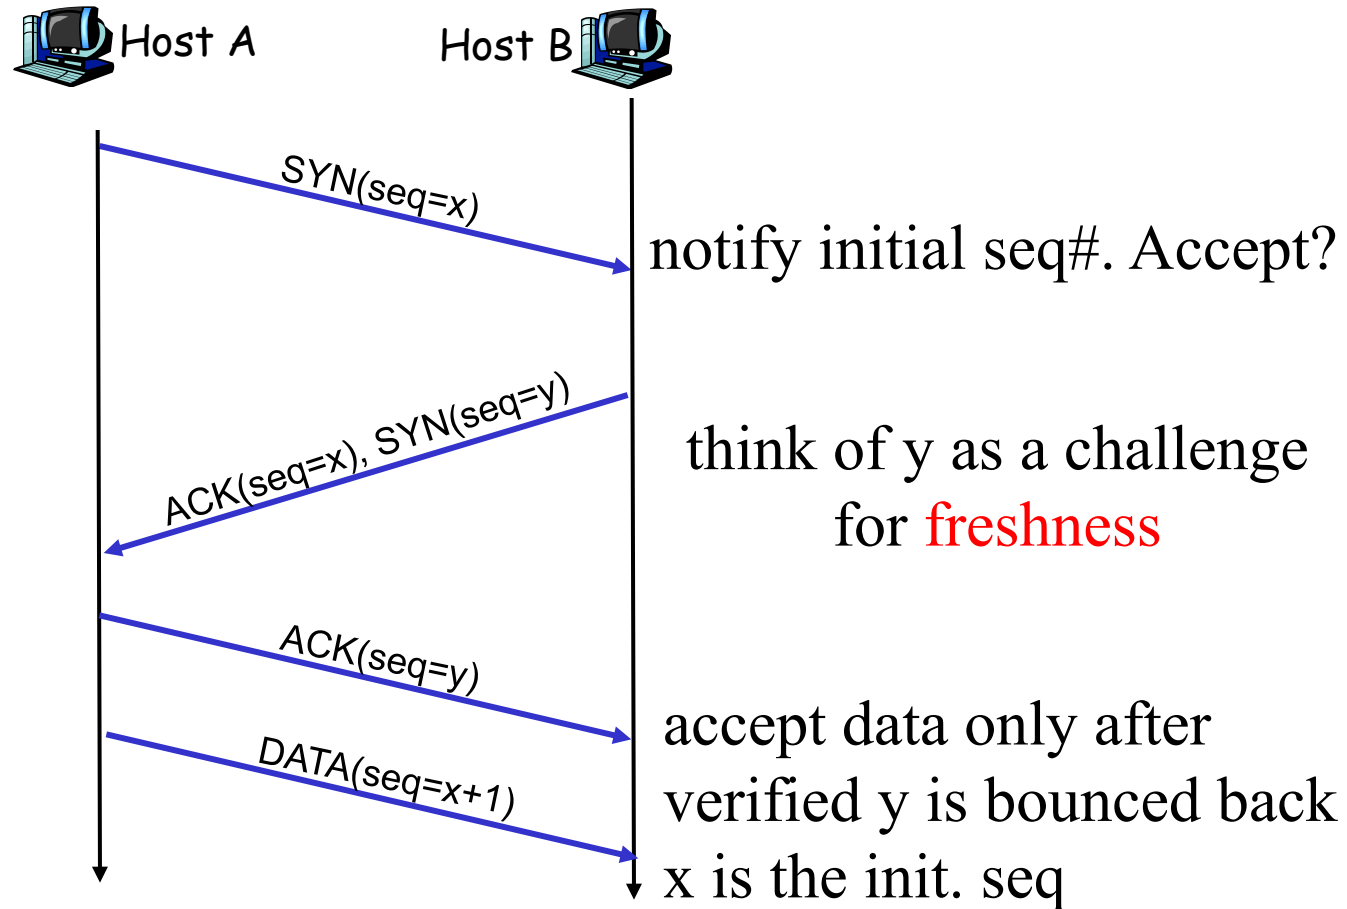
---

- Admin and recap

- Reliable data transfer


- perfect channel
- channel with bit errors
- channel with bit errors and losses
- sliding window: reliability with throughput
- connection management

## Three Way Handshake (TWH) [Tomlinson 1975]




SYN: indicates connection setup

# Scenarios with Duplicate Request/SYN Attack



# Scenarios with Duplicate Request/SYN Attack

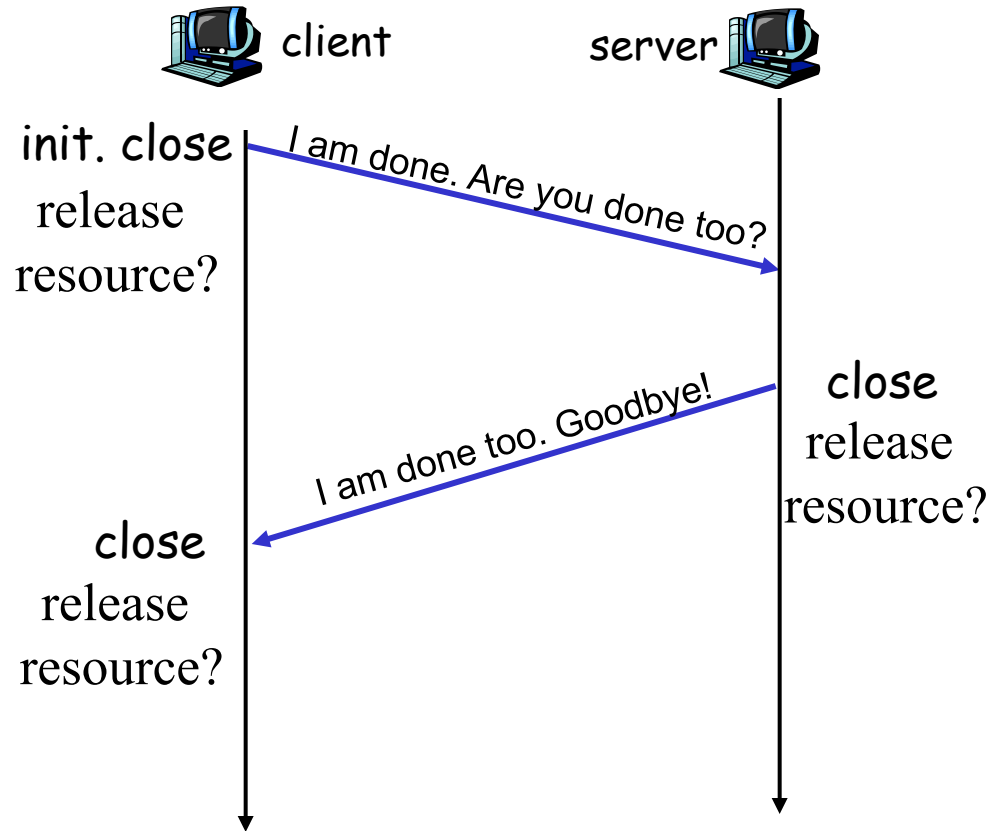


# Make "Challenge y" Robust


- ❑ To avoid that "SYNC ACK y" comes from reordering and duplication
  - for each connection (sender-receiver pair), ensuring that two identically numbered packets are never outstanding at the same time
    - network bounds the life time of each packet
    - a sender will not reuse a seq# before it is sure that all packets with the seq# are purged from the network
    - seq. number space should be large enough to not limit transmission rate

# Connection Close

- Why connection close?
  - so that each side can release resource and remove state about the connection (do not want dangling socket)



# General Case: The Two-Army Problem



The gray (blue) armies need to agree on whether or not they will attack the white army. They achieve agreement by sending messengers to the other side. If they both agree, attack; otherwise, no. Note that a messenger can be captured!

Discussion: Potential approaches to close state?

# Four Way Teardown

