
Network Layer: intro; Distance Vector Protocols

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

11/27/2018

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control-plane
 - Routing

Admin

❑ Assignment four meeting

- Today:
 - 2:30-3:30 pm
 - 5:00-6:30 pm
- Wednesday

❑ Exam 2 date?

Recap: BW Allocation Framework

$$\begin{array}{ll}
 \max & \sum_{f \in F} U_f(x_f) \\
 \text{subject to} & \sum_{f: f \text{ uses link } l} x_f \leq c_l \text{ for any link } l \\
 \text{over} & x \geq 0
 \end{array}$$

- Forward engineering: systematically design of
 - objective function
 - distributed alg to achieve objective
- Science/reverse engineering: what do TCP/Reno, TCP/Vegas achieve?

Objective	Allocation (x1, x2, x3)		
TCP/Reno	0.26	0.74	0.74
TCP/Vegas	1/3	2/3	2/3
Max throughput	0	1	1
Max-min	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{2}$
Max sum $\log(x)$	1/3	2/3	2/3
Max sum of $-1/(\text{RTT}^2 x)$	0.26	0.74	0.74

Recap: Systematic Derivation of Objective Function


□ NBS axioms

- Pareto optimality
- symmetry
- invariance of linear transformation
- independence of irrelevant alternatives

□ NBS solution

- the rate allocation point is the feasible point which maximizes


$$x_1 x_2 \cdots x_F$$



Recap: Systematic Derivation of Alg: Foundation (Strong Dual Theorem)

$$\begin{array}{ll} \max & f(x) \\ \text{subject to} & g(x) \leq 0 \\ \text{over} & x \in S \end{array}$$

$f(x)$ concave
 $g(x)$ linear
 S is a convex set



$$D(q) = \max_{x \in S} (f(x) - qg(x))$$

$-D(q)$ is called the dual;
 q (≥ 0) are called prices in economics

Recap: Primal-Dual Decomposition of Network-Wide Resource Allocation

□ SYSTEM(U):

$$\begin{array}{ll}\max & \sum_{f \in F} U_f(x_f) \\ \text{subject to} & \sum_{f: f \text{ uses link } l} x_f \leq c_l \text{ for any link } l \\ \text{over} & x \geq 0\end{array}$$

□ USER_f:

$$\begin{array}{ll}\max_{x_f} & U_f(x_f) - x_f p_f \\ \text{over} & x_f \geq 0\end{array}$$


□ NETWORK:

$$\min_{q \geq 0} \tilde{D}(q) = \sum_l q_l (c_l - \sum_{f: f \text{ uses } l} x_f)$$

TCP/Reno Dynamics

$$\Delta x_f \propto U'_f(x_f) - p_f$$

$$\Delta x = \frac{RTT}{2} x^2 \left(\frac{2}{x^2 RTT^2} - p \right)$$


$$U'_f(x_f) - p_f$$


$$\Rightarrow U'_f(x_f) = \left(\frac{\sqrt{2}}{x_f RTT} \right)^2 \quad \Rightarrow U_f(x_f) = -\frac{2}{RTT^2 x_f}$$

TCP/Vegas Dynamics

$$\Delta x_f \propto U'_f(x_f) - p_f$$

$$\Delta x = \frac{x}{RTT} \left(\frac{\alpha}{x} - (RTT - RTT_{min}) \right)$$

$$U'_f(x_f) - p_f$$


$$\Rightarrow U'_f(x_f) = \frac{\alpha}{x}$$

$$\Rightarrow U_f(x_f) = \alpha \log(x_f)$$

Outline

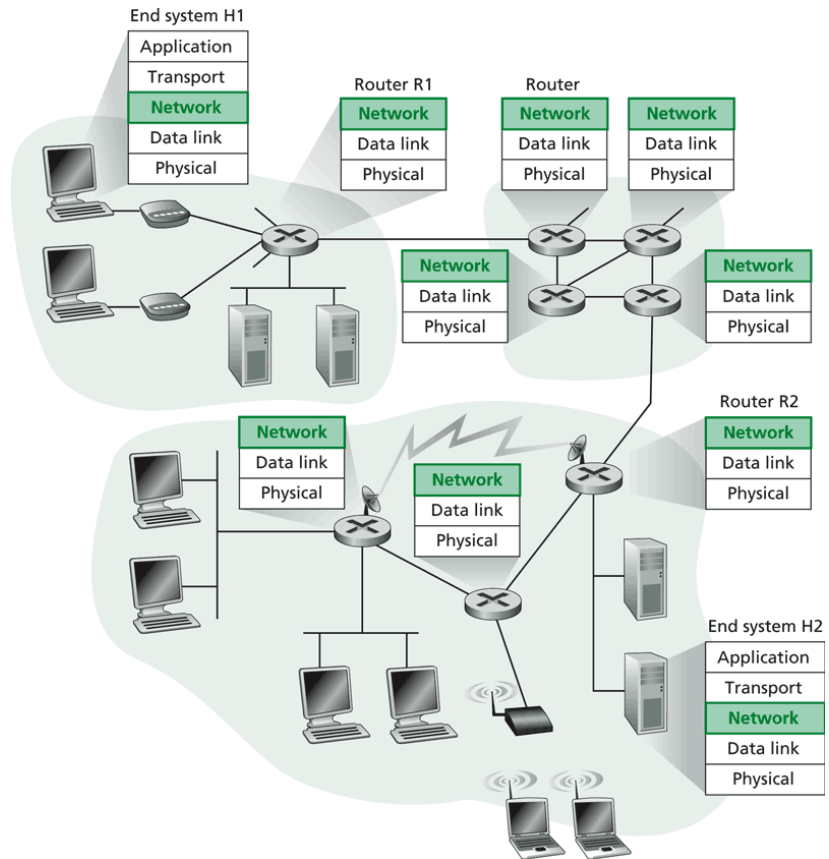
- Admin and recap
- *Network overview*

Network Layer

- ❑ Transport packets from source to destination
- ❑ Network layer in *every* host, router

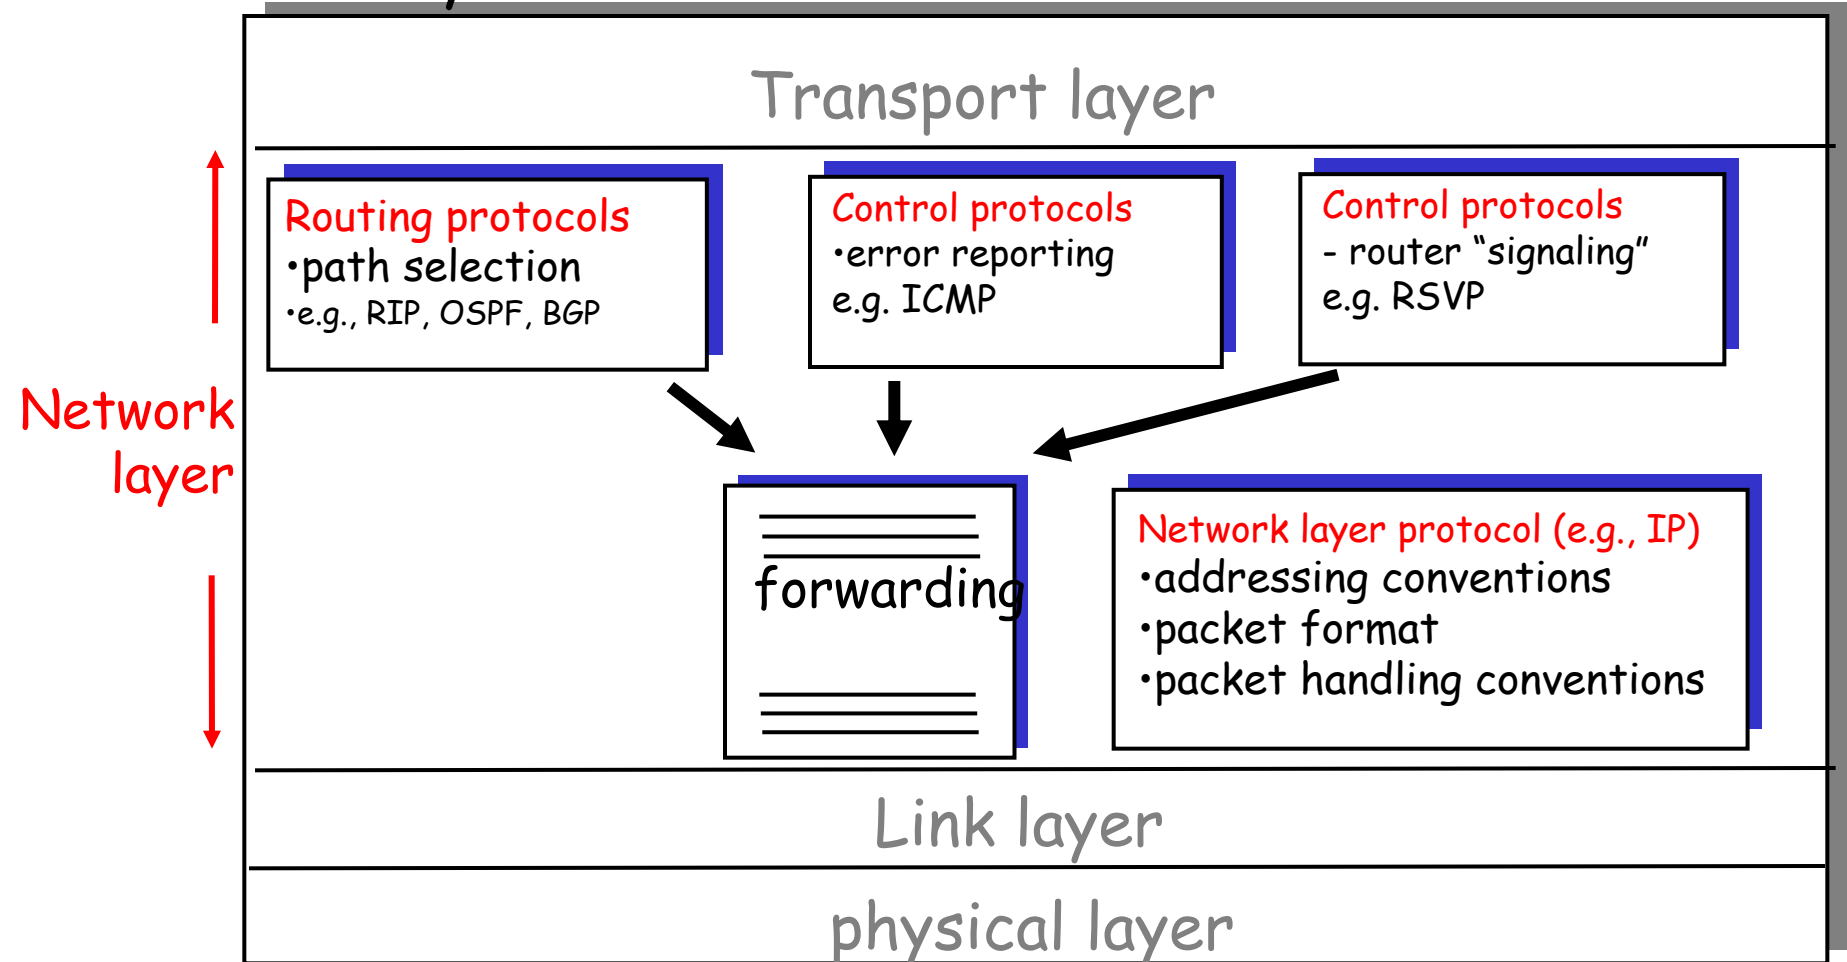
Basic functions:

- ❑ inter-networking (e.g., fragmentation/assembly)
- routing (determine route(s) taken by packets of a flow), and forwarding (move the packets along the route(s))



Current Internet Network Layer

Network layer functions:




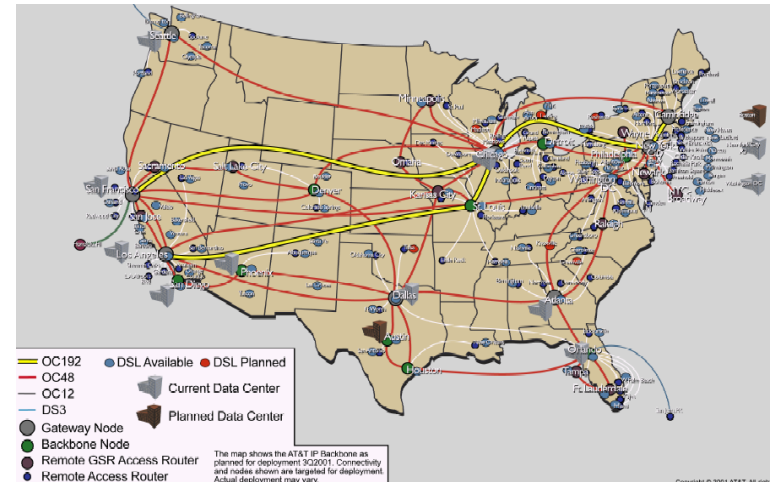
Routing: Overview

Routing

Goal: determine "good" paths (sequences of routers) thru networks from source to dest.

Graph abstraction for the routing problem:

- graph nodes are routers
- graph edges are physical links
 - links have properties: delay, capacity, \$ cost
- compute path on graph



Network Layer: Complexity Factors/Objectives

❑ For network providers

- efficiency of routes
- policy control on routes
- scalability

❑ For users

- quality of services, e.g.,
 - guaranteed bandwidth?
 - preservation of inter-packet timing (no jitter)?
 - loss-free delivery?
 - in-order delivery?

❑ Users and network may interact

Routing Design Space

- Robustness
- Optimality
- Simplicity

- Routing has a large design space
 - who decides routing?
 - source routing: end hosts make decision
 - network routing: networks make decision
 - how many paths from source s to destination d ?
 - multi-path routing
 - single path routing
 - what does routing compute?
 - network cost minimization
 - QoS aware
 - will routing adapt to network traffic demand?
 - adaptive routing
 - static routing
 - ...


Routing Design Space: Internet

- Robustness
- Optimality
- Simplicity

- Routing has a large design space
 - who decides routing?
 - source routing: end hosts make decision
 - network routing: networks make decision
 - (applications such as overlay and p2p are trying to bypass it)
 - what does routing compute?
 - network cost minimization (shortest path)
 - QoS aware
 - how many paths from source s to destination d ?
 - multi-path routing
 - single path routing (with small amount of multipath)
 - will routing adapt to network traffic demand?
 - adaptive routing
 - static routing (mostly static; adjust in larger timescale)
 - ...

Basic Formulation

- Assign link weights
- Compute shortest path



Example: Cisco Proprietary Recommendation on Assigning Link Costs

□ Link metric:

- metric = $[K1 * \text{bandwidth}^{-1} + (K2 * \text{bandwidth}^{-1}) / (256 - \text{load}) + K3 * \text{delay}] * [K5 / (\text{reliability} + K4)]$

By default, $k1=k3=1$ and $k2=k4=k5=0$. The default composite metric for EIGRP, adjusted for scaling factors, is as follows:


$$\text{EIGRP}_{\text{metric}} = 256 \times \{ [10^7 / \text{BW}_{\text{min}}] + [\text{sum_of_delays}] \}$$

BW_{min} is in kbps and the sum of delays are in 10s of microseconds.

EIGRP : Enhanced Interior Gateway Routing Protocol

Example: EIGRP Link Cost

- ❑ The bandwidth and delay for an Ethernet interface are 10 Mbps and 1 ms, respectively.
- ❑ The calculated EIGRP metric is as follows:
 - $256 \times [10^7 / \text{BWks} + \text{delayin10us}]$
 - $= 256 \times [10^7 / 10,000 + 100]$
 - $= 256 \times [1000 + 100]$
 - $= 256,000 + 25,600$
 - $= 281,600$



Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Distributed routing computation

Why Study?


- Just as Dijkstra's Shortest Path algorithm is among the most classical algorithms in algorithm design, distributed shortest path protocols provide many insights in distributed protocol design.
- Please learn not only the protocols, but also the techniques (convergence, global invariants, ...)

Outline


- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation

Basic Routing Computation Setting

- Setting: **static (positive)** costs assigned to network links
 - The static link costs may be adjusted in a longer time scale: this is called **traffic engineering**
- Goal: distributed computing to compute the **shortest path** from a source to a destination
 - Conceptually, runs for each destination separately



Intuition



$d_i \leq d_j + d_{ij}$, for
each neighbor j

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

Understanding Shortest Path and an Exercise of Primal-Dual

$$\max d_s - d_D$$

for any edge $i \rightarrow j$: $d_i \leq d_j + d_{ij}$

$$d_i \geq 0$$

$$\text{Dual: } D(x) = \max(d_s - d_D - \sum x_{ij}(d_i - d_j - d_{ij}))$$

$$= \sum x_{ij} d_{ij}$$

x_{ij} is a flow from s to D

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - *Distributed distance vector protocols*


Distance Vector Routing: Basic Idea

- Based on Bellman-Ford equation: At node i , the basic update rule

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

where


- d_i denotes the distance estimation from i to the destination,
- $N(i)$ is set of neighbors of node i , and
- d_{ij} is the distance of the direct link from i to j



Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - *distributed distance vector protocols*
 - *synchronous Bellman-Ford (SBF)*

Synchronous Bellman-Ford (SBF)




□ Nodes update in rounds:

- there is a global clock;
- at the beginning of each round, each node sends its estimate to all of its neighbors;
- at the end of the round, updates its estimation

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$


$d(0)?$



Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control-plane path
 - Routing
 - Link weights assignment
 - Routing computation
 - *distributed distance vector protocols*
 - *synchronous Bellman-Ford (SBF)*
 - SBF/∞


SBF/ ∞



□ Initialization (time 0):

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ \infty & \text{otherwise} \end{cases}$$

Example



Consider D as destination; $d(t)$ is a vector consisting of estimation of each node at round t

	A	B	C	E	D
$d(0)$	∞	∞	∞	∞	0
$d(1)$	∞	∞	2	2	0
$d(2)$	12	3	2	2	0
$d(3)$	10	3	2	2	0
$d(4)$	10	3	2	2	0

Observation: $d(0) \geq d(1) \geq d(2) \geq d(3) \geq d(4) = d^*$

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

A Nice Property of SBF: Monotonicity

- Consider two configurations $d(t)$ and $d'(t)$
- If $d(t) \geq d'(t)$
 - i.e., each node has a higher estimate in one scenario (d) than in another scenario (d'),
- then $d(t+1) \geq d'(t+1)$
 - i.e., each node has a higher estimate in d than in d' after one round of synchronous update.

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

Correctness of SBF/ ∞

- Claim: $d_i(h)$ is the length $L_i(h)$ of a shortest path from i to the destination using $\leq h$ hops
 - base case: $h = 0$ is trivially true
 - assume true for $\leq h$,
i.e., $L_i(h) = d_i(h)$, $L_i(h-1) = d_i(h-1)$, ...

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

Correctness of SBF/ ∞

□ consider $\leq h+1$ hops:

$$\begin{aligned} L_i(h+1) &= \min(L_i(h), \min_{j \in N(i)} (d_{ij} + L_j(h))) \\ &= \min(d_i(h), \min_{j \in N(i)} (d_{ij} + d_j(h))) \\ &= \min(d_i(h), d_i(h+1)) \end{aligned}$$

since $d_i(h) \leq d_i(h-1)$


$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h)) \leq \min_{j \in N(i)} (d_{ij} + d_j(h-1)) = d_i(h)$$

$$L_i(h+1) = d_i(h+1)$$

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - *Distributed distance vector protocols*
 - *synchronous Bellman-Ford (SBF)*
 - SBF/ ∞
 - SBF/-1


SBF at another Initial Configuration: SBF/-1



□ Initialization (time 0):

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ -1 & \text{otherwise} \end{cases}$$

Example



Consider D as destination

	A	B	C	E	D
d(0)	-1	-1	-1	-1	0
d(1)	6	0	0	2	0
d(2)	7	1	1	2	0
d(3)	8	2	2	2	0
d(4)	9	3	3	2	0
d(5)	10	3	3	2	0
d(6)	10	3	3	2	0

Observation: $d(0) \leq d(1) \leq d(2) \leq d(3) \leq d(4) \leq d(5) = d(6) = d^*$

Correctness of SBF/-1

- SBF/-1 converges due to monotonicity
- Remaining question:
 - Can we guarantee that SBF/-1 converges to shortest path?

Correctness of SBF/-1

- Common between SBF/ ∞ and SBF/-1: they solve the Bellman equation

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

where $d_D = 0$.

- We have proven SBF/ ∞ is the shortest path solution.
- SBF/-1 computes shortest path if Bellman equation has a unique solution.



$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

Uniqueness of Solution to BE

□ Assume another solution d , we will show that $d = d^*$

case 1: we show $d \geq d^*$

Since d is a solution to BE, we can construct paths as follows: for each i , pick a j which satisfies the equation; since d^* is shortest, $d \geq d^*$



$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

Uniqueness of Solution to BE

Case 2: we show $d \leq d^*$

assume we run SBF with two initial configurations:

- one is d
- another is $SBF/\infty (d^\infty)$,

-> monotonicity and convergence of SBF/∞ imply that $d \leq d^*$

Summary: "Extreme" SBF Initial States

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ \infty & \text{otherwise} \end{cases}$$

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ -1 & \text{otherwise} \end{cases}$$

- Nice properties of both cases
 - Monotonicity
 - Convergence

Discussion

- Will SBF converge under any non-negative initial conditions?

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - *Distributed distance vector protocols*
 - synchronous Bellman-Ford (SBF)
 - *asynchronous Bellman-Ford (ABF)*

Asynchronous Bellman-Ford (ABF)

- No notion of global iterations
 - each node updates at its own pace
- Asynchronously each node i computes


$$d_i = \min_{j \in N(i)} (d_{ij} + d_j^i)$$

using last received value d_j^i from neighbor j .

- Asynchronously node j sends its estimate to its neighbor i :
 - We assume that there is an upper bound on the delay of estimate packet

ABF: Example

Below is just one step! The protocol repeats forever!




		distance tables from neighbors			computation			E's routing table	distance table E sends to its neighbors	
destinations	$d_E()$	A	B	D	A	B	D			
	A	0	7	∞	10	15	∞	A: 10		A: 10
	B	7	0	∞	17	8	∞	B: 8		B: 8
	C	∞	1	2	∞	9	4	D: 4		C: 4
	D	∞	∞	0	∞	∞	2	D: 2		D: 2
		10	8	2						E: 0

next hop distance

Asynchronous Bellman-Ford (ABF)


- ABF will eventually converge to the shortest path
 - links can go down and come up - but if topology is stabilized after some time t and connected, ABF will eventually converge to the shortest path !

ABF Convergence Proof Complexity: Complex System State



What is system state?

System State



three types of distance state from node j :

- d_j : current distance estimate state at node j
- d_j^i : last d_j that neighbor i received
- \bar{d}_j^i : those d_j that are still in transit to neighbor i

ABF Convergence Proof: The Sandwich Technique

□ Basic idea:

- bound system state using extreme states

□ Extreme states:

- SBF/∞ ; call the sequence $U()$
- $SBF/-1$; call the sequence $L()$

ABF Convergence

- Consider the time when the topology is stabilized at time 0
- $U(0)$ and $L(0)$ provide upper and lower bounds at time 0 on all corresponding elements of states
 - $L_j(0) \leq d_j \leq U_j(0)$ for all d_j state at node j
 - $L_j(0) \leq d_j^i \leq U_j(0)$
 - $L_j(0) \leq \text{update messages } d_j^i \leq U_j(0)$


ABF Convergence

□ d_j

- after at least one update at node j :
 d_j falls between
 $L_j(1) \leq d_j \leq U_j(1)$

□ d_j^i :

- eventually all d_j^i that
are only bounded by
 $L_j(0)$ and $U_j(0)$ are
replaced with in
 $L_j(1)$ and $U_j(1)$



Distributed, Asynchronous, Routing Protocol: Summary of Features

❑ Distributed

- each node communicates its routing table to its directly-attached neighbors

❑ Iterative

- continues periodically or when link changes, e.g. detects a link failure

❑ Asynchronous

- nodes need *not* exchange info/iterate in lock step!

❑ Convergence

- in finite steps, independent of initial condition if network is connected

Distributed, Asynchronous, Routing Protocol: Summary of Analytical Technique


- Tool box: a key technique for analyzing convergence (**liveness**) of distributed protocols: **monotonicity** and the **bounding-box (sandwich) theorem**
 - Consider two configurations $d(t)$ and $d'(t)$:
 - if $d(t) \leq d'(t)$, then $d(t+1) \leq d'(t+1)$
 - Identify two extreme configurations to sandwich any real configurations

Outline

- ❑ Admin and recap
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - Distance vector protocols (distributed computing)
 - synchronous Bellman-Ford (SBF)
 - asynchronous Bellman-Ford (ABF)
 - *properties of DV*

Properties of Distance-Vector Algorithms

- Good news propagate fast



Properties of Distance-Vector Algorithms

❑ Bad news propagate slowly


	A	B	C	D	E	
	●	●	●	●	●	
		1	2	3	4	Initially
A-B link down		3	2	3	4	After 1 exchange
		3	4	3	4	After 2 exchanges
		5	4	5	4	After 3 exchanges
		5	6	5	6	After 4 exchanges
		7	6	7	6	After 5 exchanges
		7	8	7	8	After 6 exchanges
			⋮			
		∞	∞	∞	∞	

❑ This is called the *counting-to-infinity* problem

❑ Q: what causes counting-to-infinity?

Counting-To-Infinity is Because of Routing Loop

- Counting-to-infinity is caused by a routing loop, which is a **global state** (consisting of the nodes' local states) at a global moment (observed by an oracle) such that there exist nodes A, B, C, ... E such that A (locally) thinks B as next hop, B thinks C as next hop, ... E thinks A as next hop



Discussion


- ❑ Why avoid routing loops is hard?
- ❑ Any proposals to avoid distributed routing loops?

Outline

- ❑ Admin and recap
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - Distance vector protocols (distributed computing)
 - synchronous Bellman-Ford (SBF)
 - asynchronous Bellman-Ford (ABF)
 - properties of DV
 - distributed protocols w/ safety (loop prevention)
 - *reverse poison/split horizon*

The Reverse-Poison (Split-horizon) Hack

If the path to dest is through neighbor h, report ∞ to neighbor h for dest.



		distance tables from neighbors			computation			E's distance table	distance table E sends to its neighbors		
destinations	D ^E ()	A	B	D	A	B	D		To A	To B	To D
	A	0	7	∞	1	15	∞	1, A	A: ∞	A: 1	A: 1
	B	7	0	∞	8	8	∞	8, B	B: 8	B: ∞	B: 8
	C	∞	1	2	∞	9	4	4, D	C: 4	C: 4	C: ∞
	D	∞	∞	0	∞	∞	2	2, D	D: 2	D: 2	D: ∞
		1	8	2					E: 0	E: 0	E: 0

$c(E,A)$ $c(E,B)$ $c(E,D)$

distance through neighbor


Reverse-Poison Example

r Exercise: Can Reverse-poison guarantee no loop for this network?

A	B	C	D	E	
●	●	●	●	●	
	1	2	3	4	Initially
	3	2	3	4	After 1 exchange
	3	4	3	4	After 2 exchanges
	5	4	5	4	After 3 exchanges
	5	6	5	6	After 4 exchanges
	7	6	7	6	After 5 exchanges
	7	8	7	8	After 6 exchanges
		⋮			
	∞	∞	∞	∞	

DV+RP => RIP (Routing Information Protocol)

- ❑ Included in BSD-UNIX Distribution in 1982
- ❑ Link cost: 1
- ❑ Distance metric: # of hops
- ❑ Distance vectors
 - exchanged every 30 sec via Response Message (also called **advertisement**) using UDP
 - each advertisement: route to up to 25 destination nets



RIP: Link Failure and Recovery


If no advertisement heard after 180 sec -->
neighbor/link declared dead

- routes via neighbor invalidated
- new advertisements sent to neighbors
- neighbors in turn send out new advertisements (if tables changed)
- link failure info quickly propagates to entire net
- reverse-poison used to prevent **ping-pong** loops
- set infinite distance = 16 hops (why?)

General Routing Loops and Reverse-poison


r Exercise: Can Reverse-poison guarantee no loop for this network?

A	B	C	D	E	
•	•	•	•	•	
	1	2	3	4	Initially
	3	2	3	4	After 1 exchange
	3	4	3	4	After 2 exchanges
	5	4	5	4	After 3 exchanges
	5	6	5	6	After 4 exchanges
	7	6	7	6	After 5 exchanges
	7	8	7	8	After 6 exchanges
	⋮				
	∞	∞	∞	∞	



General Routing Loops and Reverse-poison

- ❑ Reverse-poison removes two-node loops but may not remove more-node loops



- r Unfortunate timing can lead to a loop
 - When the link between C and D fails, C will set its distance to D as ∞
 - A receives the bad news (∞) from C, A will use B to go to D
 - A sends the news to C
 - C sends the news to B

Backup Slides




Routing Design Space: User-based, Multipath, Adaptive

- Robustness
- Optimality
- Simplicity

- Routing has a large design space
 - who decides routing?
 - source routing: end hosts make decision
 - network routing: networks make decision
 - how many paths from source s to destination d ?
 - multi-path routing
 - single path routing
 - what does routing compute?
 - network cost minimization
 - QoS aware
 - will routing adapt to network traffic demand?
 - adaptive routing
 - static routing
 - ...

User Optimal, Multipath, Adaptive

- User optimal: users pick the shortest routes (selfish routing)



Braess's paradox

Price of Anarchy

For a network with **linear** latency functions



total latency of user (selfish) routing for
given traffic demand

$$\leq \mathbf{4/3}$$

total latency of network optimal routing
for the traffic demand

Price of Anarchy

- r For any network with continuous, non-decreasing latency functions \rightarrow


total latency of user (selfish) routing
for given traffic demand

\leq


total latency of network optimal routing
for **twice** traffic demand

Assigning Link Weight: Dynamic Link Costs

- Assign link costs to reflect current traffic



Link costs reflect current traffic intensity



Solution: Link costs are a combination of current traffic intensity (dynamic) and topology (static). To improve stability, the static topology part should be large. Thus less sensitive to traffic; thus non-adaptive.